

OBJECTIFY

REAL-TIME OBJECT DETECTION
USING DEEP LEARNING WITH PYTHON & PYTORCH

PREPARED BY

Olwin Christian /U2141230136
Deepak Soni /U2141230282

Department of Computer
Science & Engineering, I.I.T.E,
Indus University, Ahmedabad

INTERNAL GUIDE

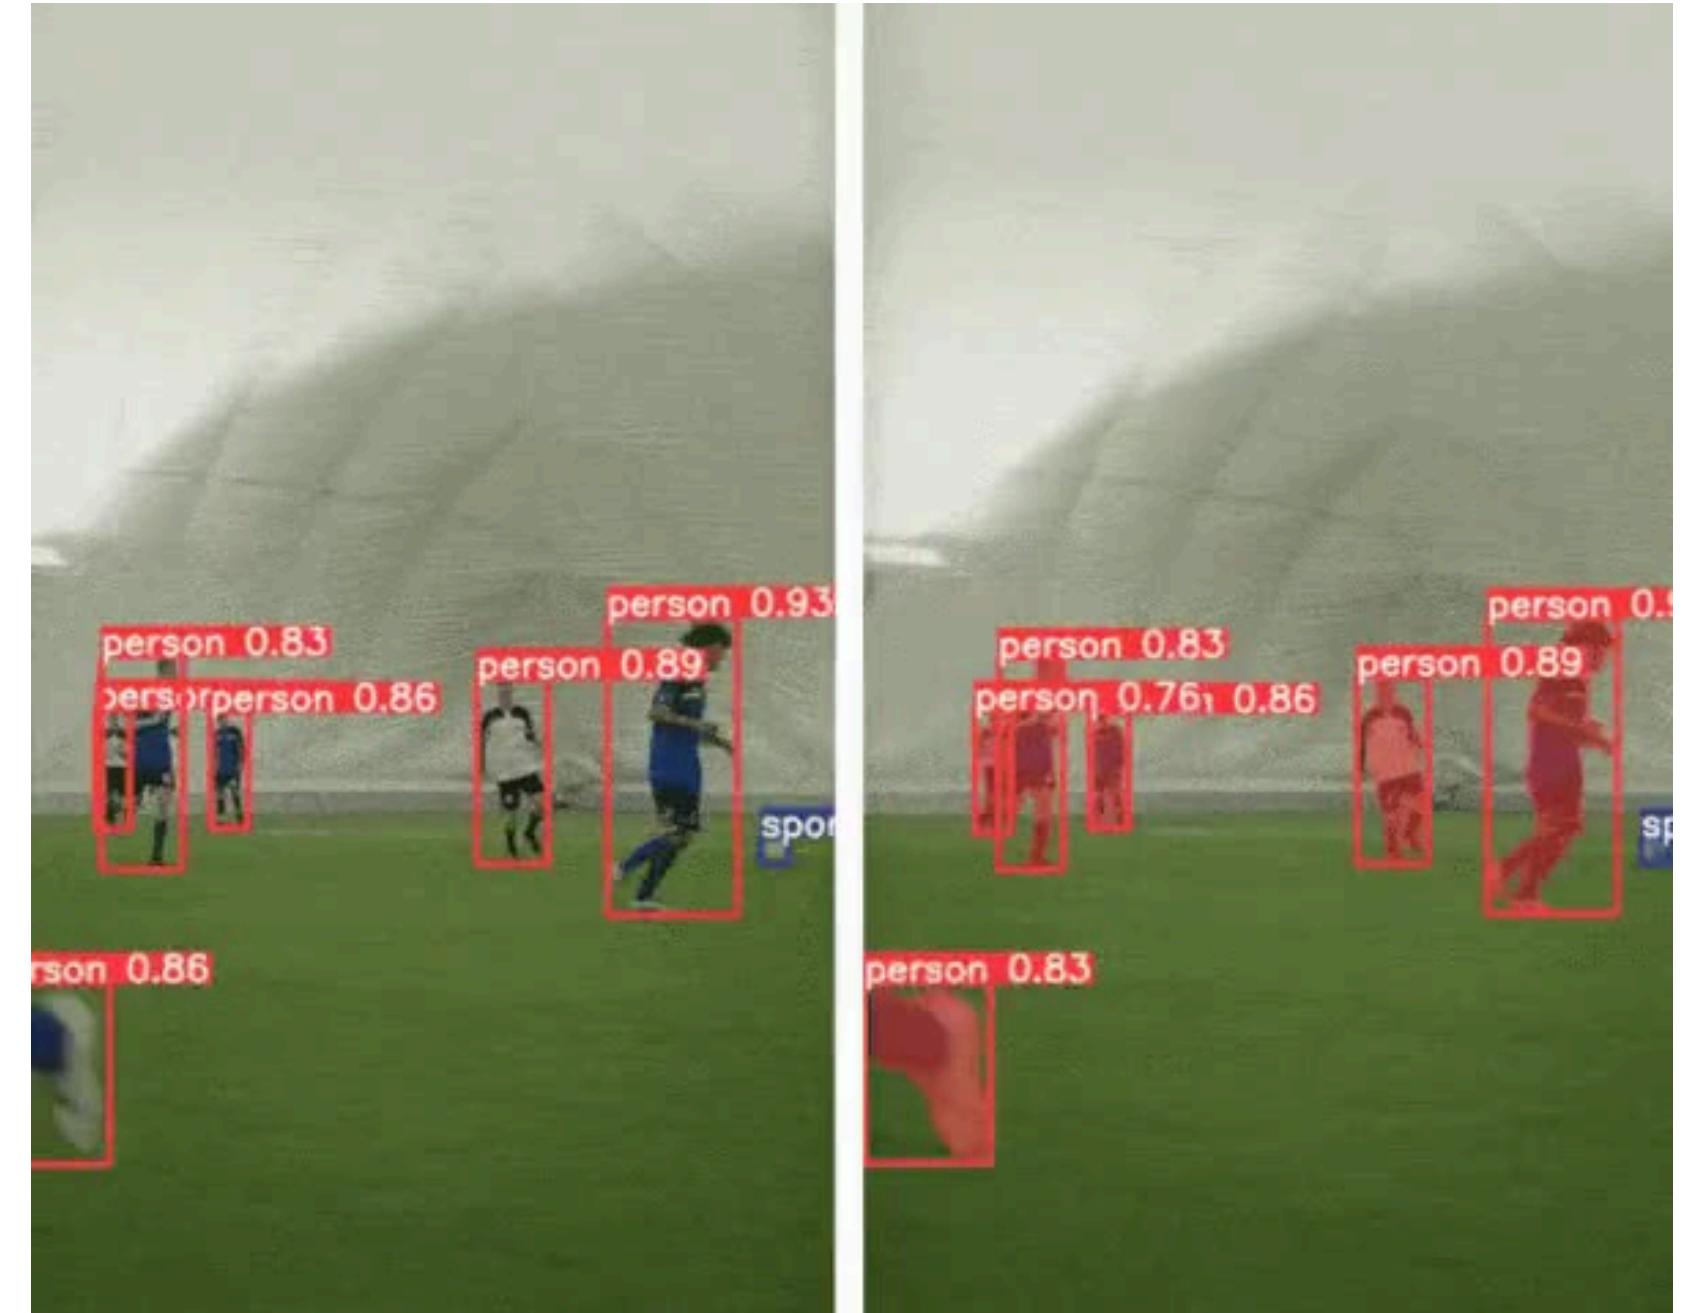
Dr. Sheetal Pandya
Assistant Professor

Department of Computer
Science & Engineering, I.I.T.E,
Indus University, Ahmedabad

PROJECT OVERVIEW

The Objectify project develops an advanced model to analyze football video clips, aiming to identify and track *players*, the *ball*, and *referees* in real time. Using deep learning frameworks like Faster R-CNN and YOLOv8, the project focuses on efficiently processing video frames to **deliver** precise object detection with bounding boxes and class labels.

Key elements include model training and testing on custom-annotated datasets in Google Colab using tools like PyTorch and Roboflow.



LITERATURE REVIEW

The Objectify project builds upon key research advancements in deep learning and computer vision for object detection.

- Deep Learning in Object Detection: Faster R-CNN, introduced by Shaoqing Ren et al., uses Region Proposal Networks integrated with CNNs to enhance both speed and accuracy, establishing a benchmark in object detection.
- YOLO (You Only Look Once): Joseph Redmon et al.'s YOLO framework changed real-time detection by evaluating images in a single step, achieving faster detection rates. YOLOv8, the latest version, improves performance for dynamic settings like sports.
- Transfer Learning for Object Detection: Research on transfer learning highlights its efficiency in training models on limited annotated datasets. Fine-tuning pre-trained models saves time and maintains high accuracy for specific detection tasks.

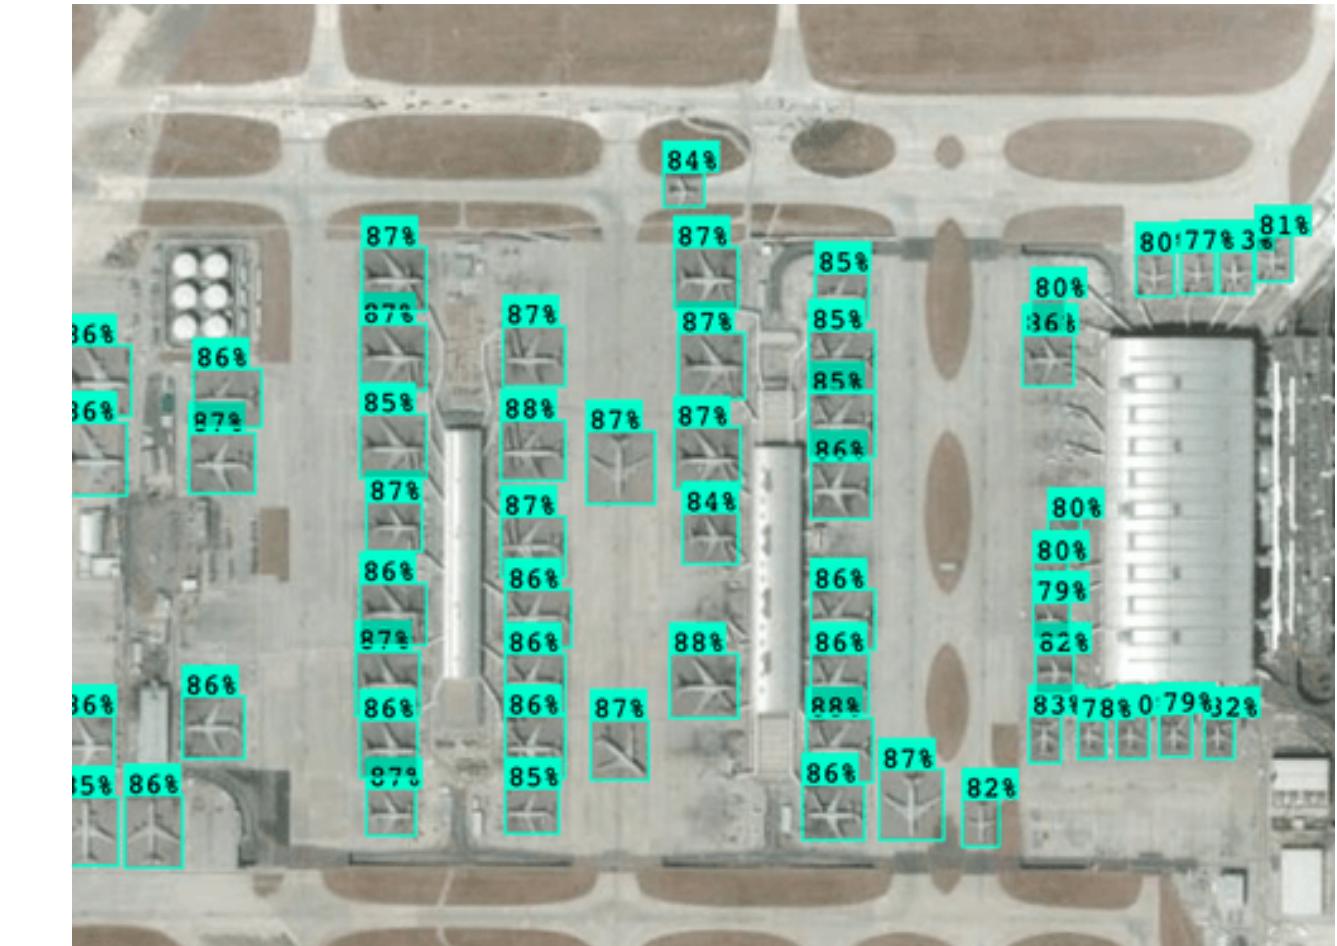
In summary, the Objectify project leverages these advancements to develop a real-time detection model optimized for football video analysis, guiding effective application of object detection in live, dynamic environments.

APPLICATIONS

The object detection model YOLOv8, developed in this project has several practical and impactful applications across various industries including, Autonomous Smart Solutions.



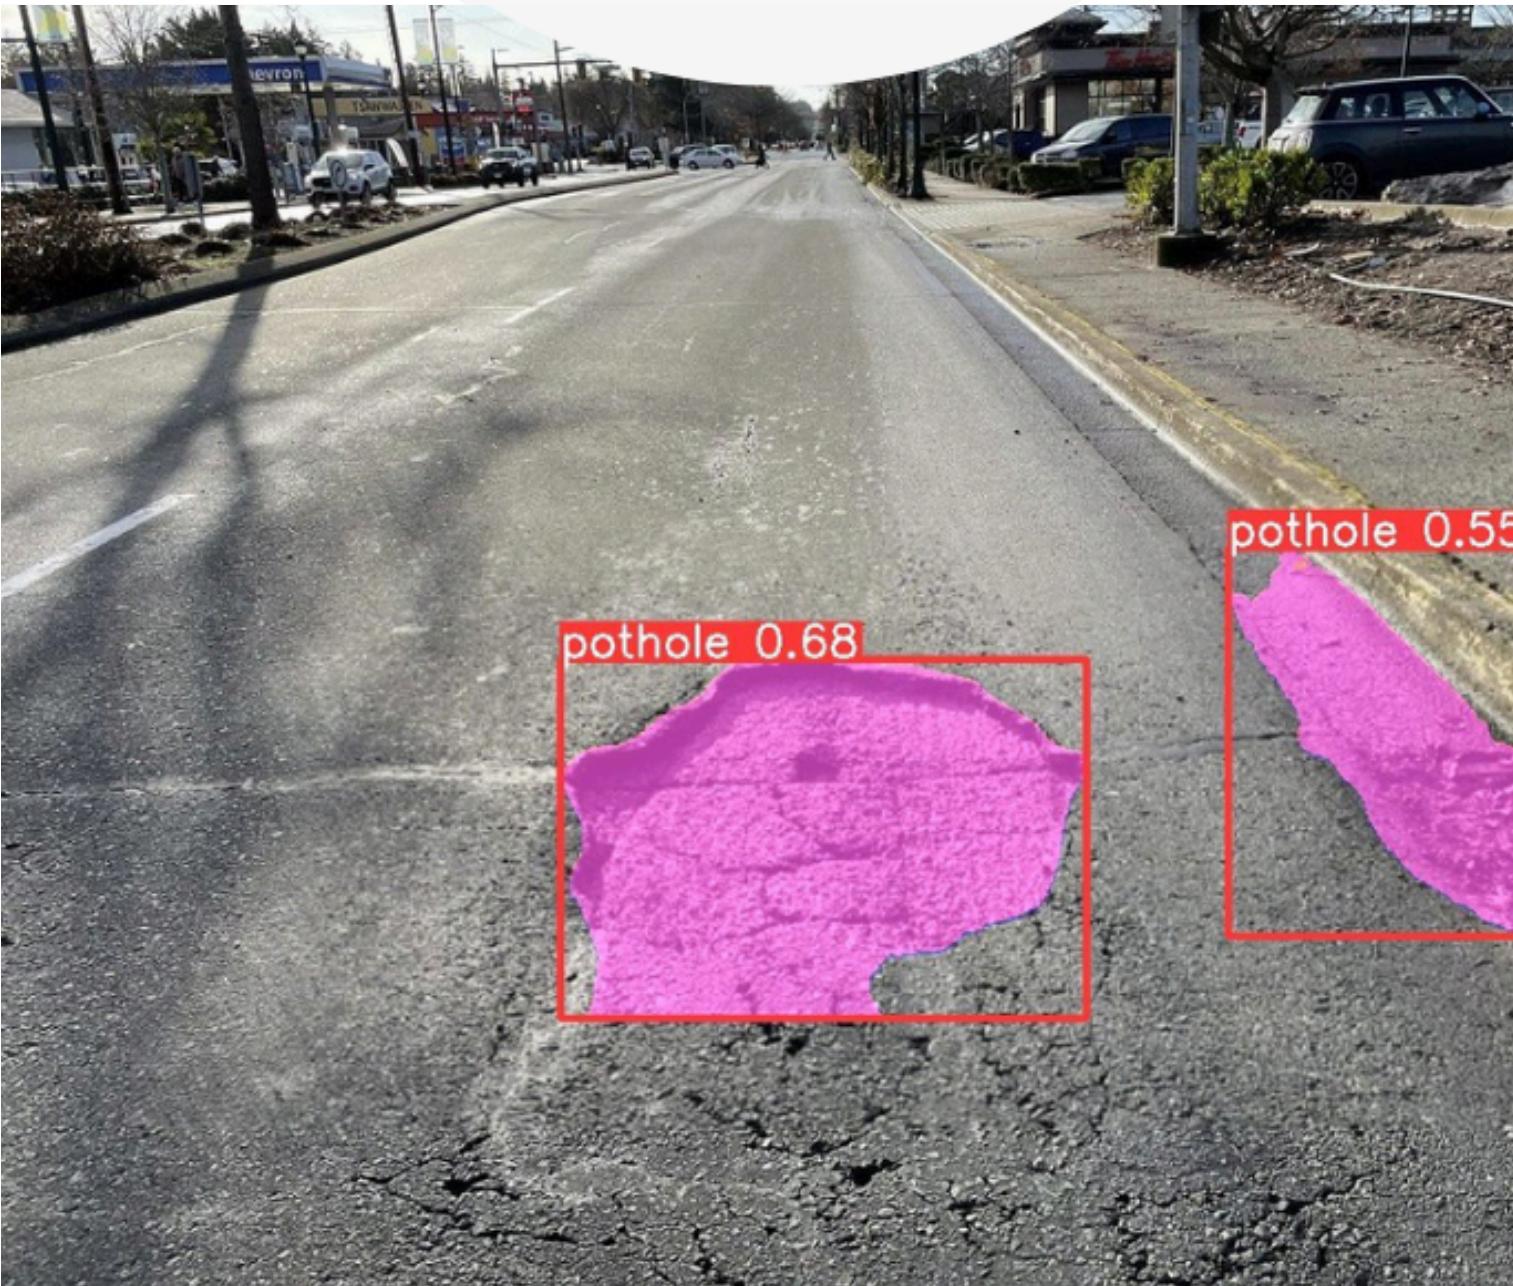
Autonomous Inventory Detection



Autonomous Airfield Surveillance

APPLICATIONS

The object detection model developed in this project has several practical and impactful applications across various industries including implementing, Smart City.

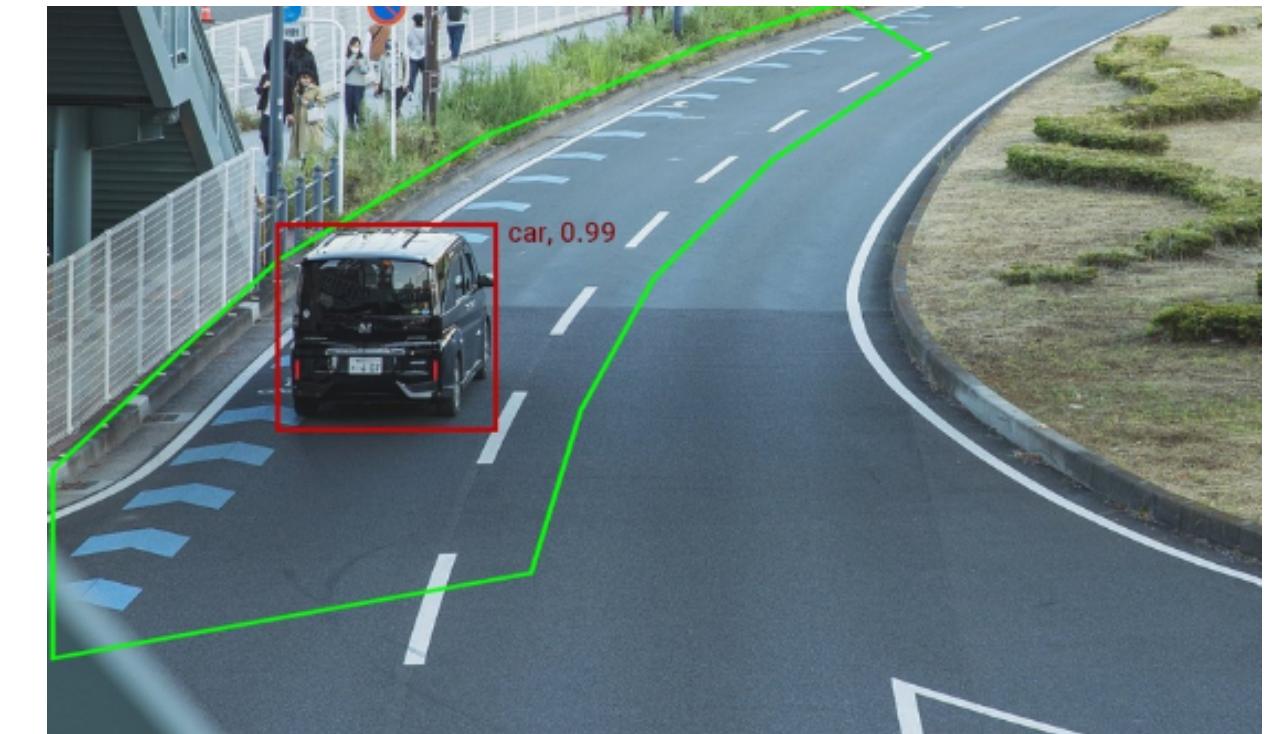


APPLICATIONS

The object detection model developed in this project has several practical and impactful applications across various industries including, Autonomous Detection and Identification.



License Plate Detection



Anomaly Vehicle Detection

COMPONENTS

Roboflow

Roboflow is used for **dataset management**, streamlining the **annotation** and **preparation** of the Football dataset for this project. It allows us to efficiently **annotate** images, **create training and validation splits**, and generate **customized datasets**, **improving** the training process for our object detection model.

Netron

Netron is a powerful **visualization** tool that enables us to **inspect** and **analyze** the **architecture** of our trained models. It provides a user-friendly **interface** to view model **layers**, **shapes**, and **parameters**, helping us understand the model's **structure** and **performance**. This insight is crucial for **fine-tuning** and **optimizing** our object detection algorithms.

TensorBoard

TensorBoard is a **visualization** tool that assists in **monitoring** and **analyzing** the training process. It provides insights into model performance **metrics**, **loss graphs**, and **training progress**, making it easier to **diagnose** issues and **optimize** training strategies.

COMPONENTS

Python Programming

The **core** of the project is developed using **Python**, a versatile programming language widely used in machine learning and deep learning projects. Python facilitates **model** training, data handling, and **deploying** object detection models.

PyTorch Framework

The project utilizes **PyTorch**, an open-source deep learning framework that allows for **flexibility** and **speed** in model building and training. PyTorch is used for developing, training, and deploying deep learning models.

Convolutional Neural Networks (CNNs)

CNNs are the **backbone** of the project's object detection process. These networks automatically learn spatial hierarchies of features through **back propagation**, which is crucial for **identifying** objects in **images** and **videos**.

Some Practical Applications of CNN

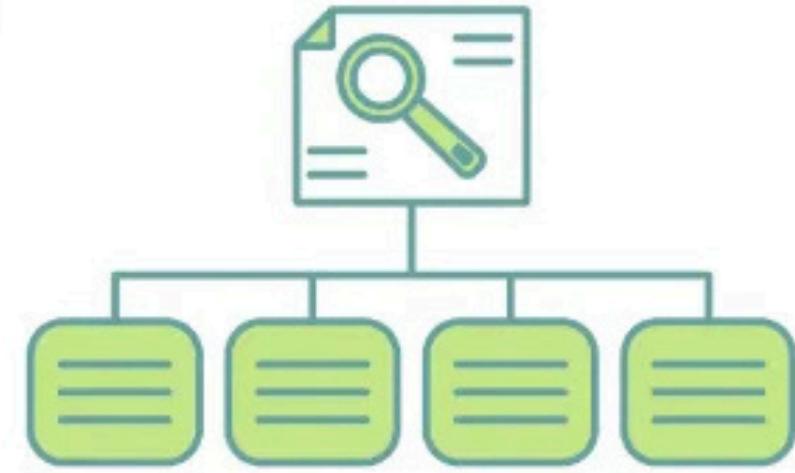
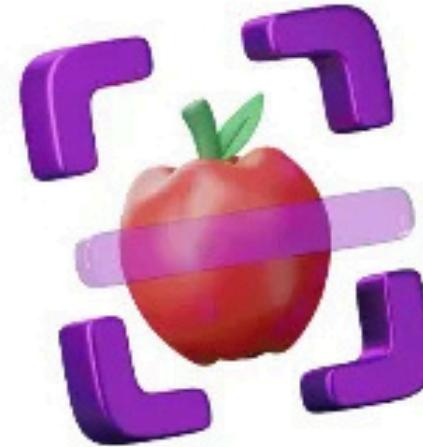
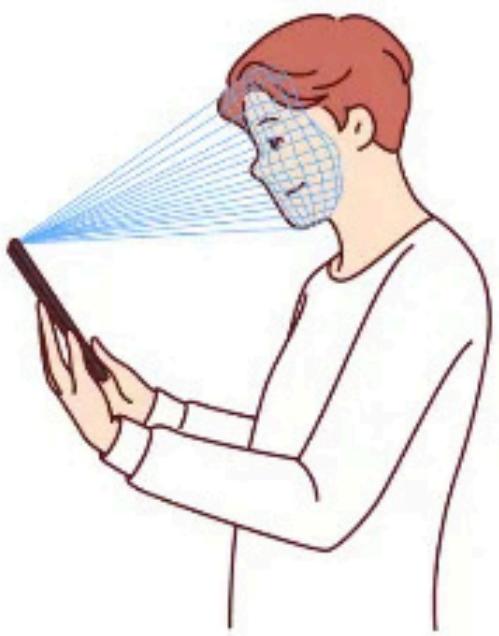


Image classification



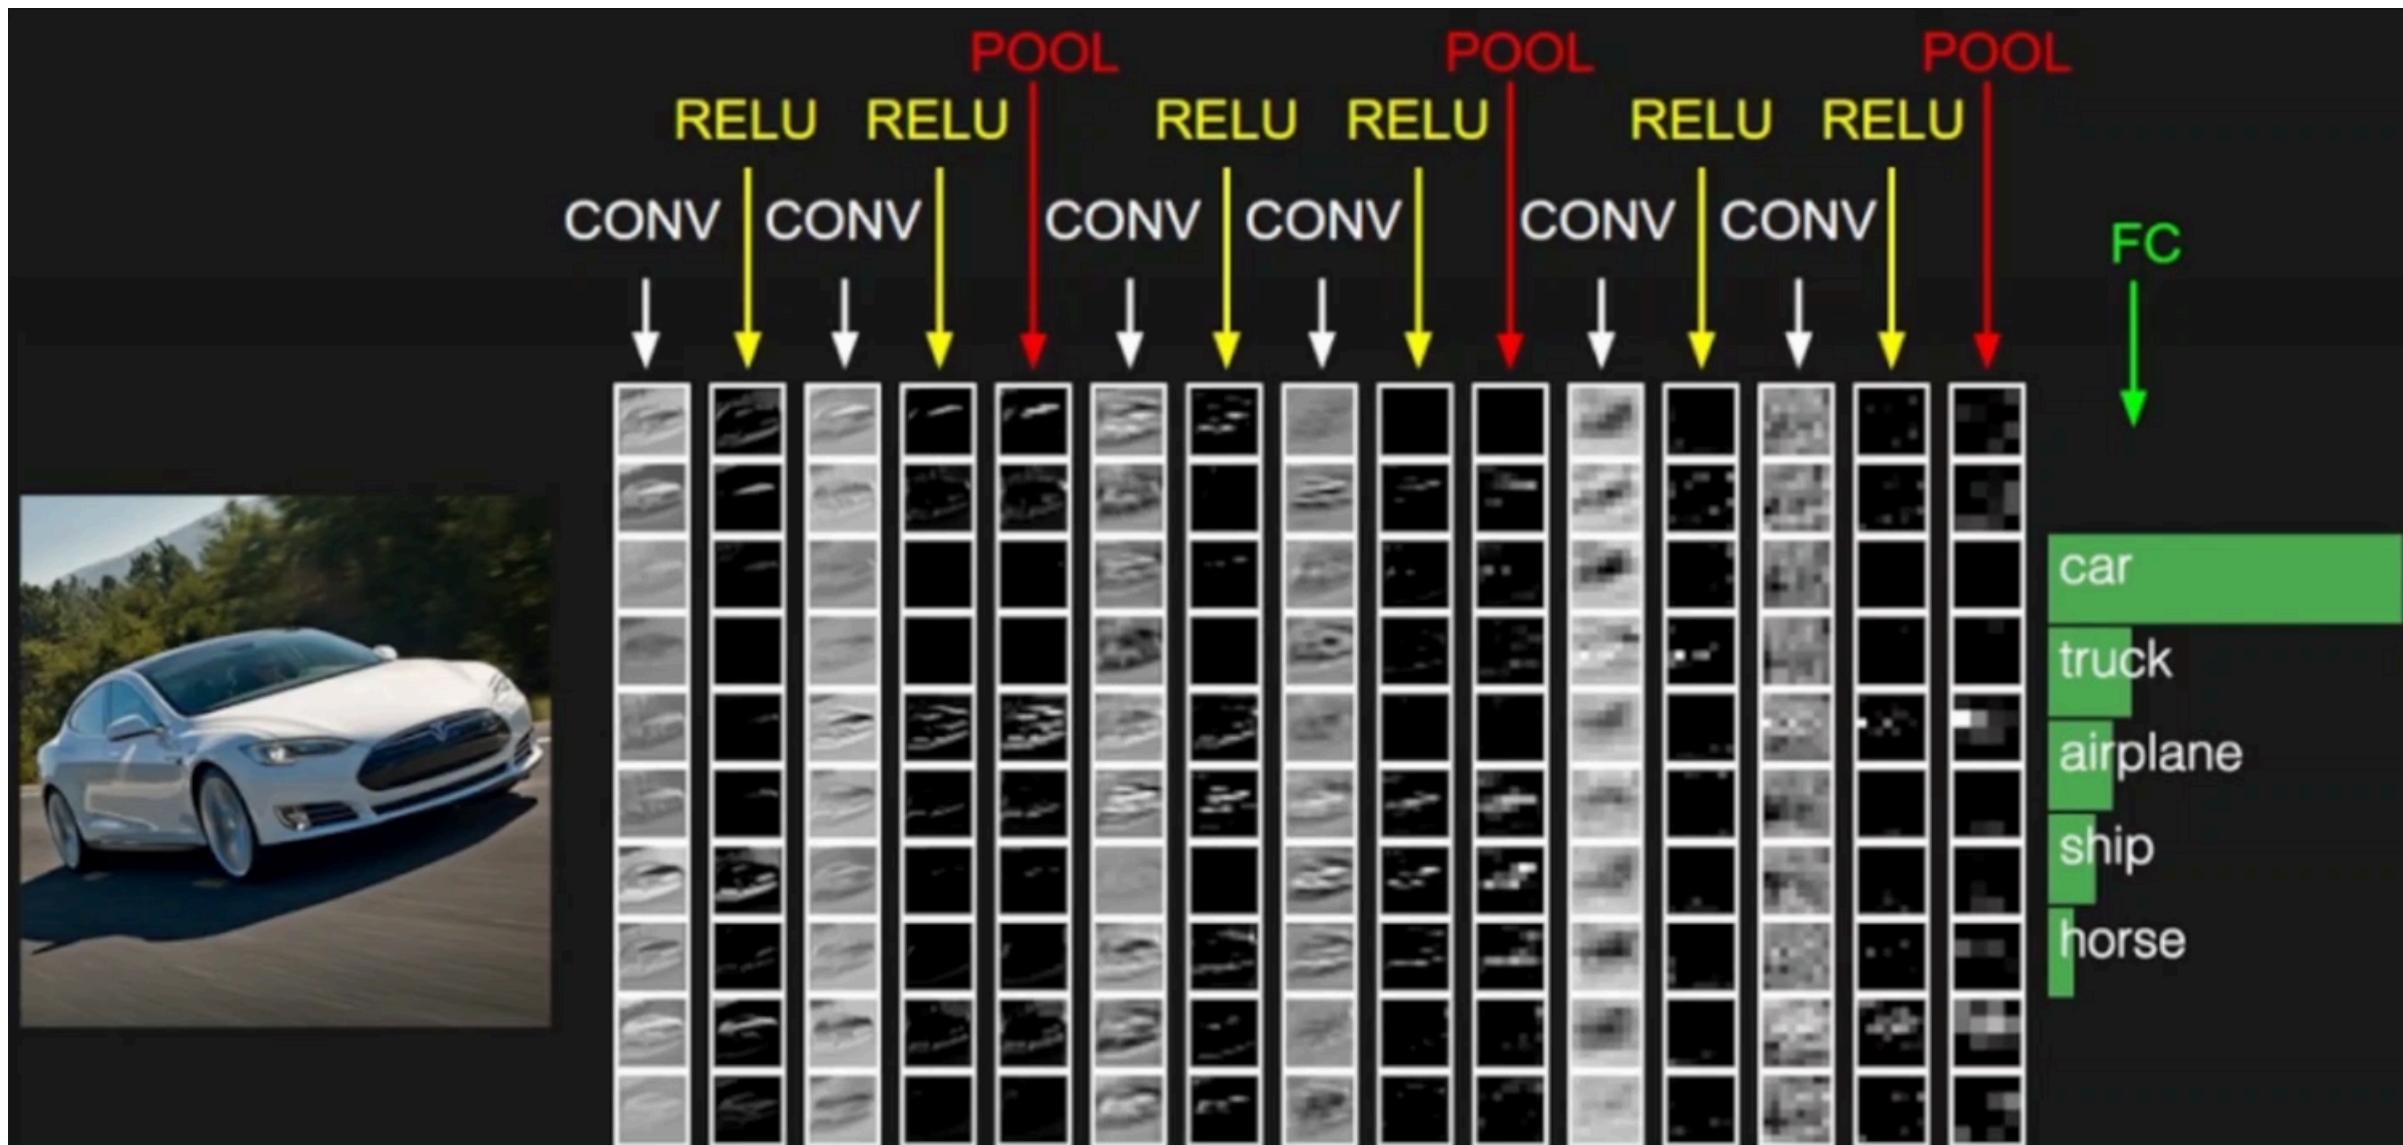
Object detection

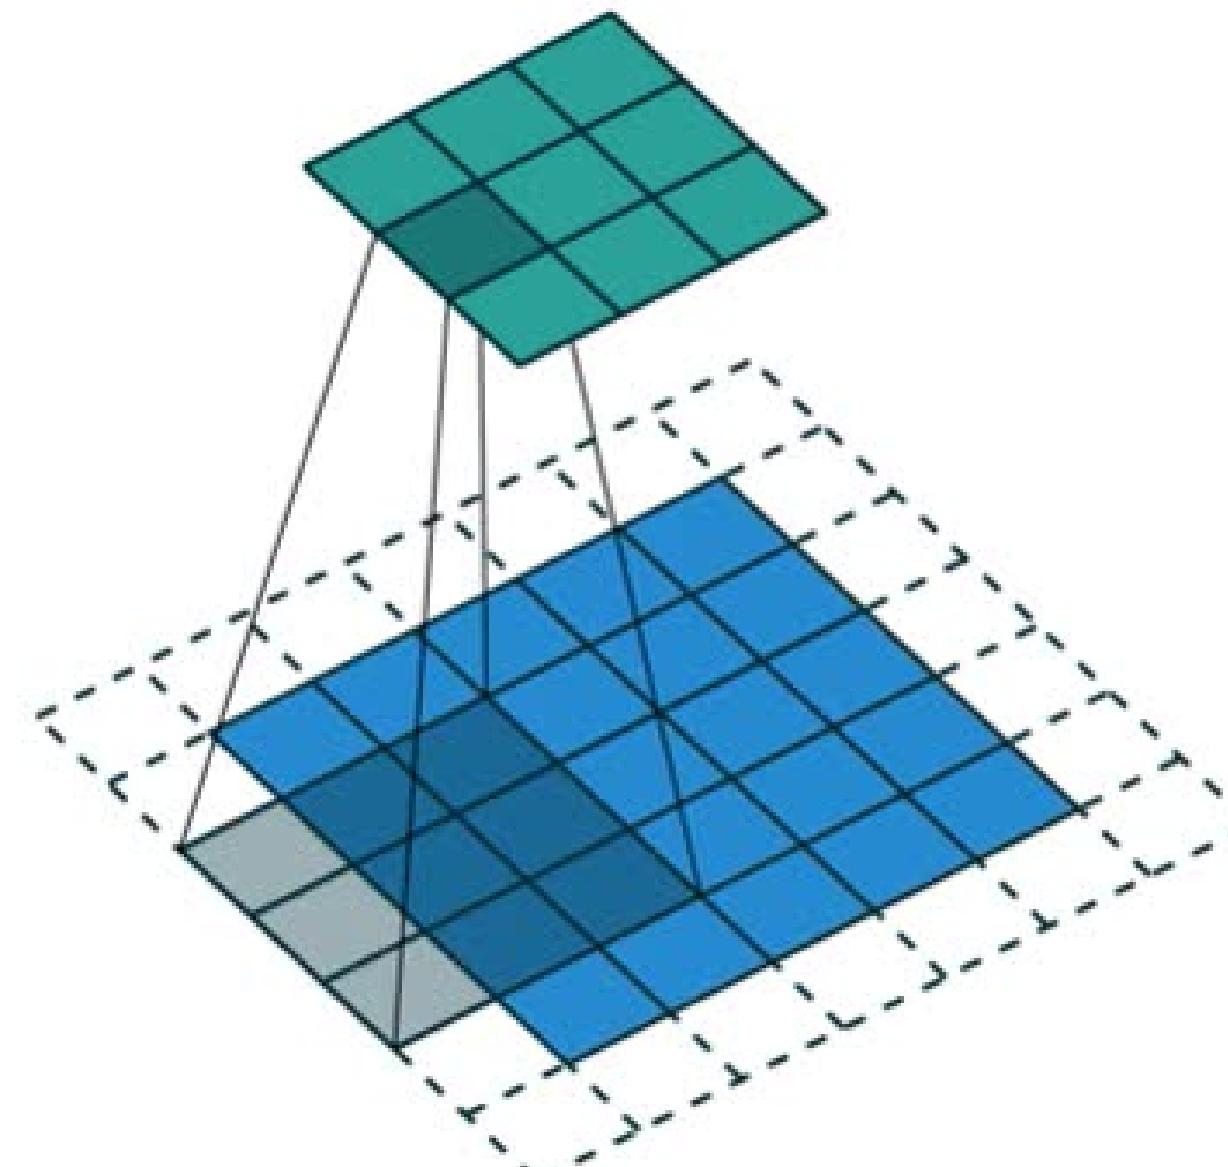


Facial recognition

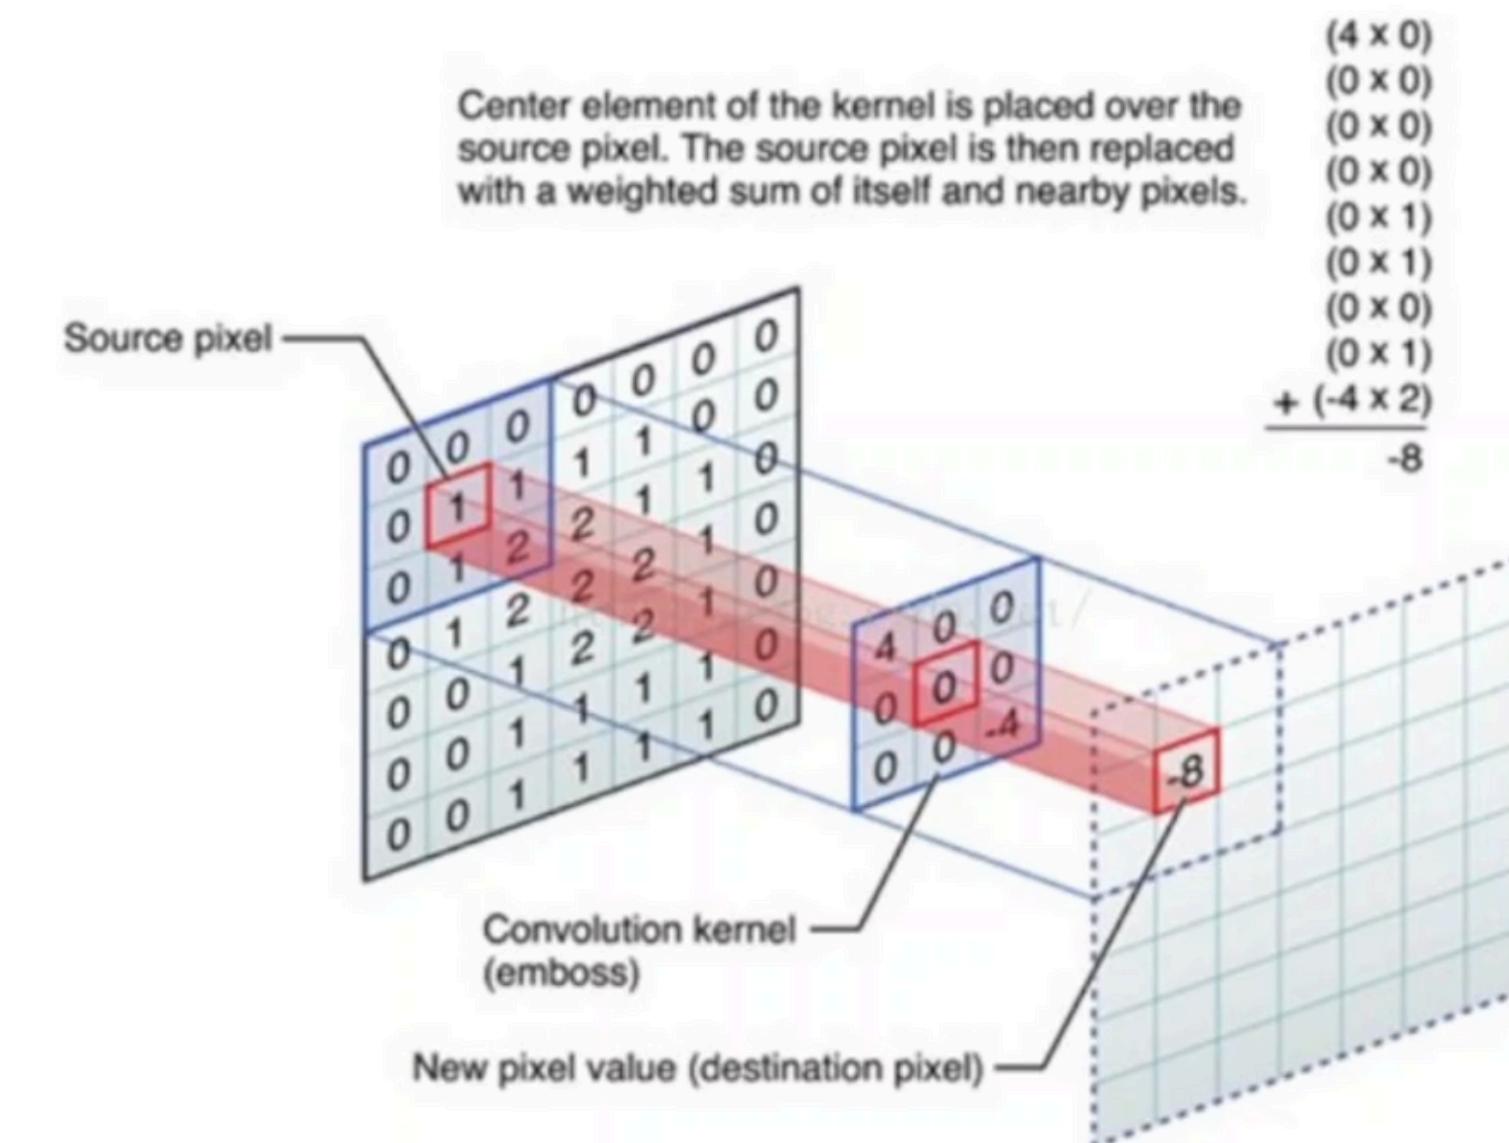
CNN ARCHITECTURE

- CONV: Convolutional kernel layer
- RELU: Activation function layer
- POOL: Pooling layer for dimension reduction
- FC: Fully connection layer





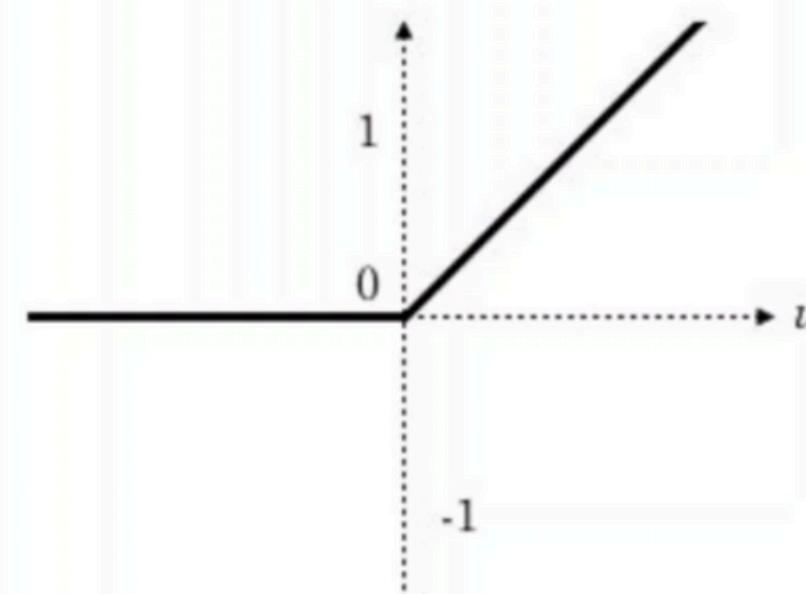
Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.



Activation Layer (ReLU)

rectified linear function, $f(x) = \max(0, x)$

$$f(u) = \max(0, u)$$



Pooling

max pooling

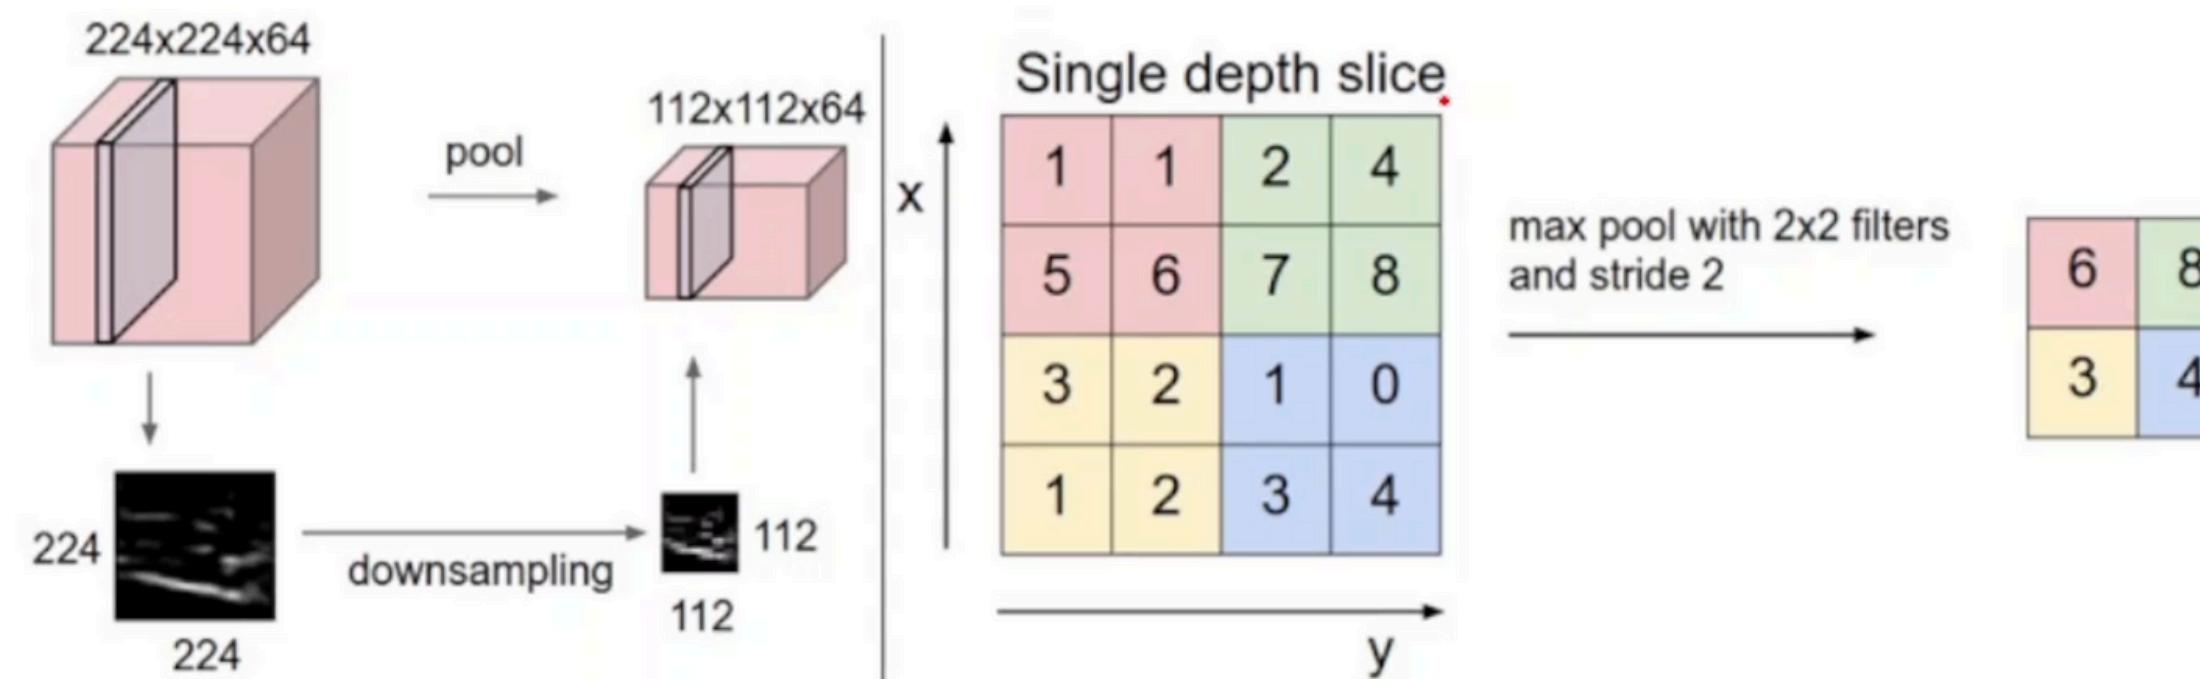
20	30
112	37

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

average pooling

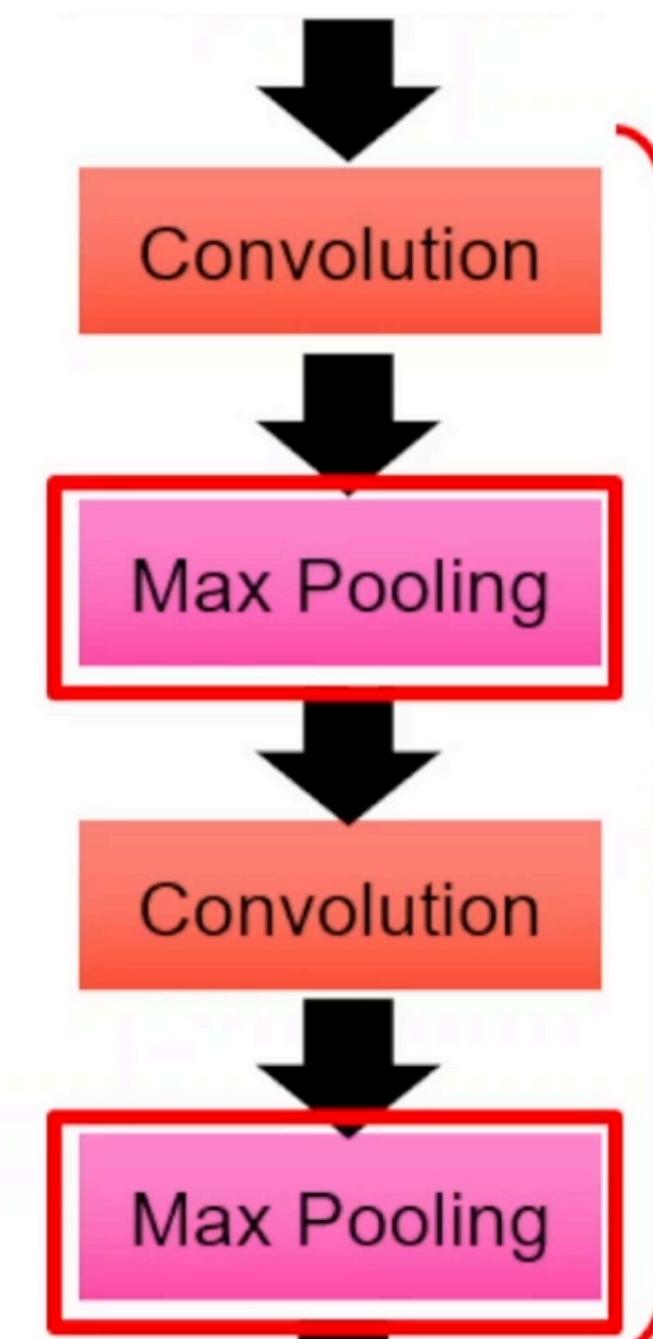
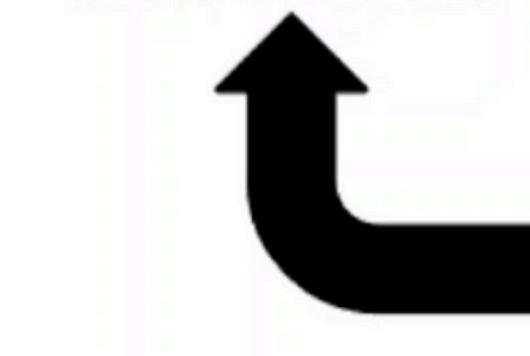
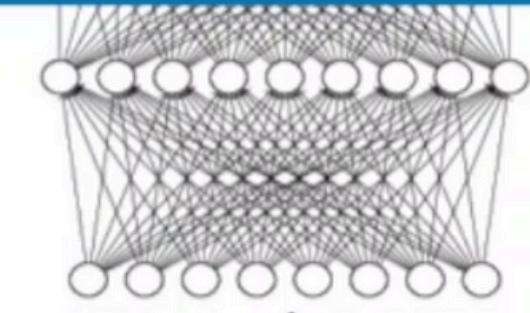
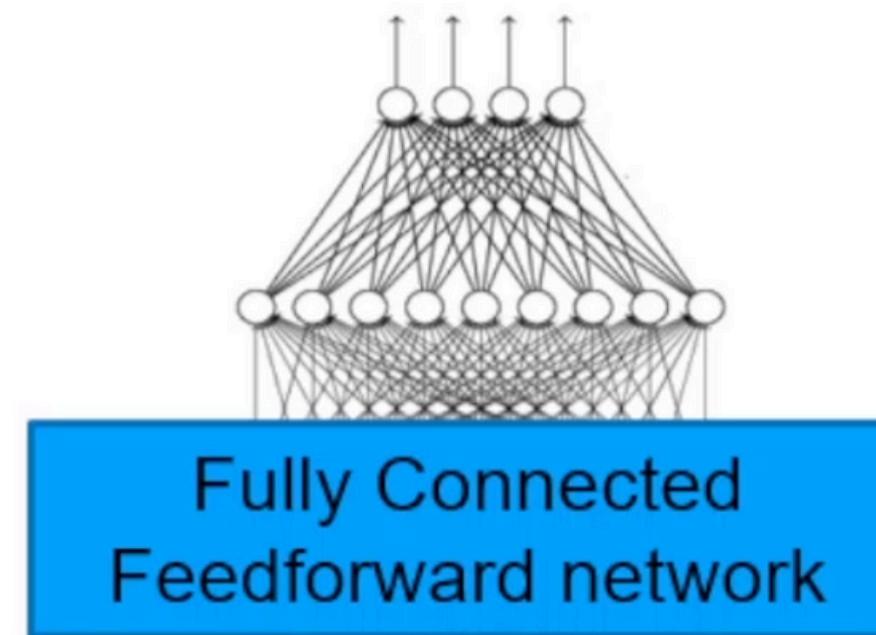
13	8
79	20

Pooling layer



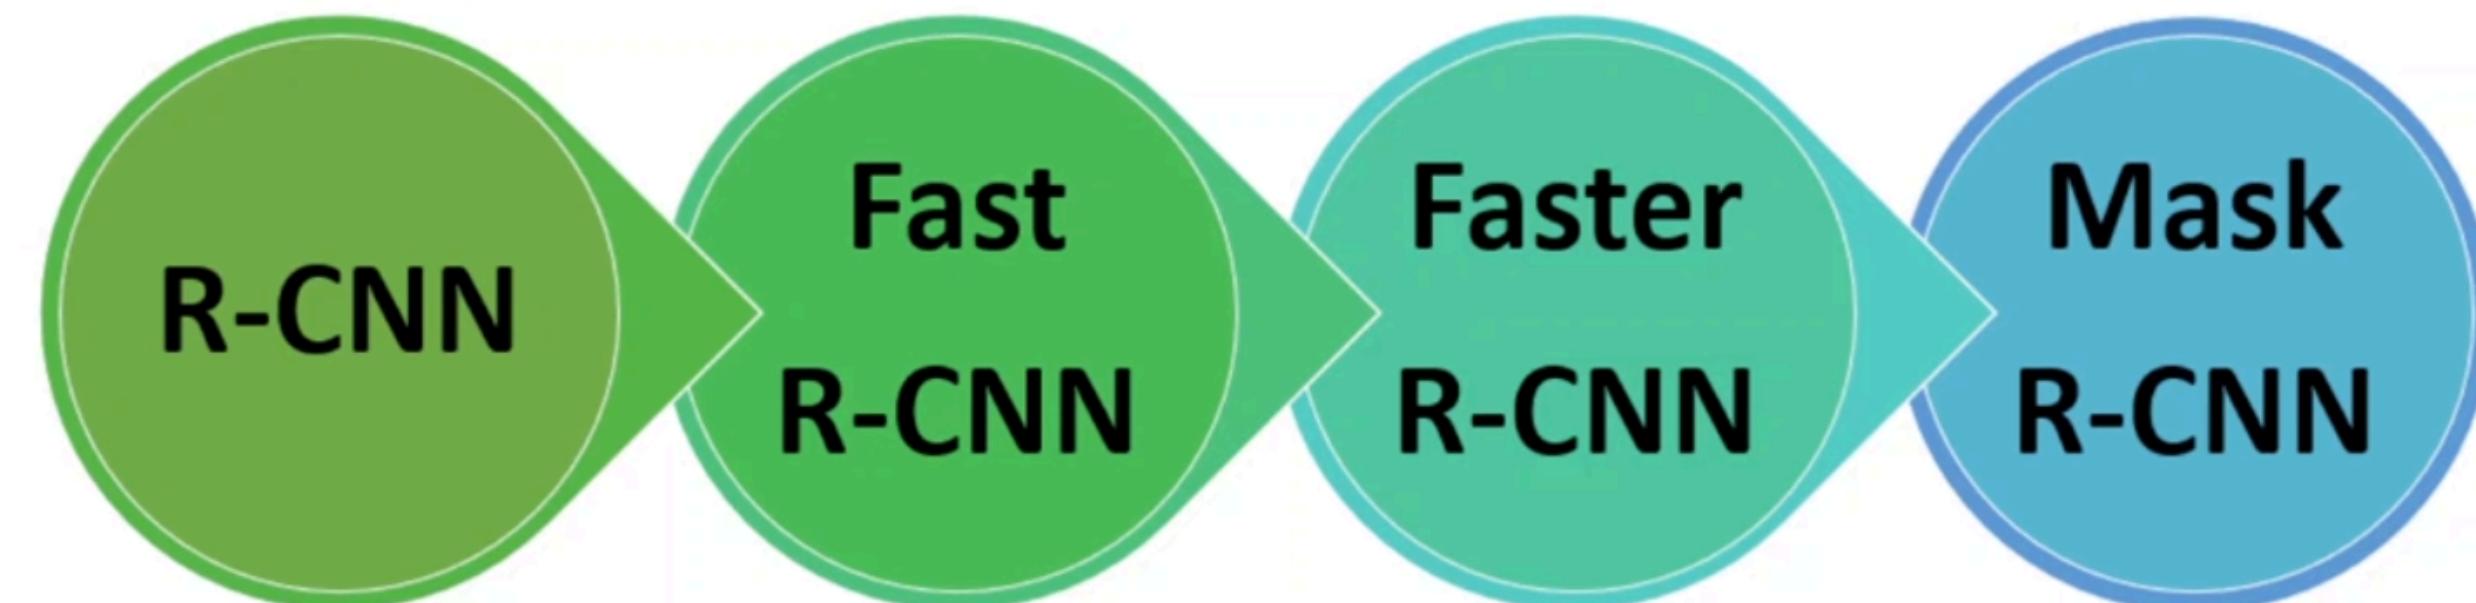
Pooling layer downsamples the volume spatially, independently in each depth slice of the input volume. **Left:** In this example, the input volume of size $[224 \times 224 \times 64]$ is pooled with filter size 2, stride 2 into output volume of size $[112 \times 112 \times 64]$. Notice that the volume depth is preserved. **Right:** The most common downsampling operation is max, giving rise to **max pooling**, here shown with a stride of 2. That is, each max is taken over 4 numbers (little 2×2 square).

cat dog



Flattened

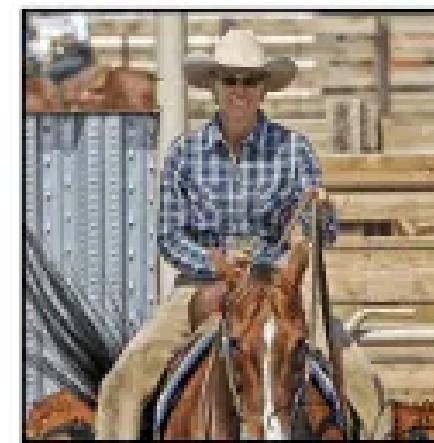
R-CNN FAMILY



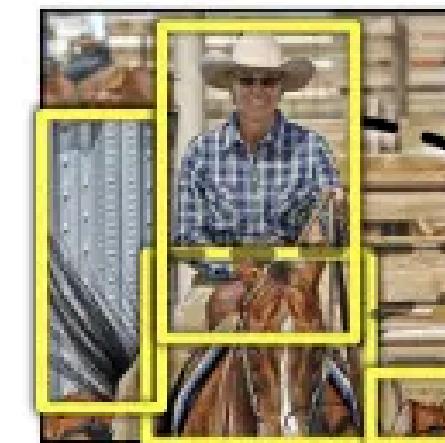
R-CNN

Uses Selective Search to propose regions, then classifies each region using a CNN.

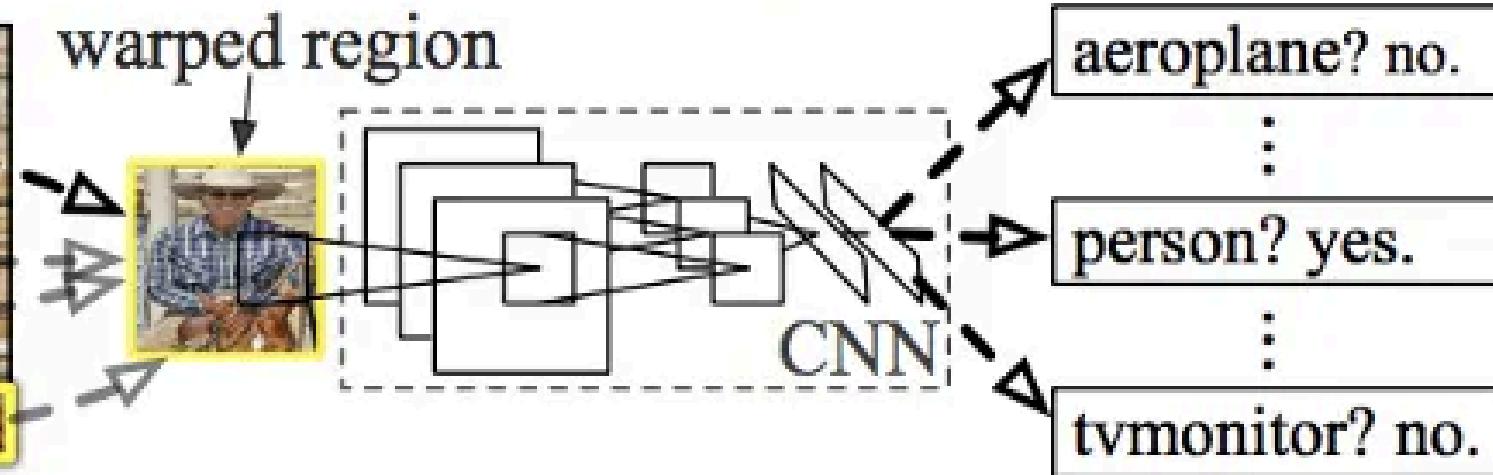
R-CNN: *Regions with CNN features*



1. Input
image



2. Extract region
proposals (~2k)



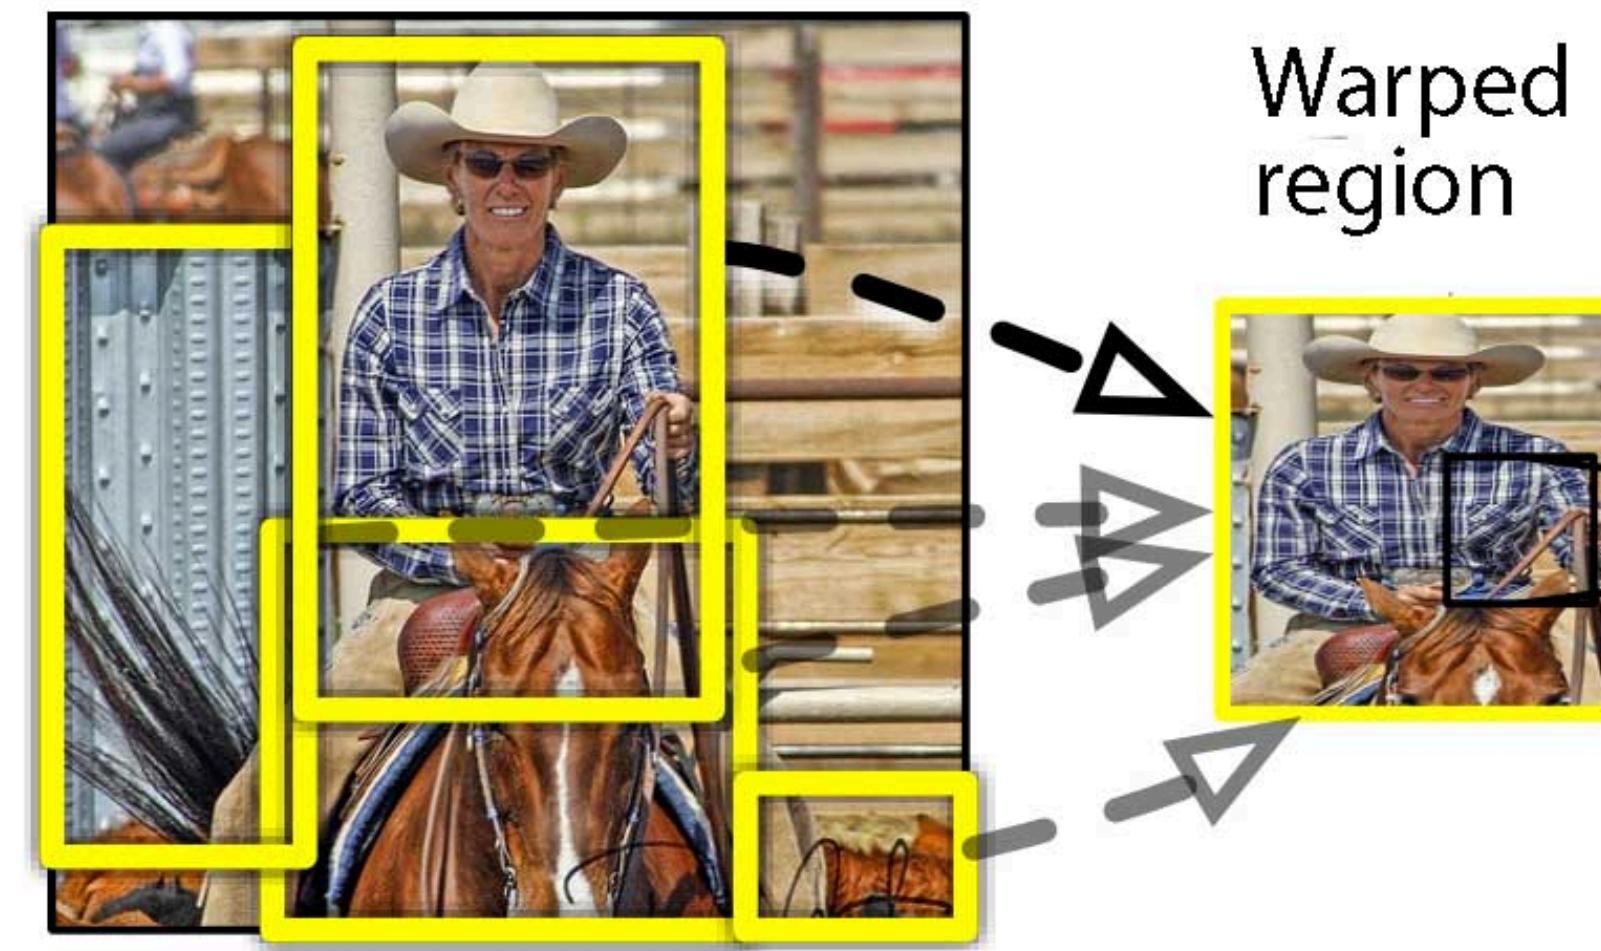
3. Compute
CNN features

4. Classify
regions

For an image with multiple objects, it proposes regions and classifies them as *aeroplane*, *person*, or *tvmonitor*.

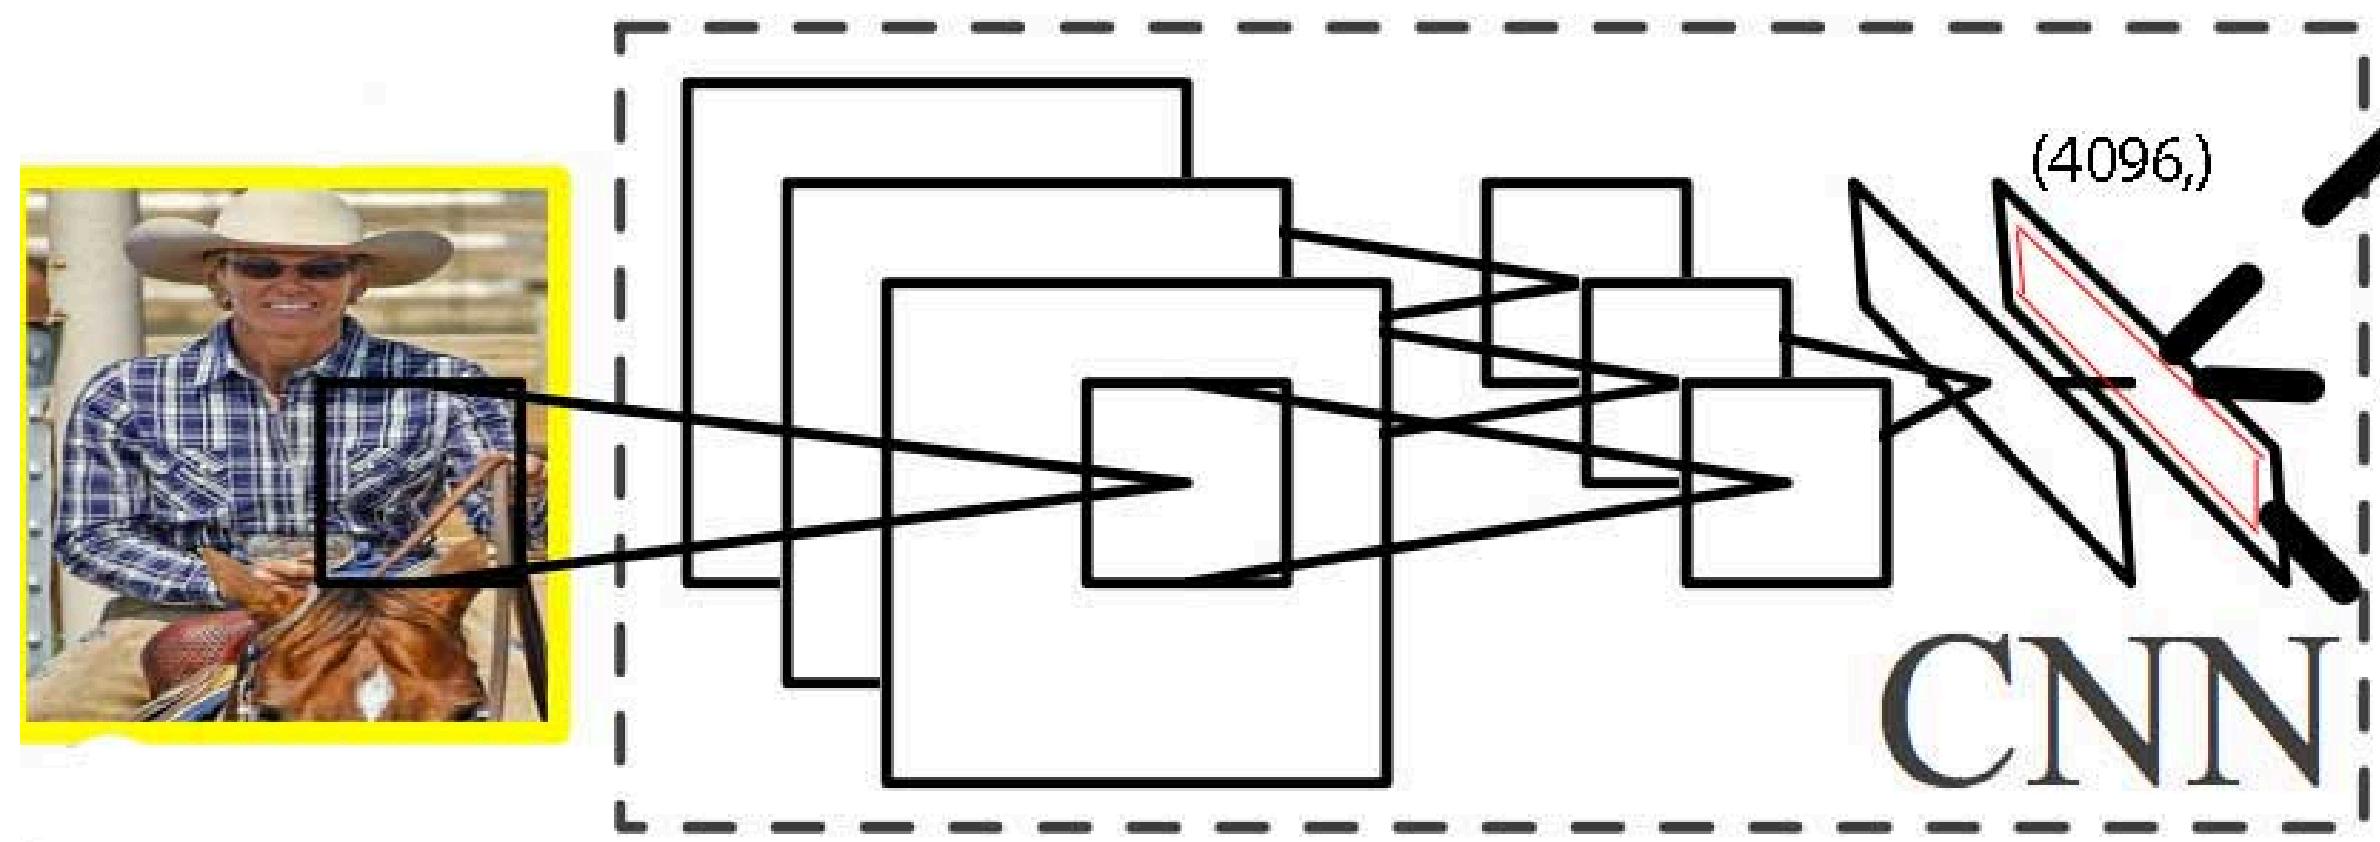
R-CNN – WARPED REGION

In R-CNN, after *selecting regions of interest (RoIs)*, each region is *resized to a fixed size of 227 x 227 x 3* before being processed by a *Convolutional Neural Network (CNN)*. This resizing ensures *consistent input dimensions* for accurate object extraction from the *identified areas*. This process is known as, *Warping*.



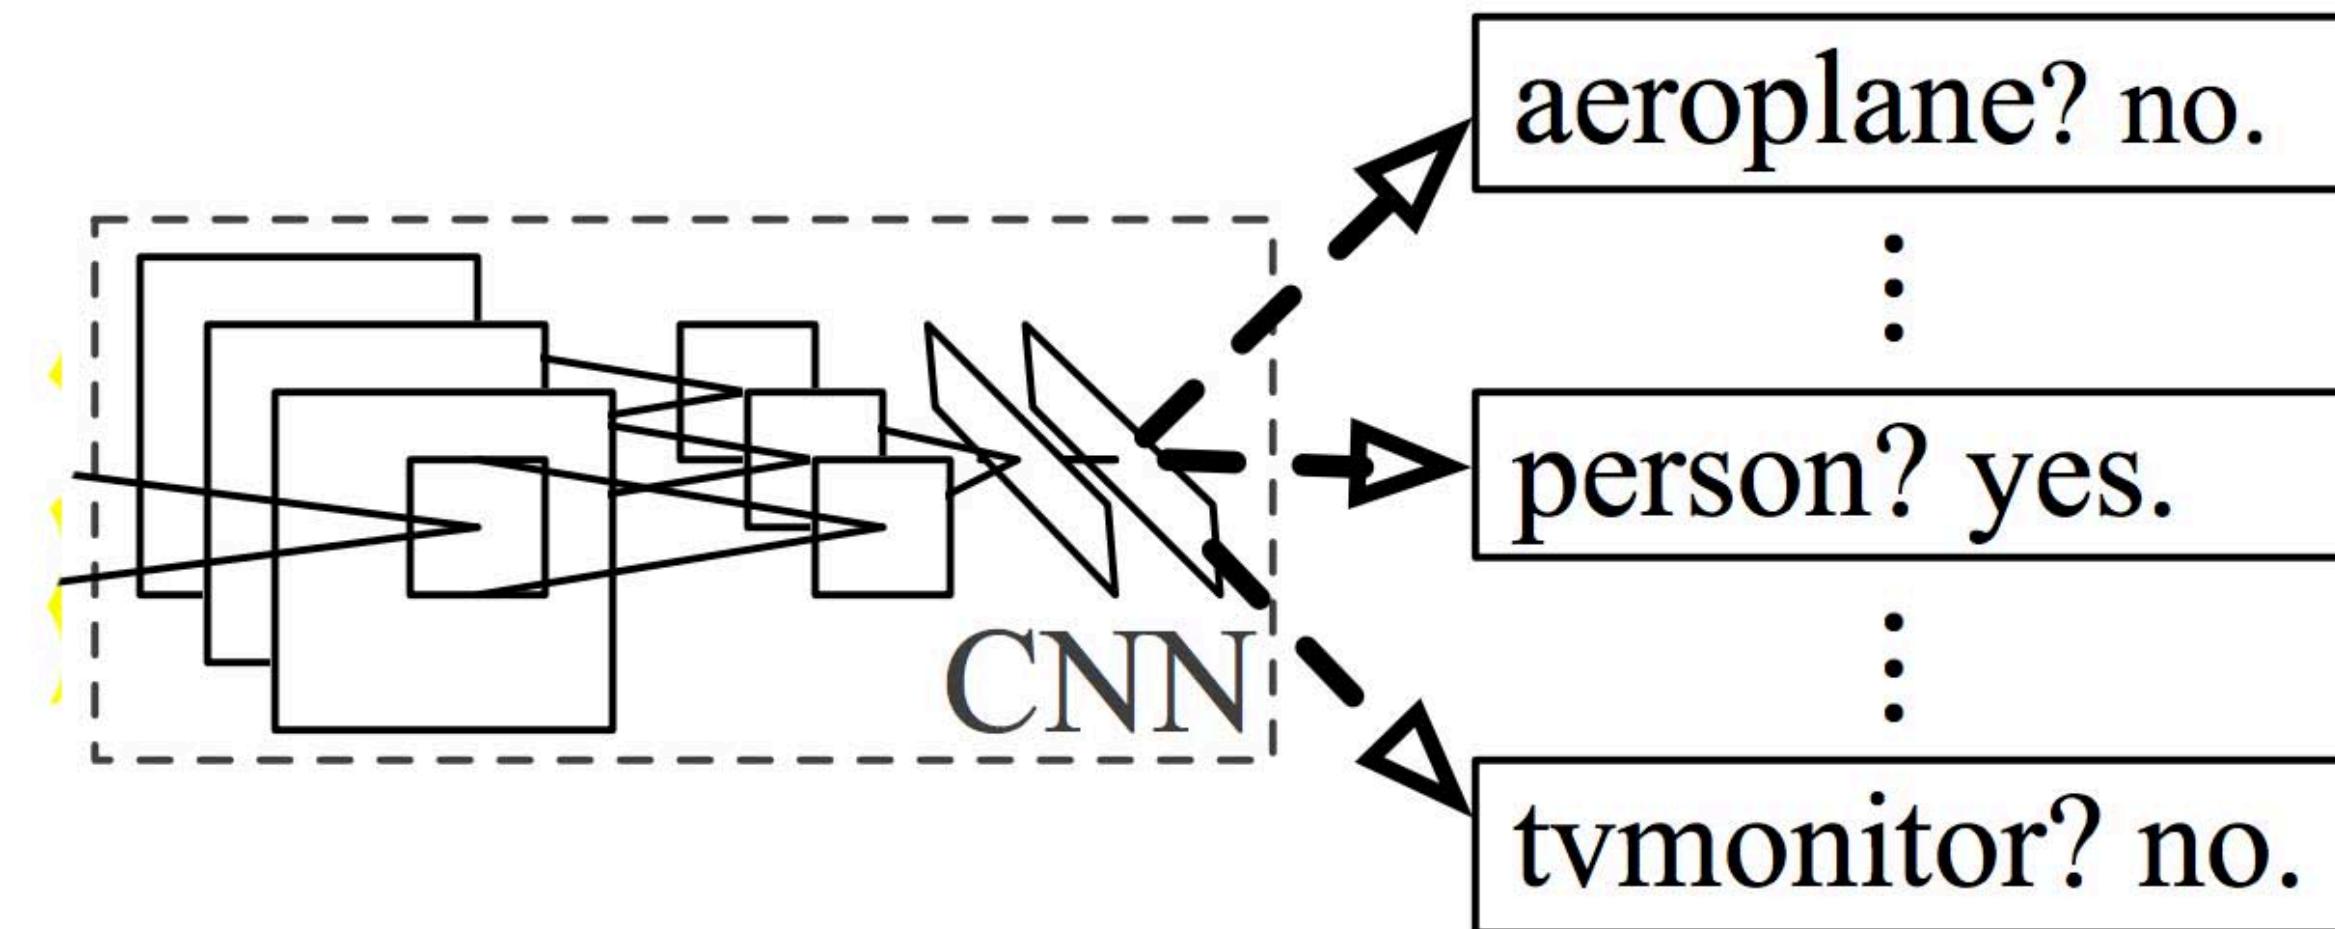
R-CNN – FEATURE EXTRACTION

A wrapped input for CNN will be processed to extract the object of size **4096** dimensions.

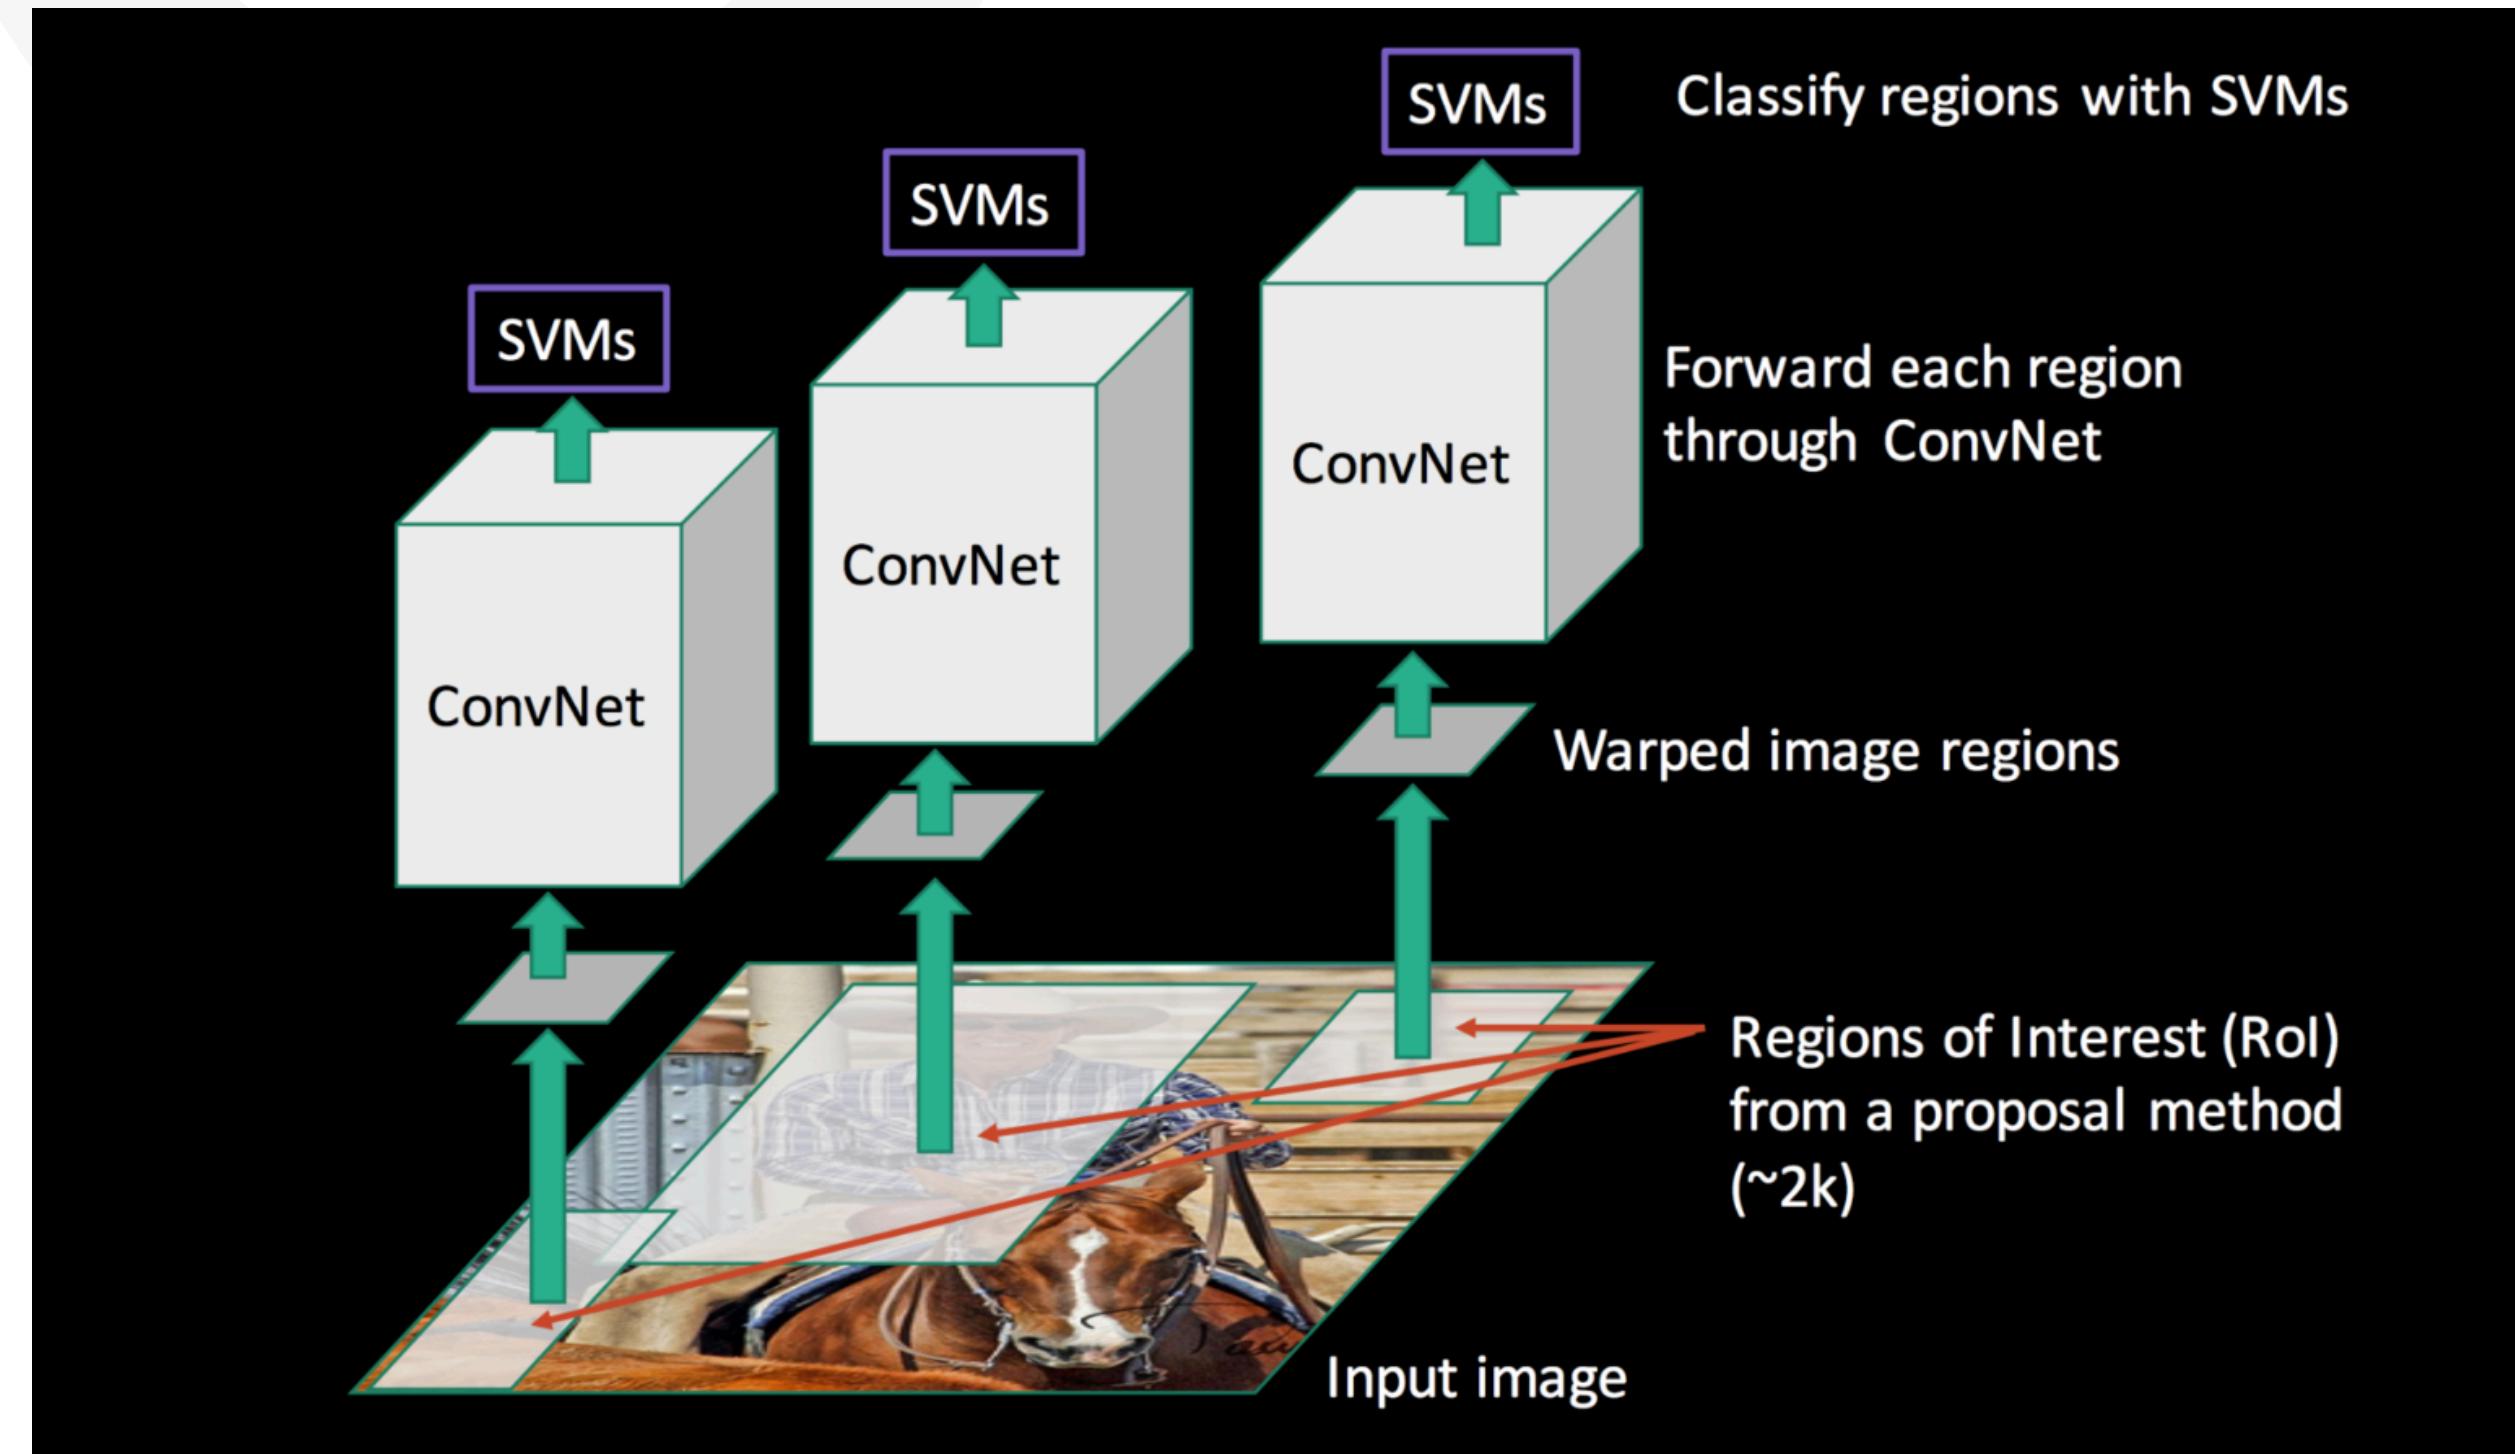


R-CNN – CLASSIFICATION

The basic R-CNN consists of an **SVM** classifier to **segregate** different objects into their class.



R-CNN – ARCHITECTURE



FAST R-CNN

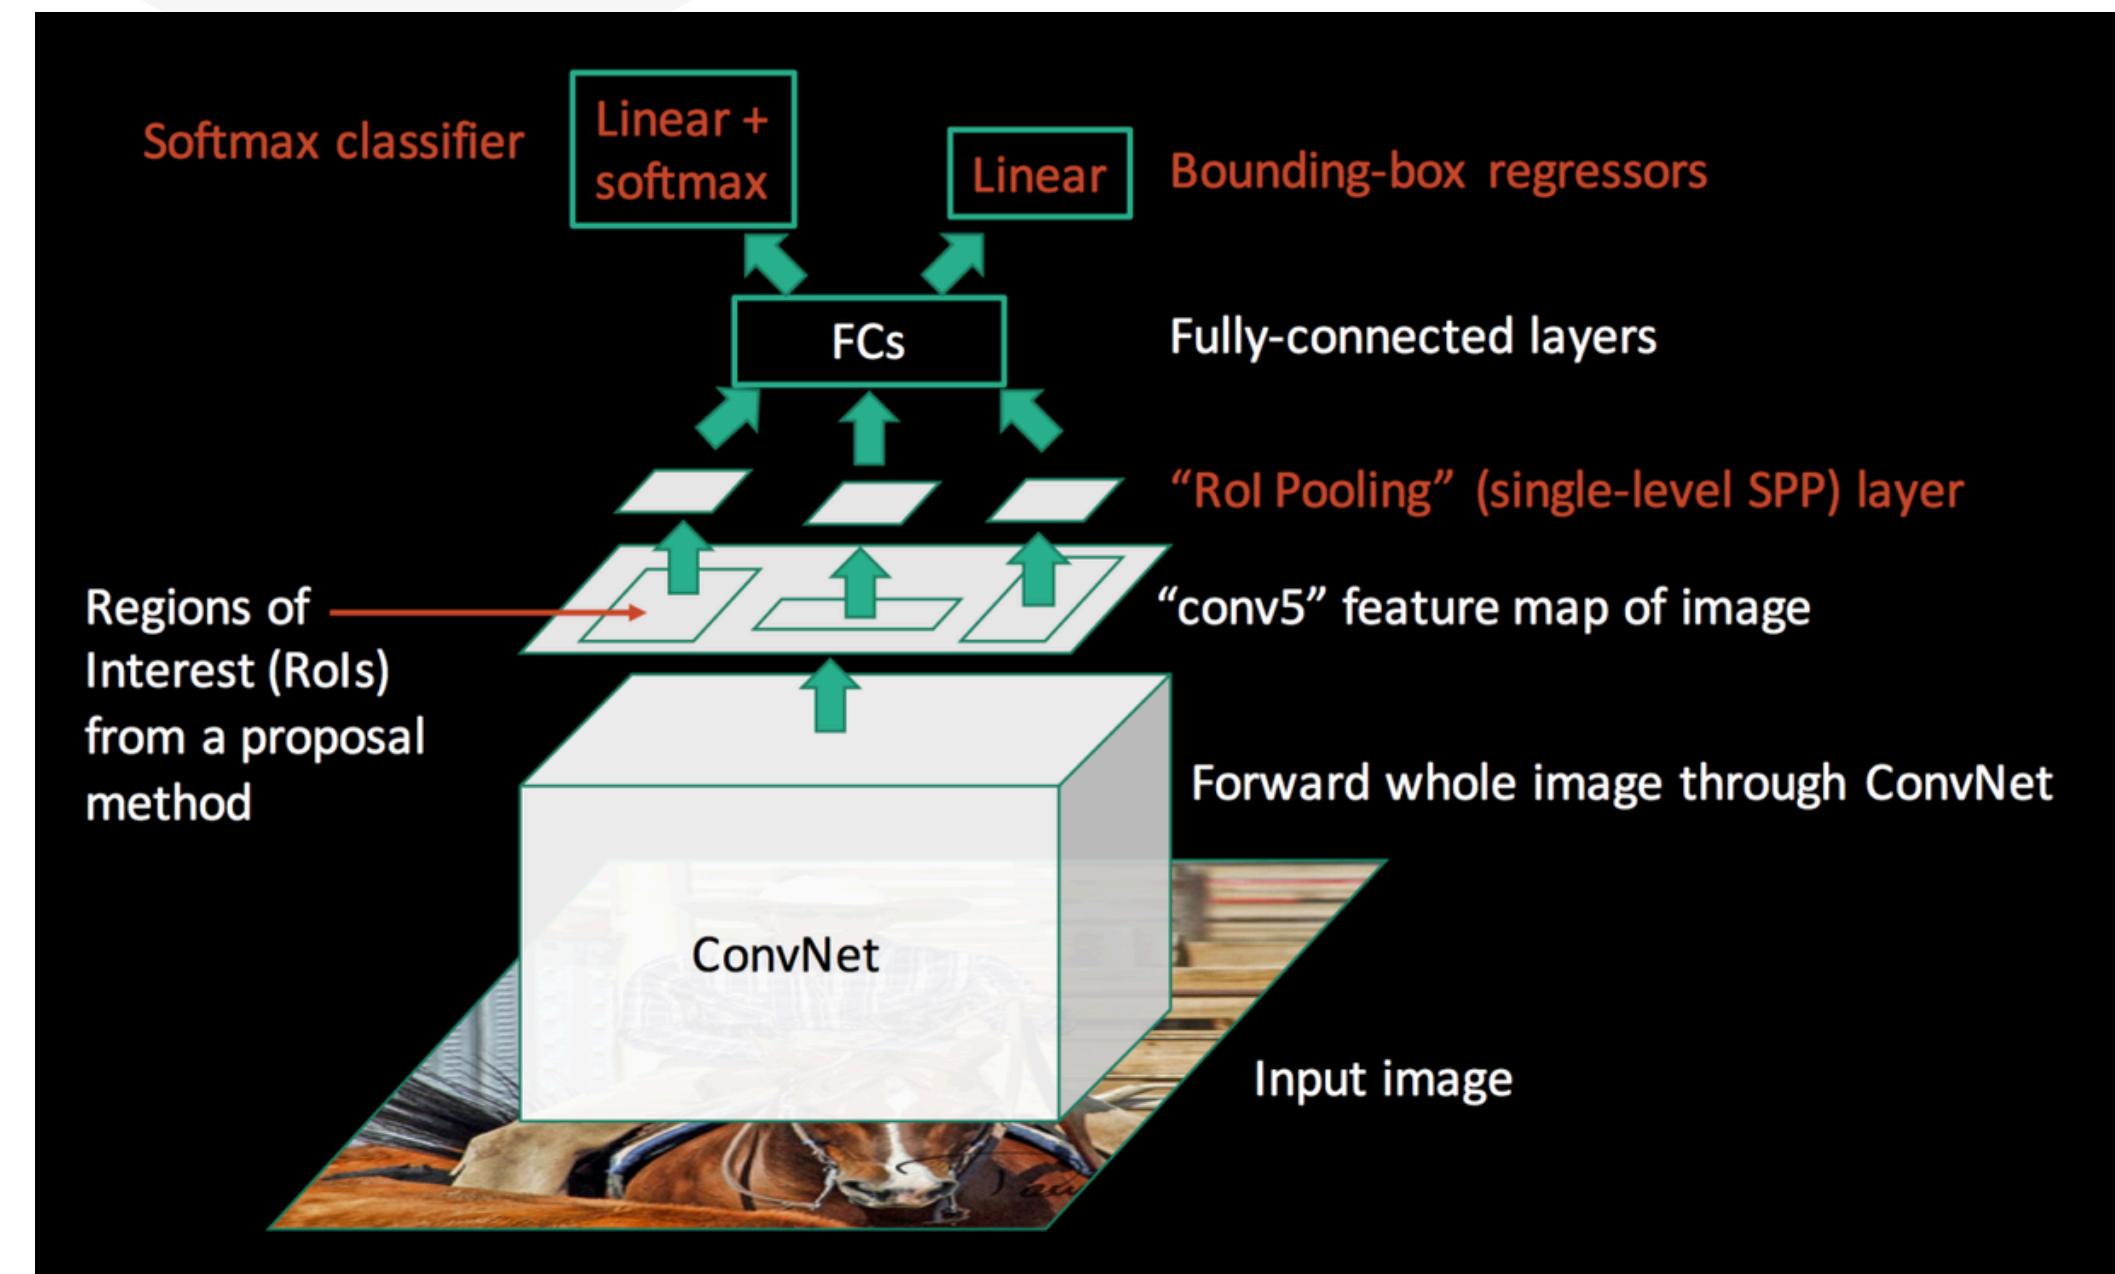
In Fast R-CNN instead of performing maximum pooling, we perform ROI pooling for utilising a Single Feature Map for all the regions. This warps ROIs into one single layer; the ROI pooling layer uses max pooling to convert the features.

We can clearly see the effects of Maximum Pooling layer black colour is more highlighted than other light colours which also helps in finding the dark coloured objects in the image.



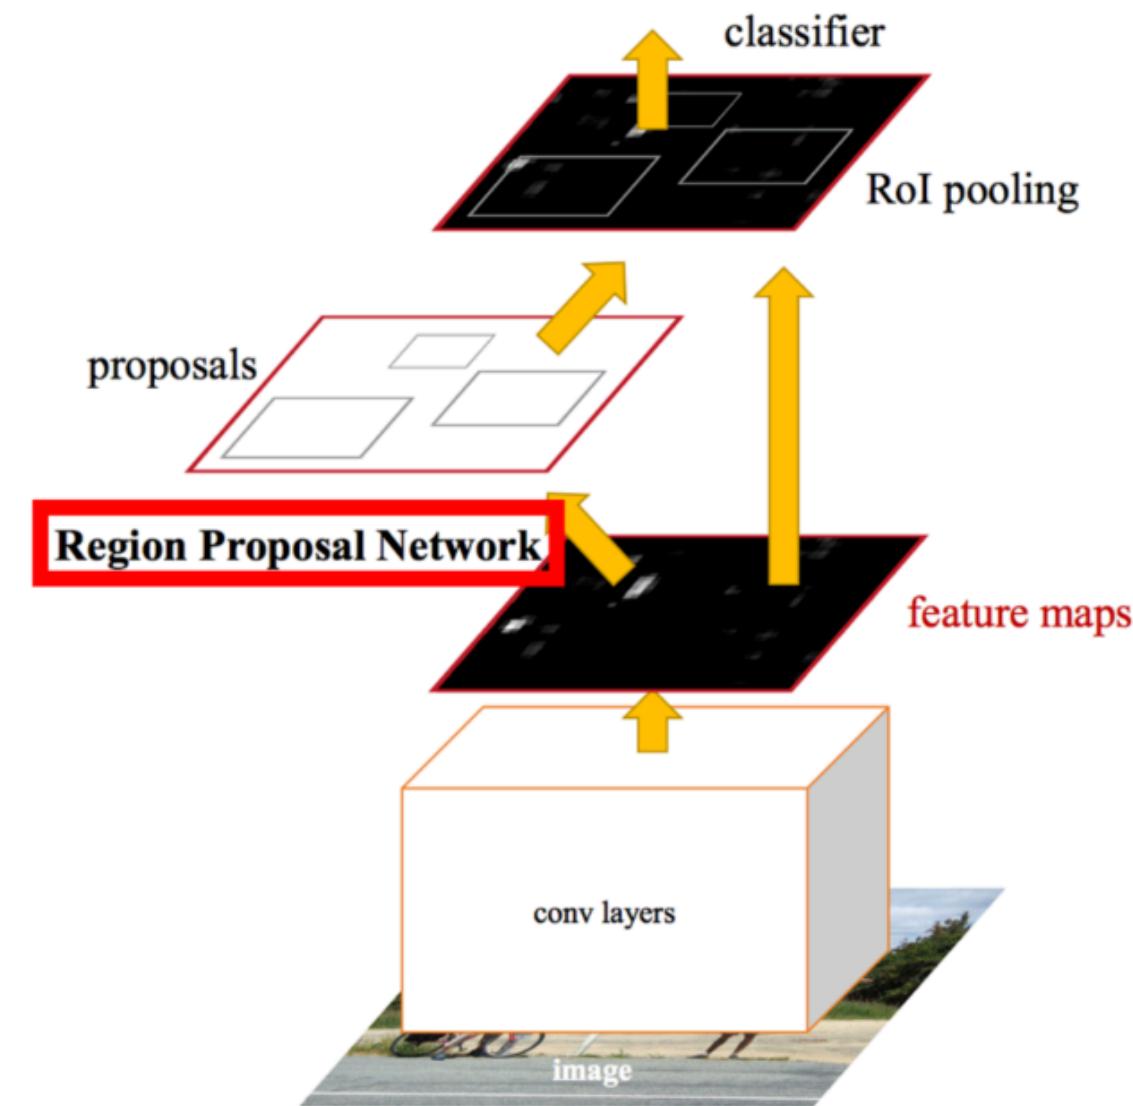
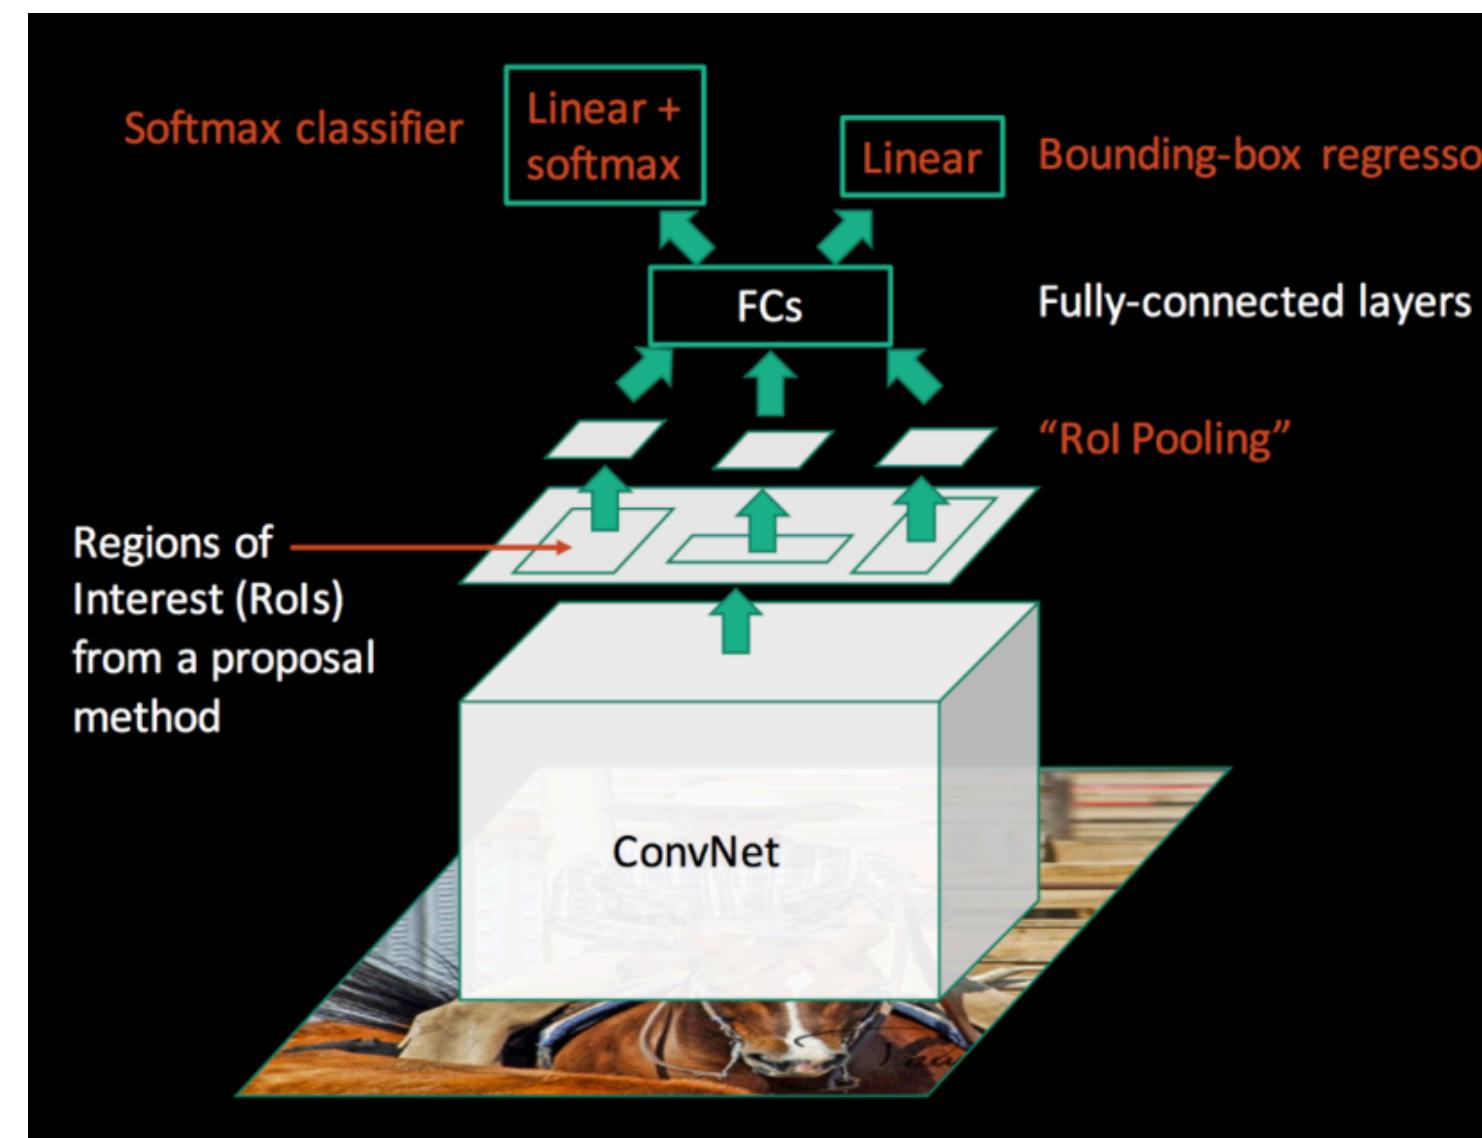
FAST R-CNN

Since Max Pooling is also working here, that's why we can consider Fast R-CNN as an upgrade of the SPPNet. Instead of generating layers in a pyramid shape, it generates only one layer.

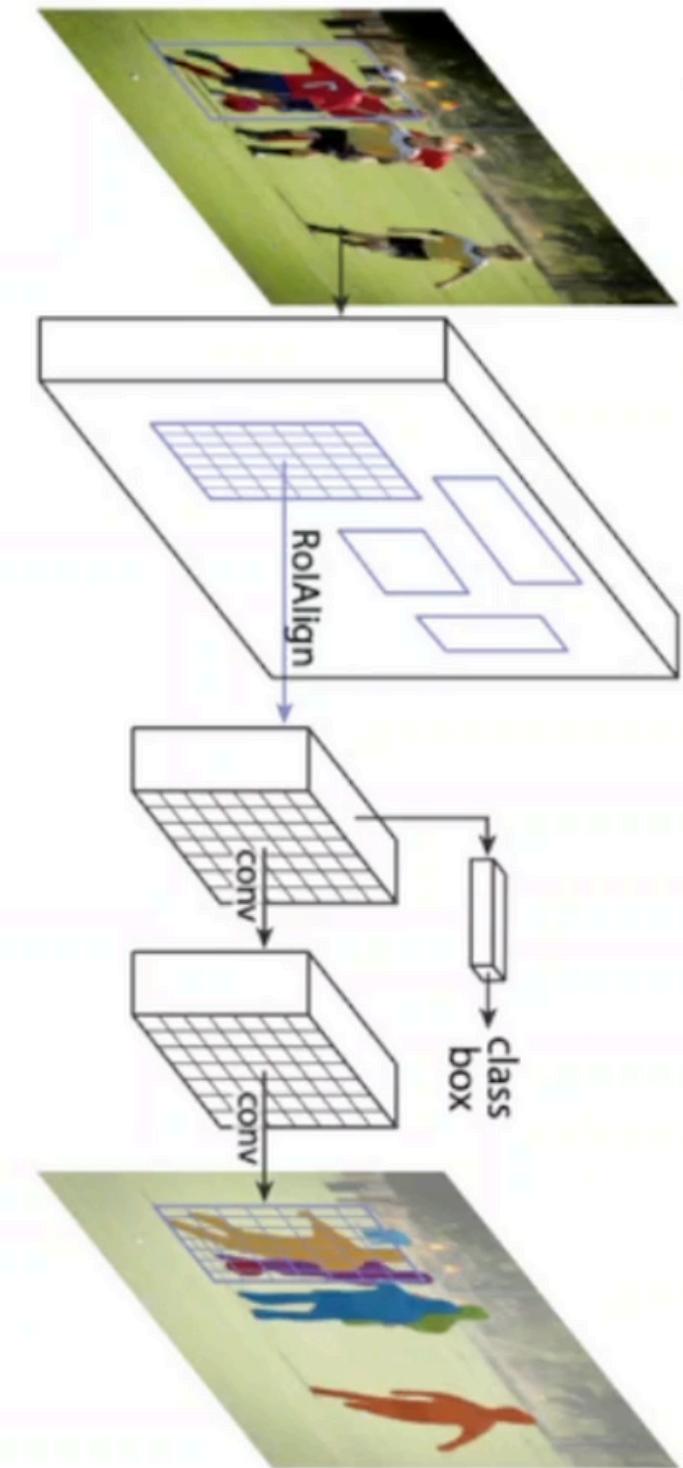
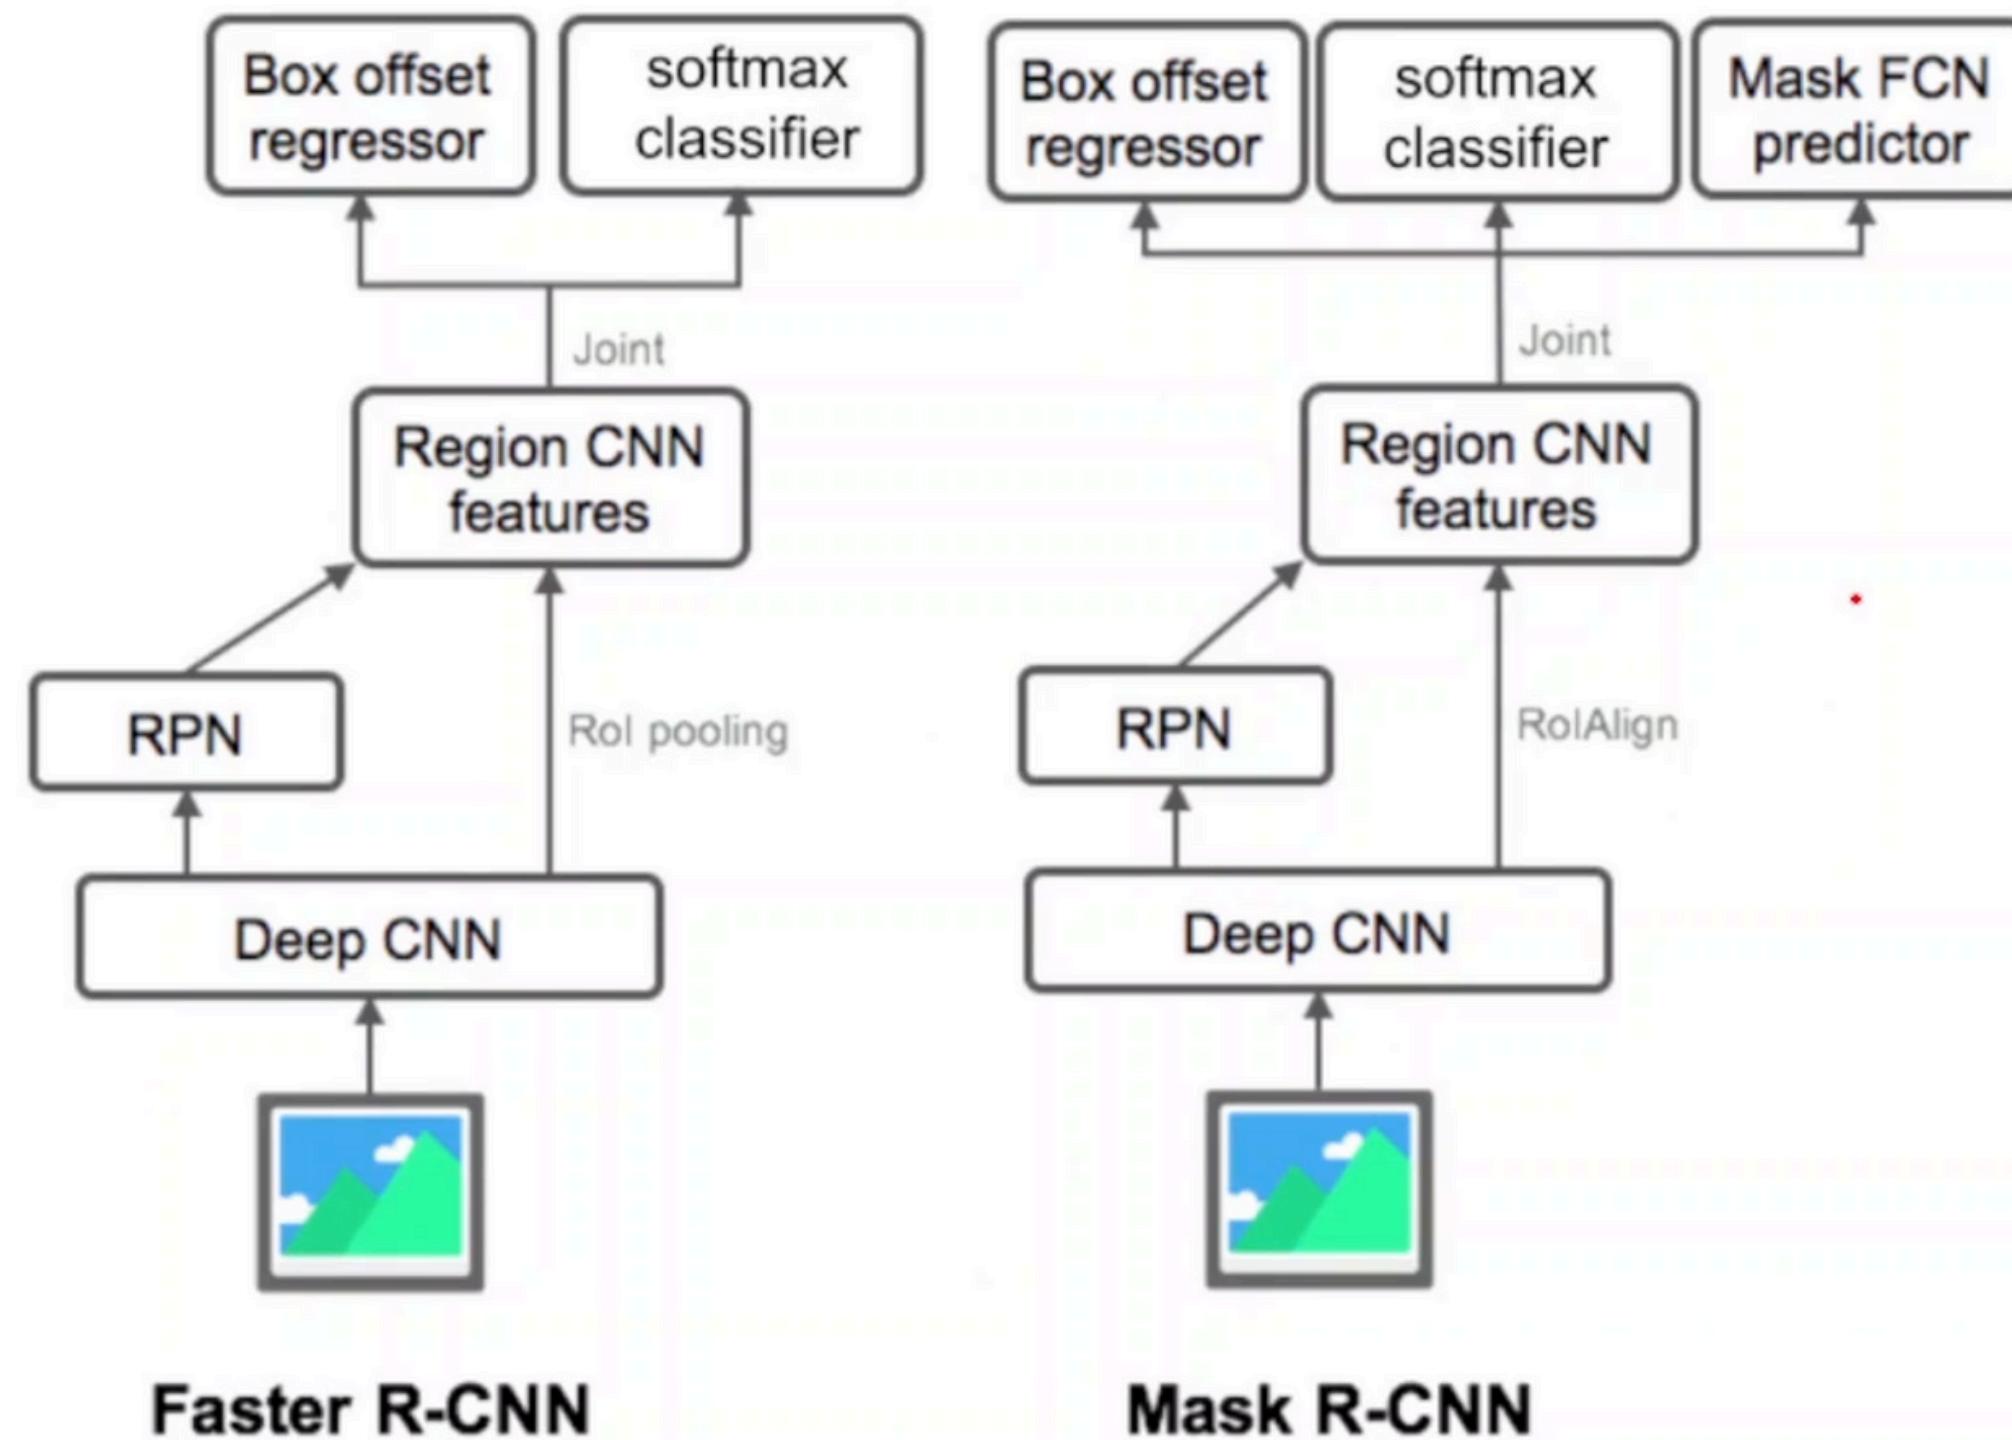


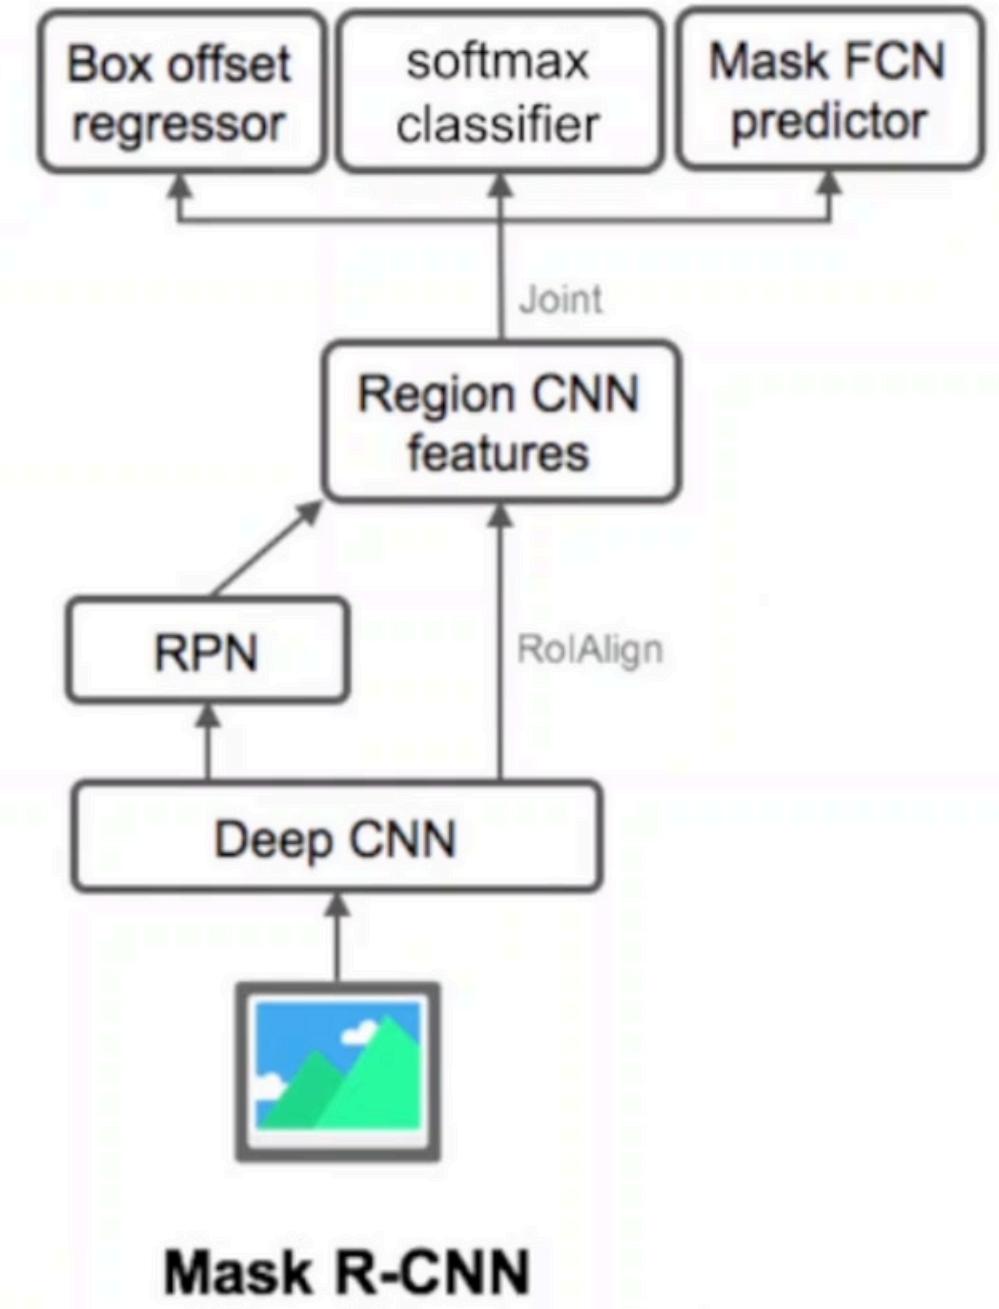
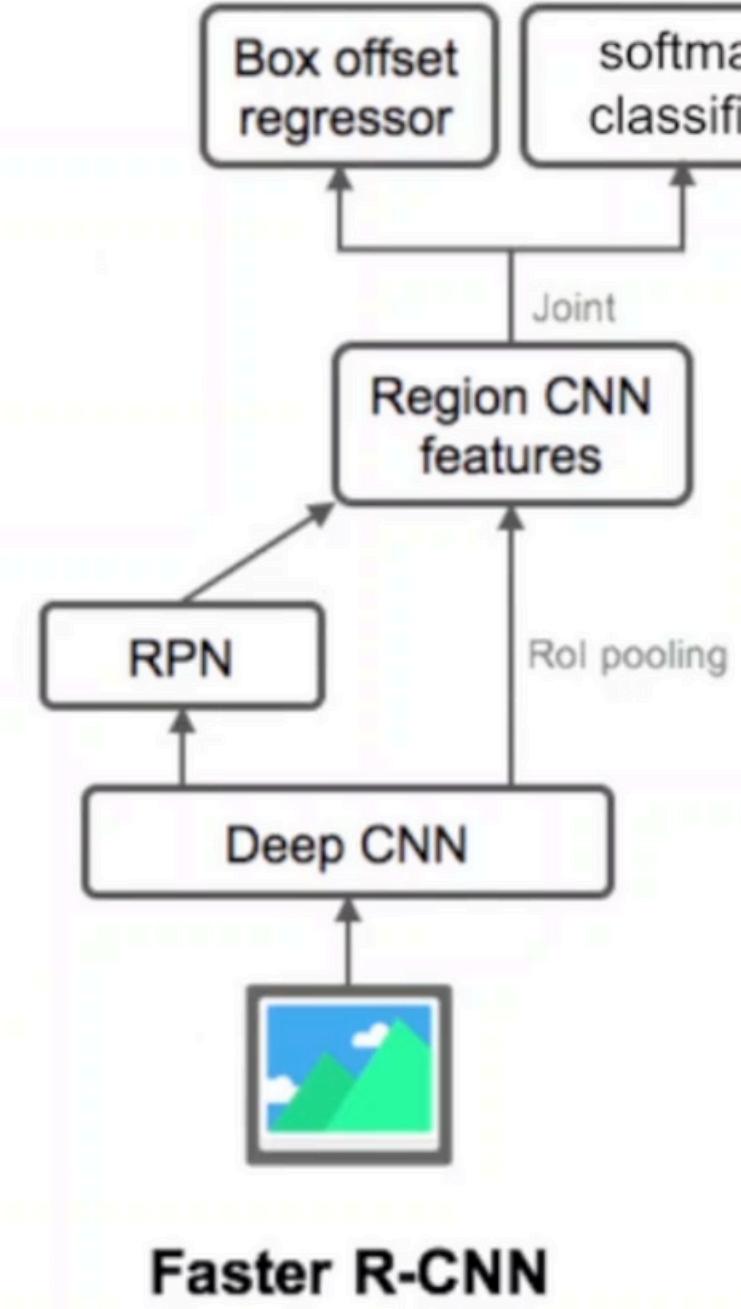
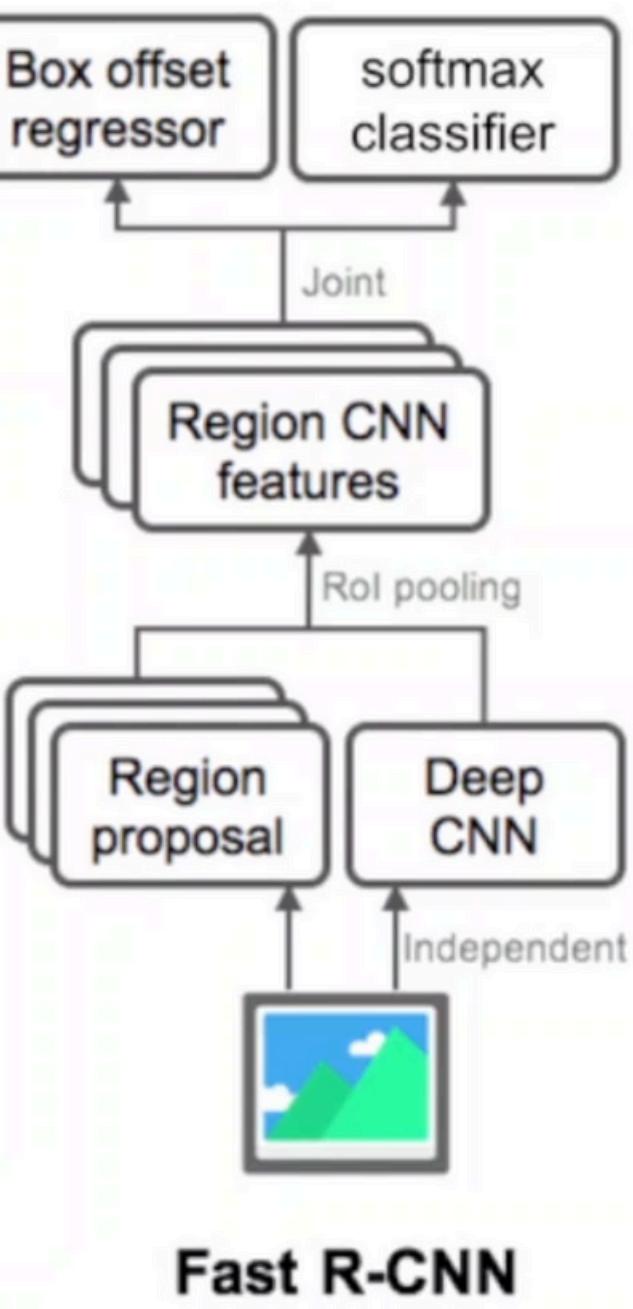
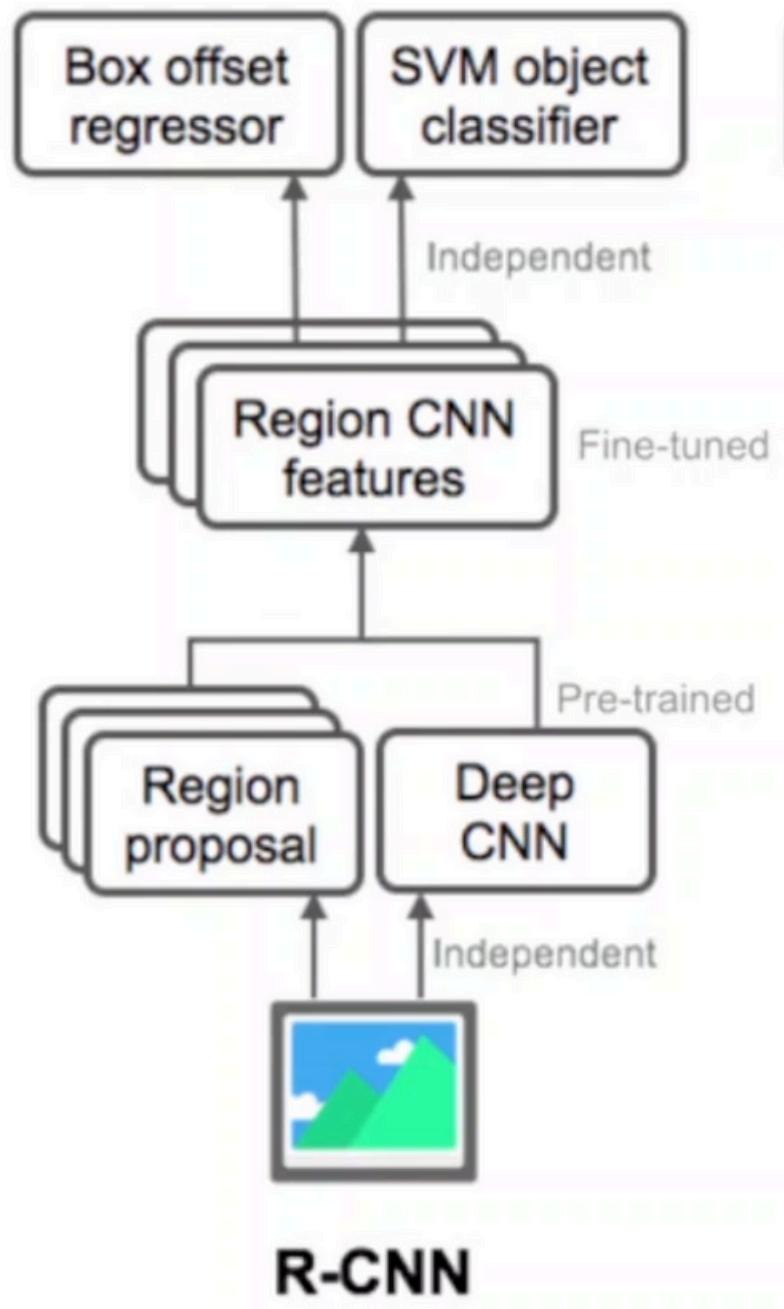
FASTER R-CNN

So far, we've noted that SPPNet and Fast R-CNN lack methods for selecting regions of interest, which distinguishes Fast R-CNN from Faster R-CNN. Faster R-CNN introduces a Region Proposal Network (RPN) that generates region proposals. The RPN takes the feature map as input during training and outputs these proposals, which are then sent to the ROI pooling layer for further processing.



MASK R-CNN



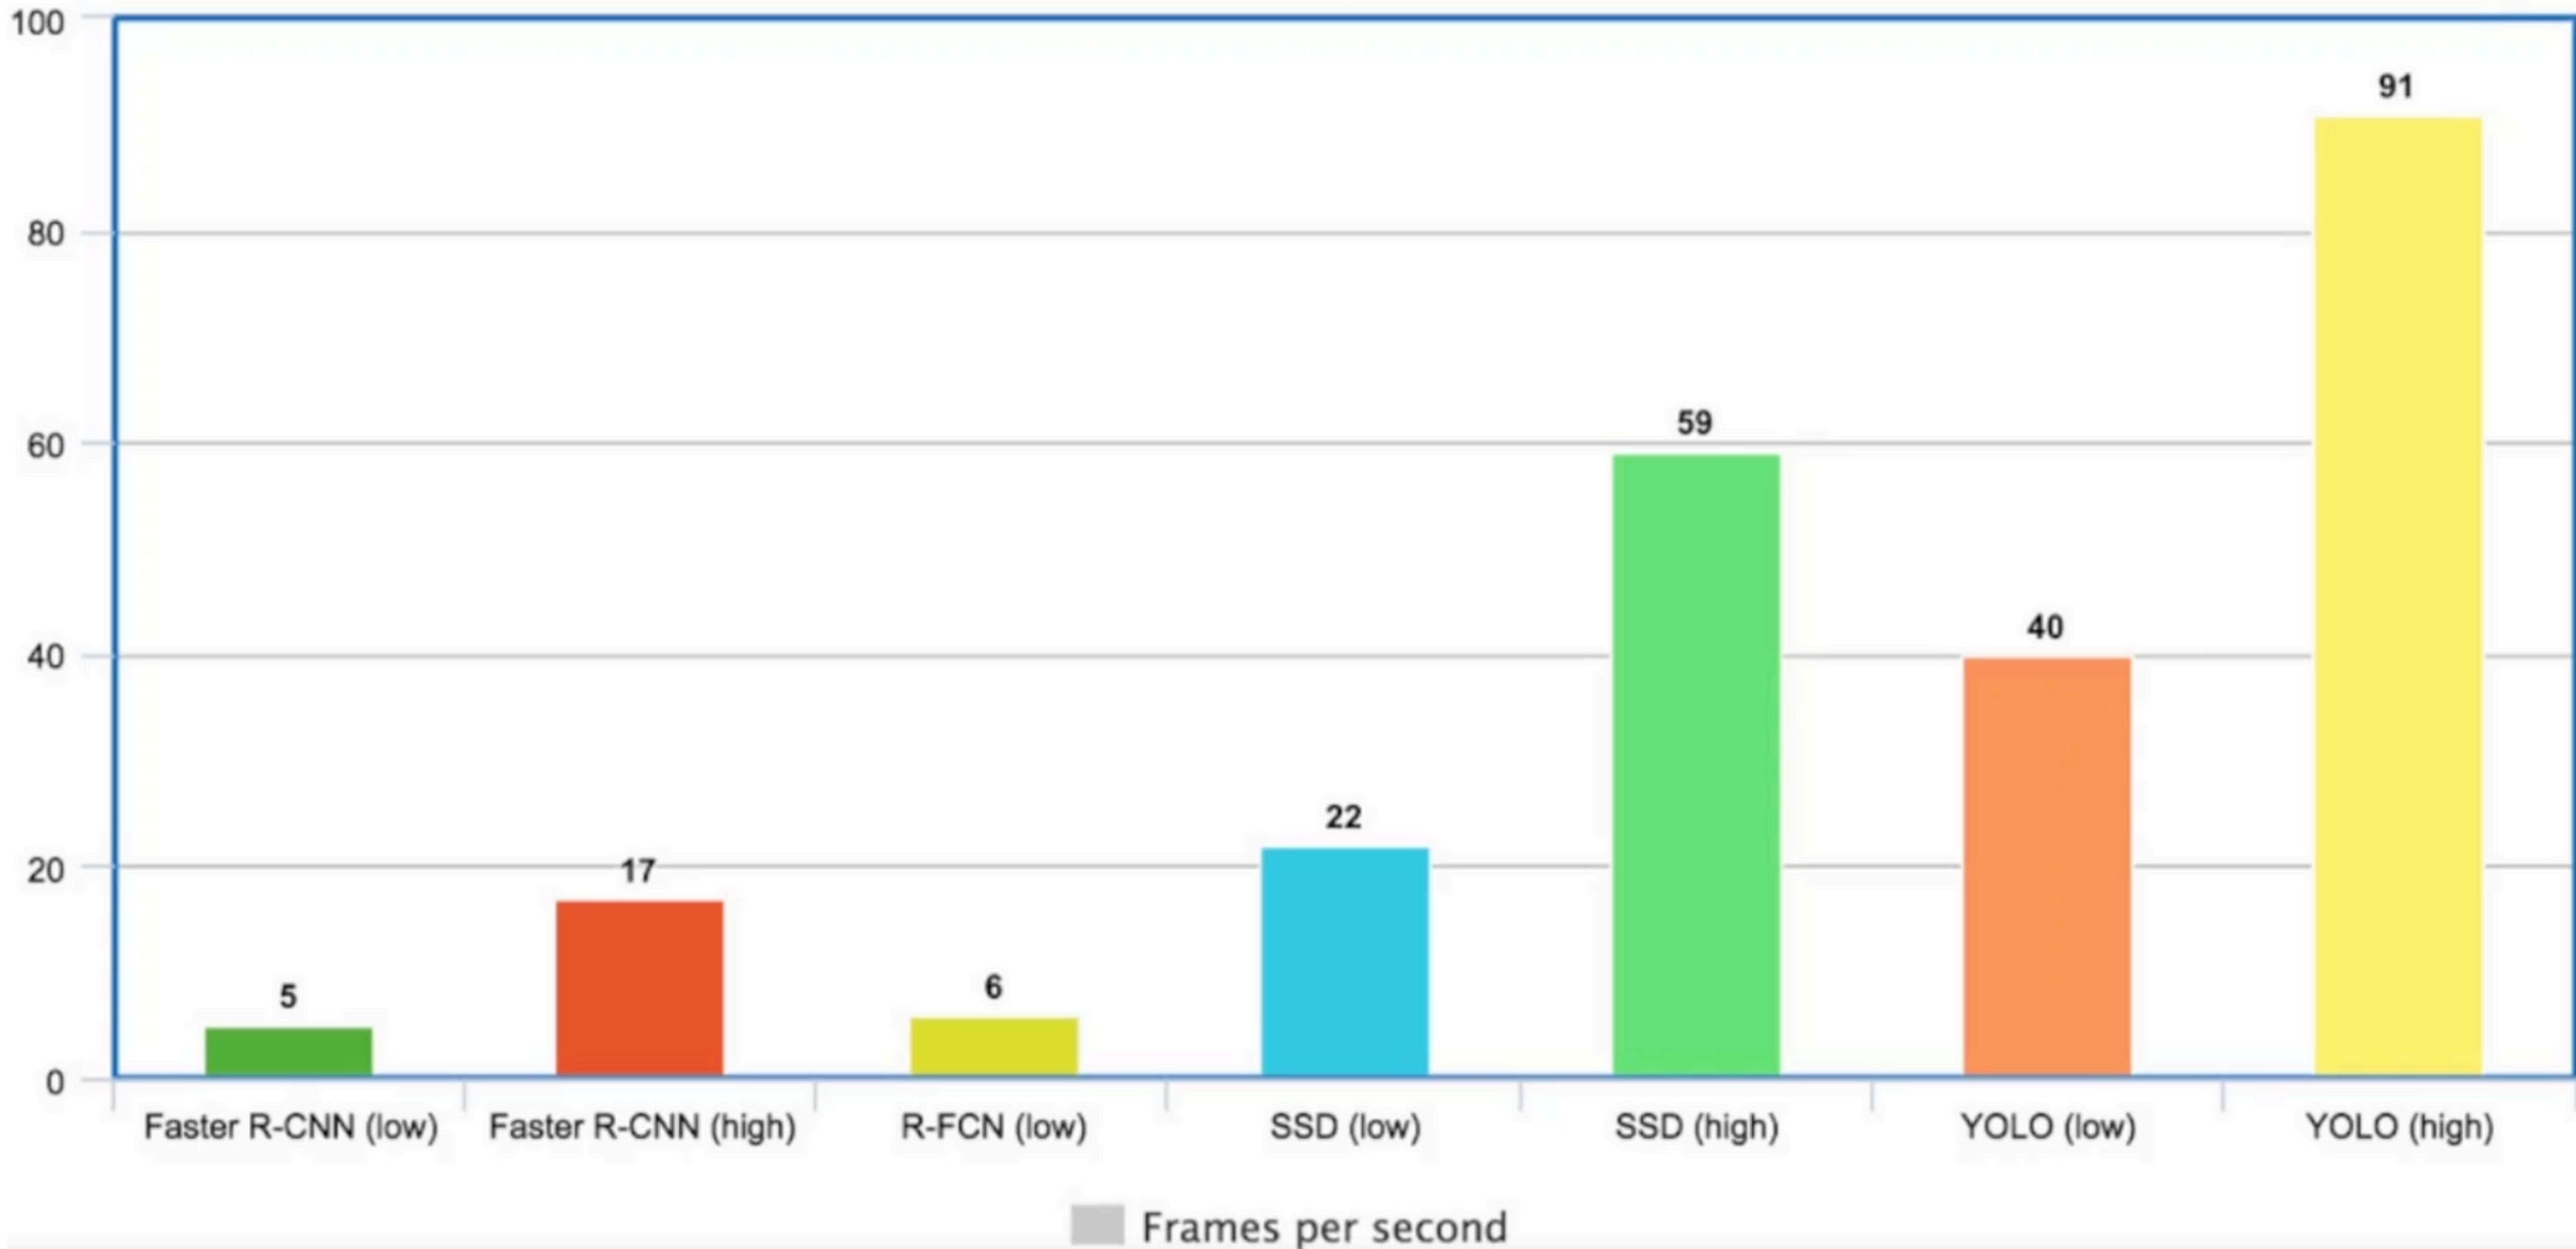


YOLO

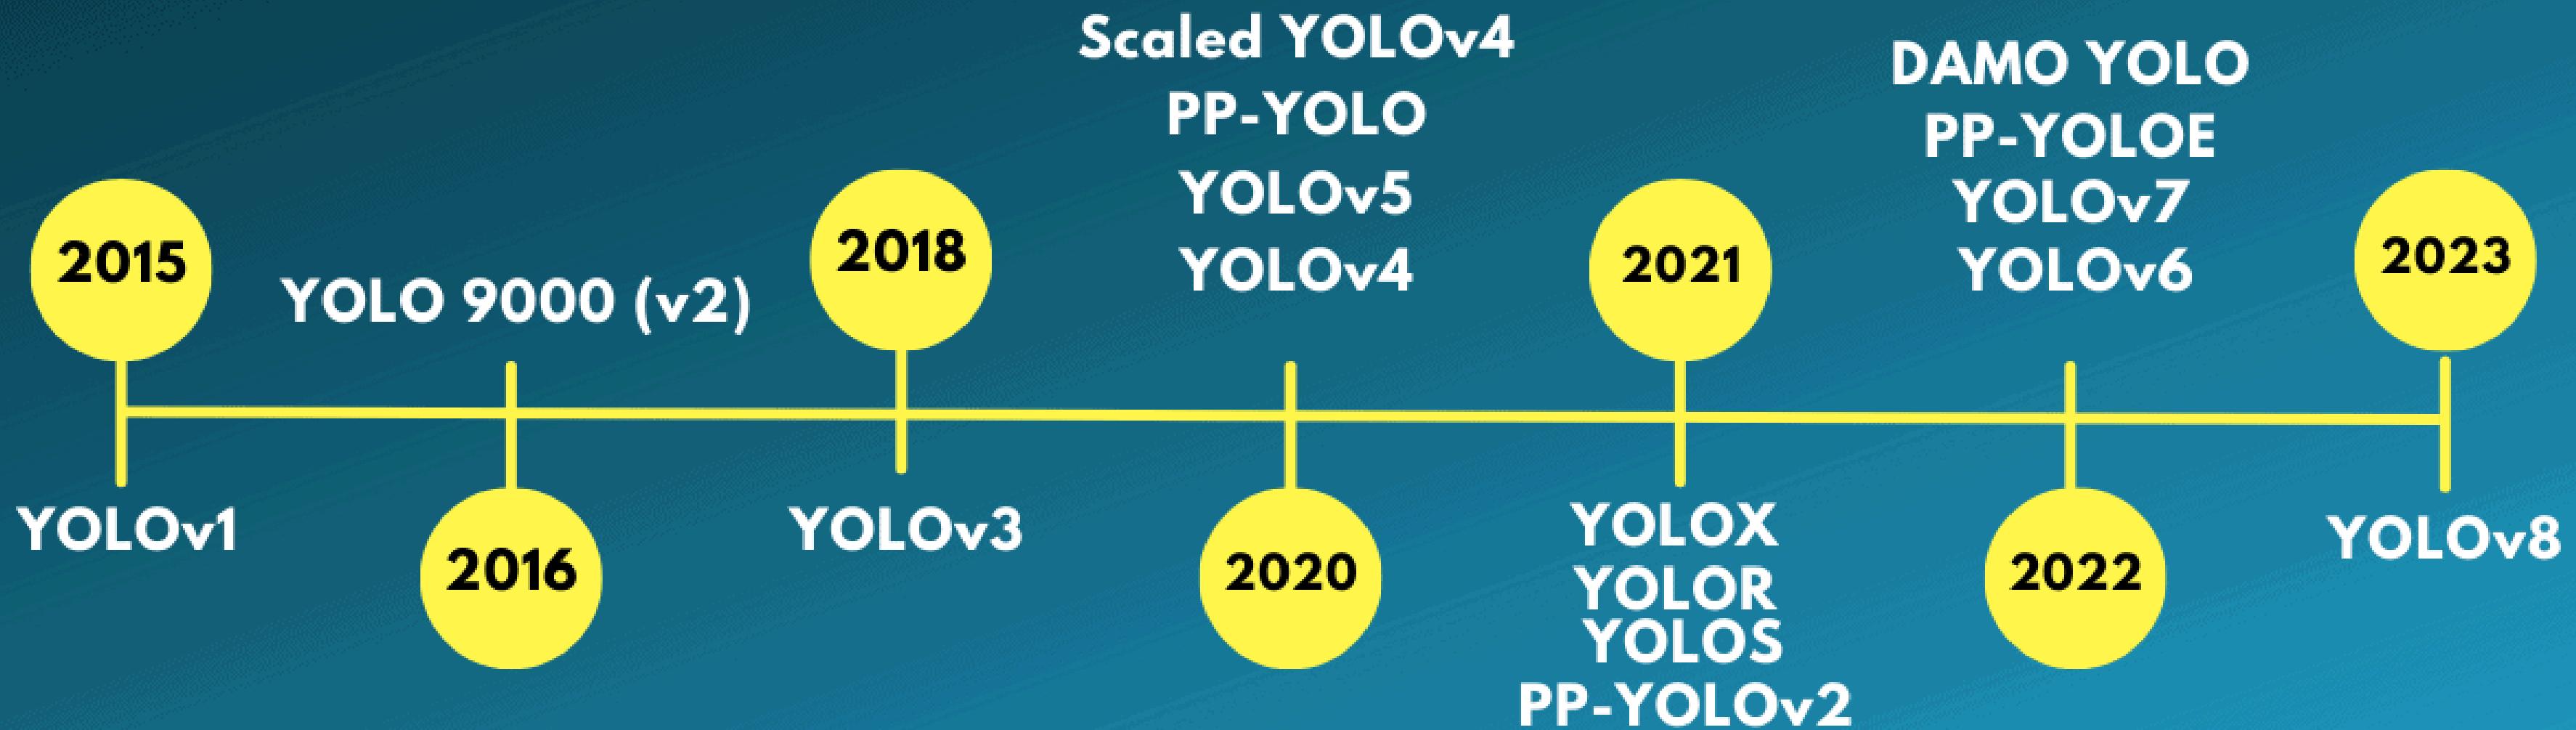
YOLO (You Only Look Once) has gained immense popularity in the field of object detection due to several key factors.

- **Speed:** YOLO is renowned for its real-time processing capabilities, allowing it to detect objects quickly by predicting bounding boxes and class probabilities from a *single forward pass* of the network.
- **Detection Accuracy:** Despite its speed, YOLO maintains high detection accuracy, effectively identifying objects in various conditions, which is crucial for applications requiring reliable performance.
- **Good Generalization:** YOLO exhibits strong generalization capabilities, enabling it to perform well across different datasets and object classes without extensive retraining.
- **Open Source:** Being an open-source framework allows developers and researchers to access the model easily, modify it for specific applications, and contribute to its continuous improvement, fostering a vibrant community around it.

These features collectively make YOLO a leading choice for object detection tasks across various domains.



YOLO Object Detection Models Timeline



YOLOV8

YOLOv8, developed by Ultralytics, is the latest advancement in the YOLO series, supporting object detection, image classification, and instance segmentation. Building on YOLOv5's legacy, YOLOv8 introduces enhanced architecture and a user-friendly experience.

- **Versatile Capabilities**

Supports object detection, image classification, and instance segmentation.

- **Developed by Ultralytics**

Creators of the groundbreaking YOLOv5, now with enhanced architecture and user experience.

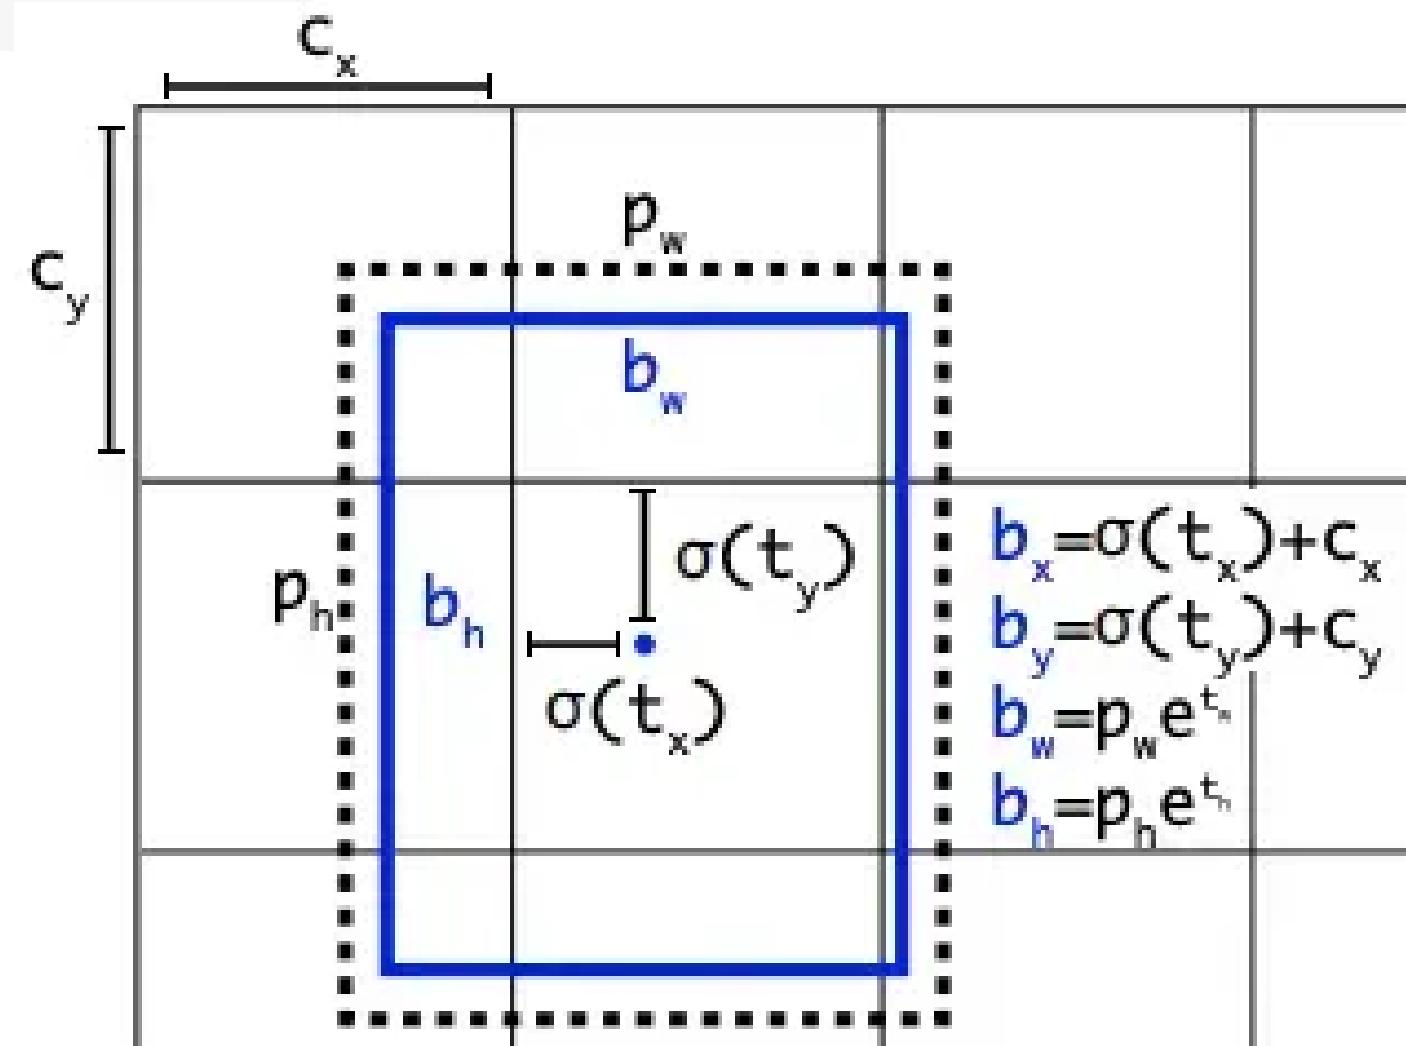
- **Community-Driven and Actively Supported**

Ongoing improvements, new features, and long-term support in collaboration with the community.

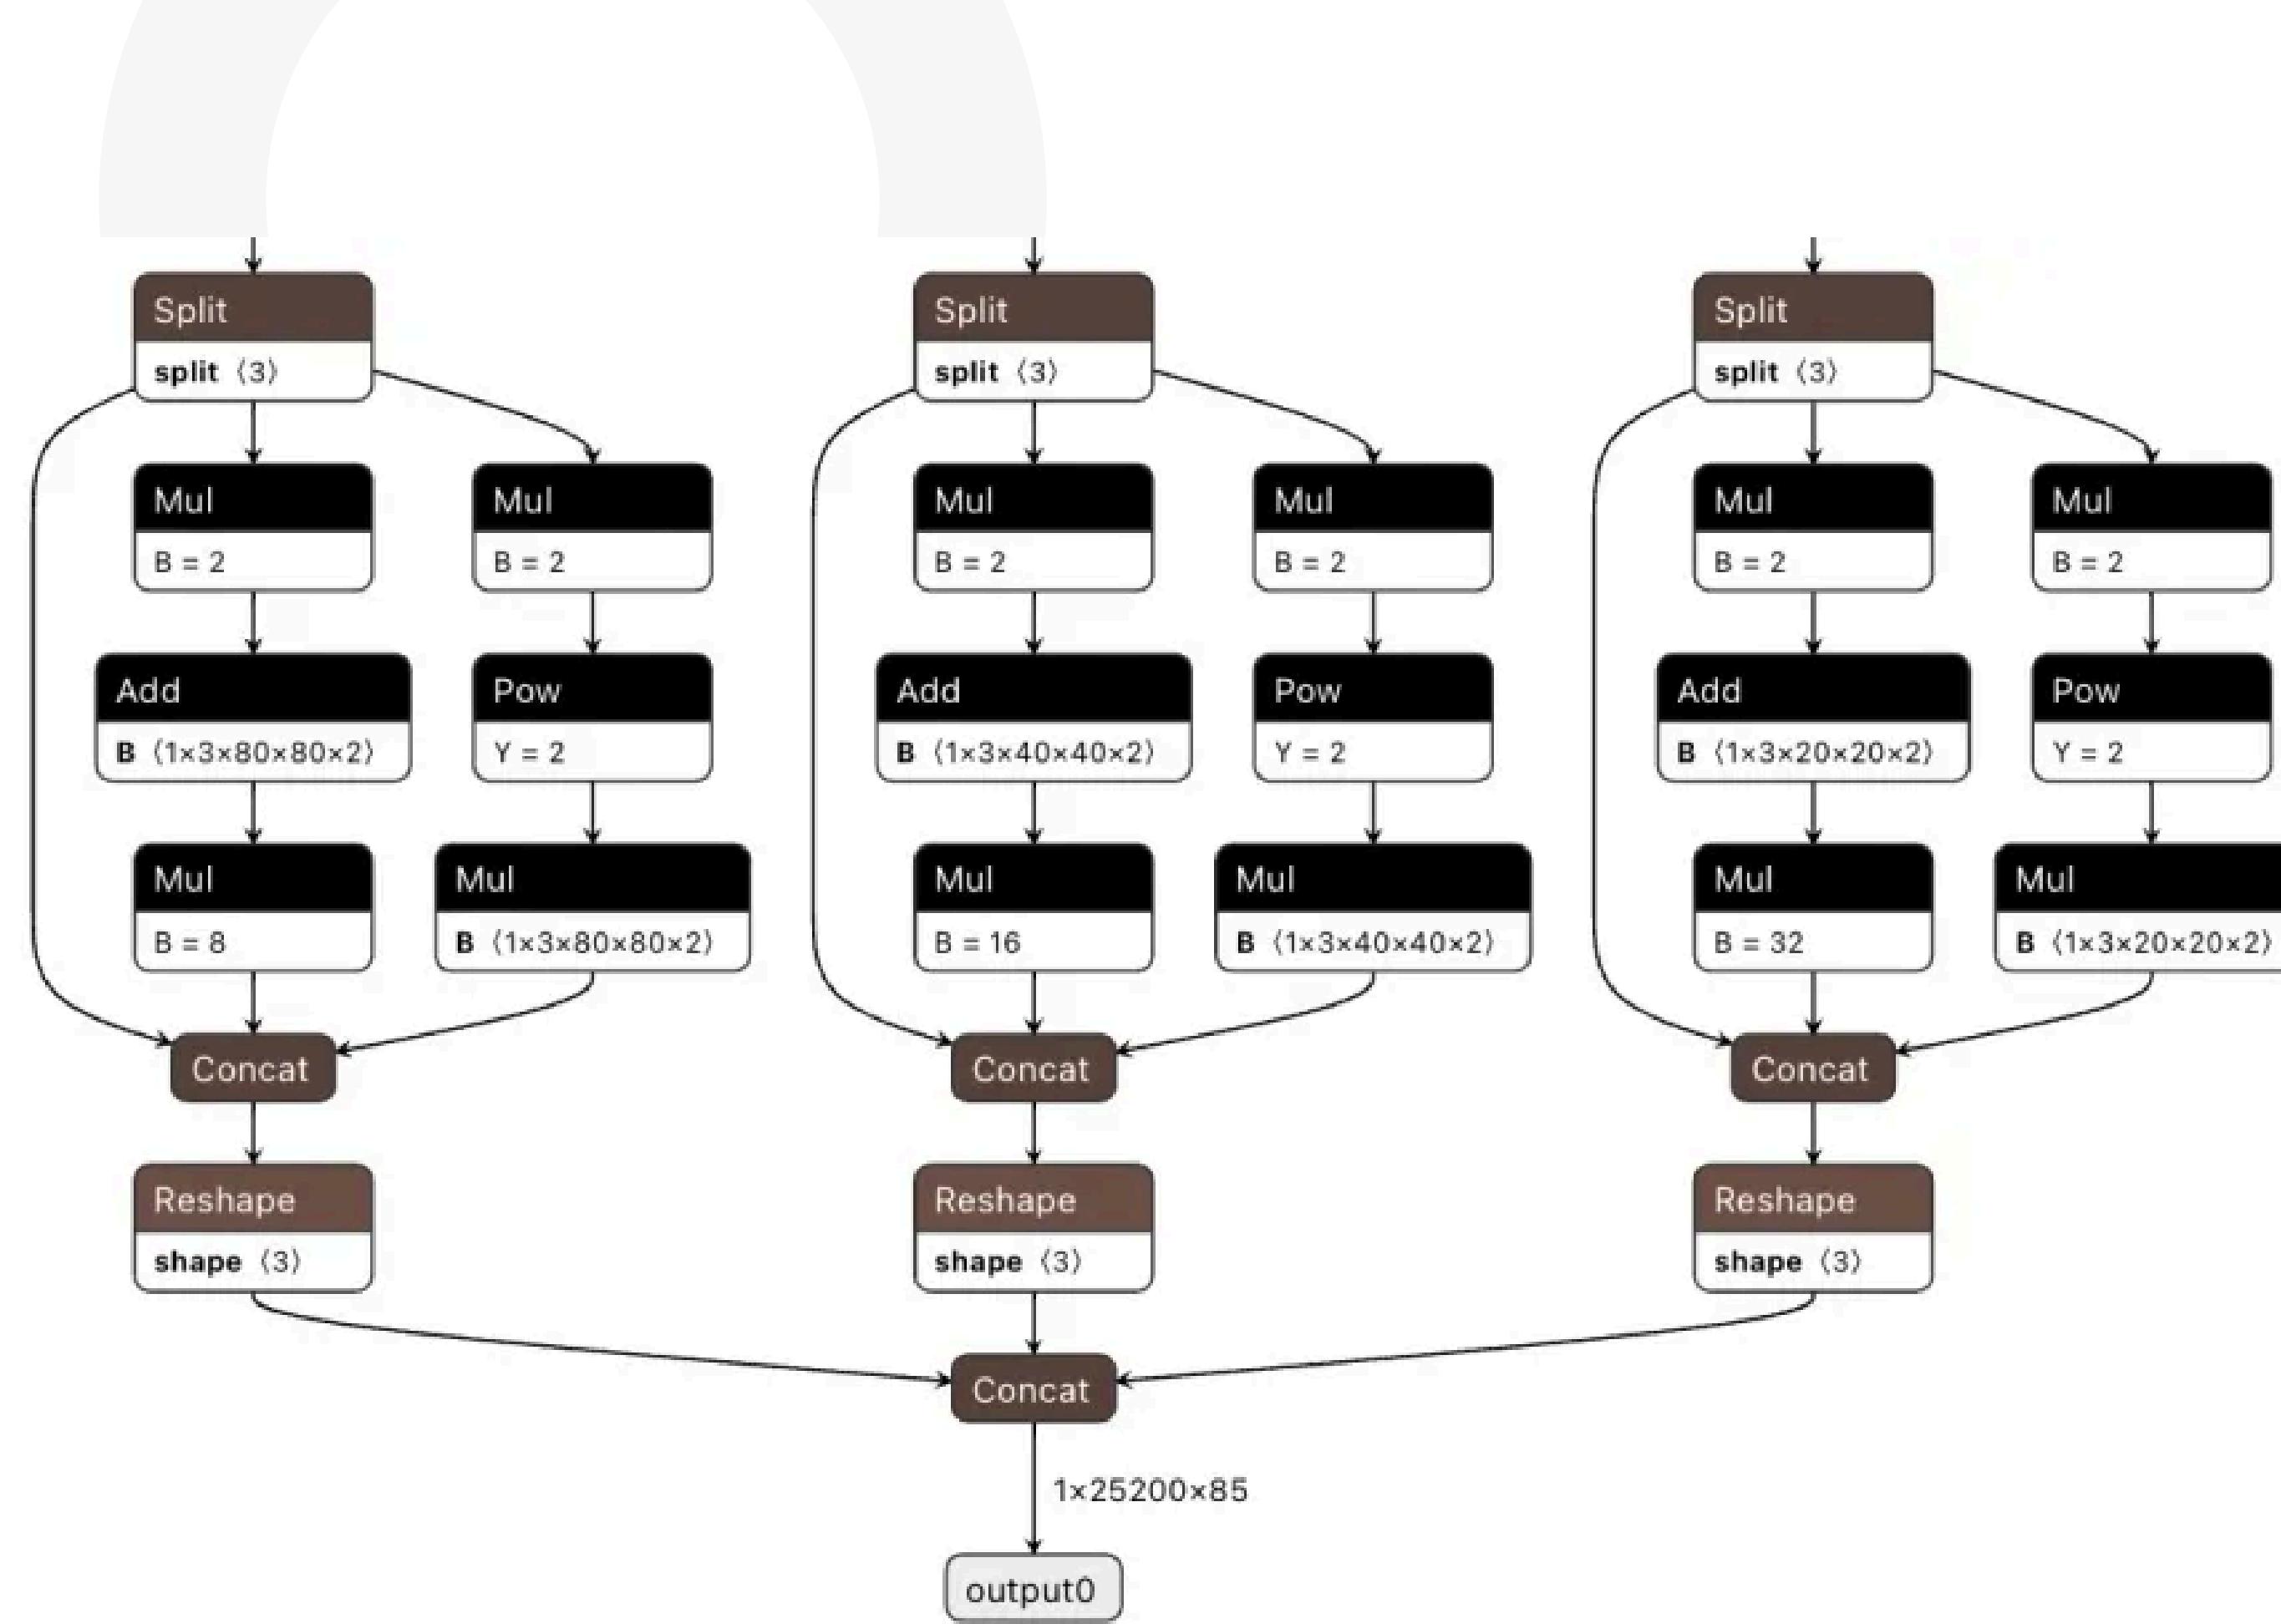
In summary, YOLOv8 excels in object detection and segmentation, combining advanced capabilities, intuitive design, and active community engagement, making it a powerful tool for diverse applications.

YOLOV8 – ANCHOR FREE DETECTION

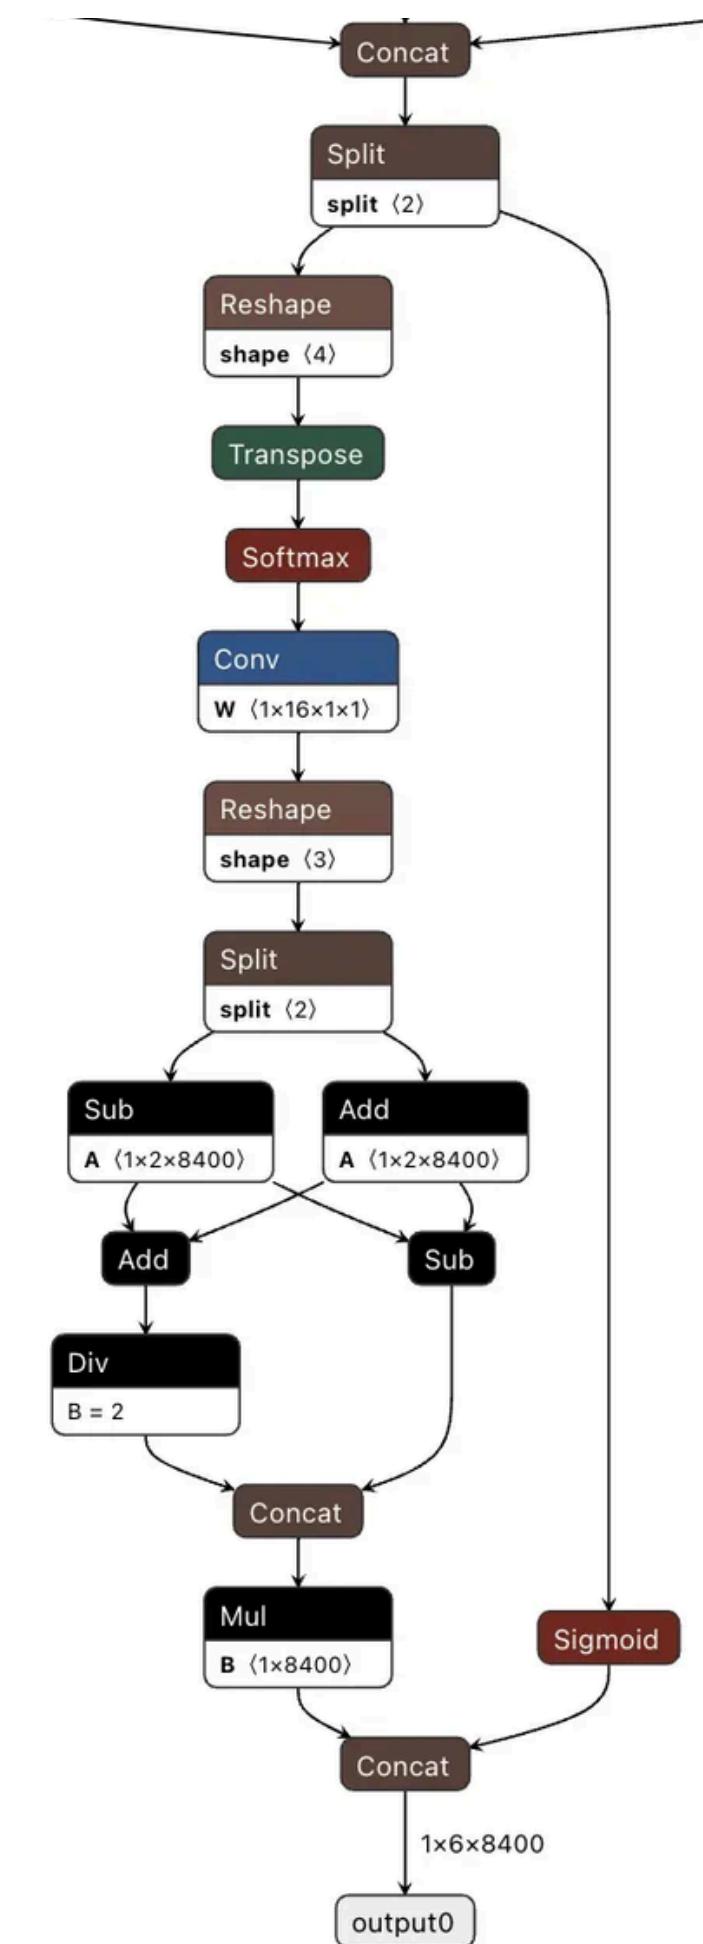
YOLOv8 is an anchor-free model. This means it predicts directly the center of an object instead of the offset from a known Anchor Box.



Visualization of an Anchor Box in YOLO



The detection head of YOLOv5, visualized in [netron.app](#)



The detection head of YOLOv8, visualized in netron.app

Object Detection Performance Comparison

(YOLOv8 vs YOLOv5)

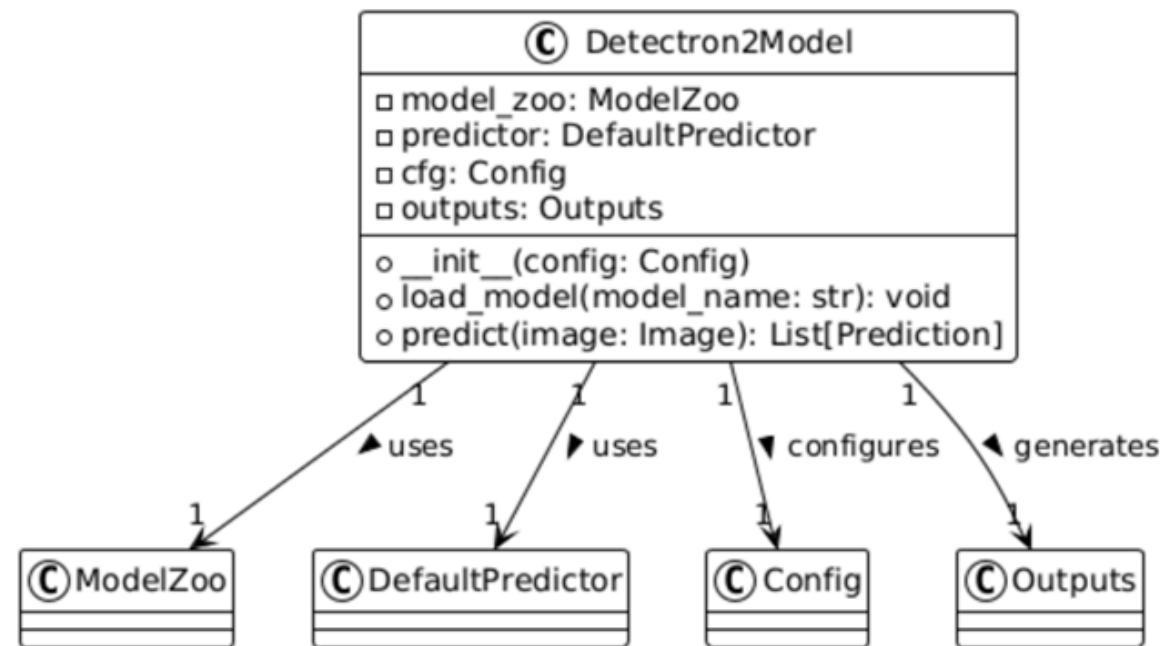
Model Size	YOLOv5	YOLOv8	Difference
Nano	28	37.3	+33.21%
Small	37.4	44.9	+20.05%
Medium	45.4	50.2	+10.57%
Large	49	52.9	+7.96%
Xtra Large	50.7	53.9	+6.31%

*Image Size = 640

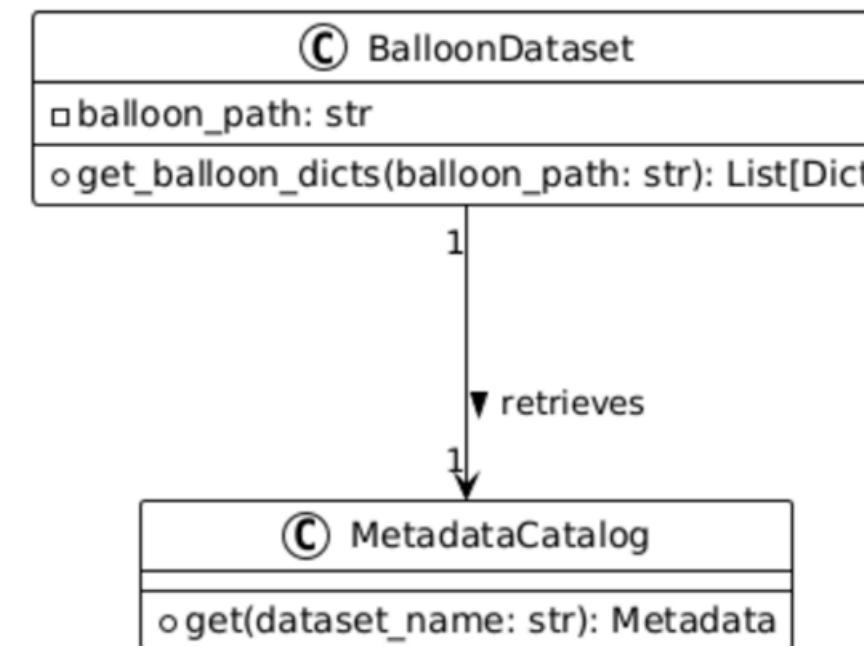
SYSTEM DESIGN

Class Diagram

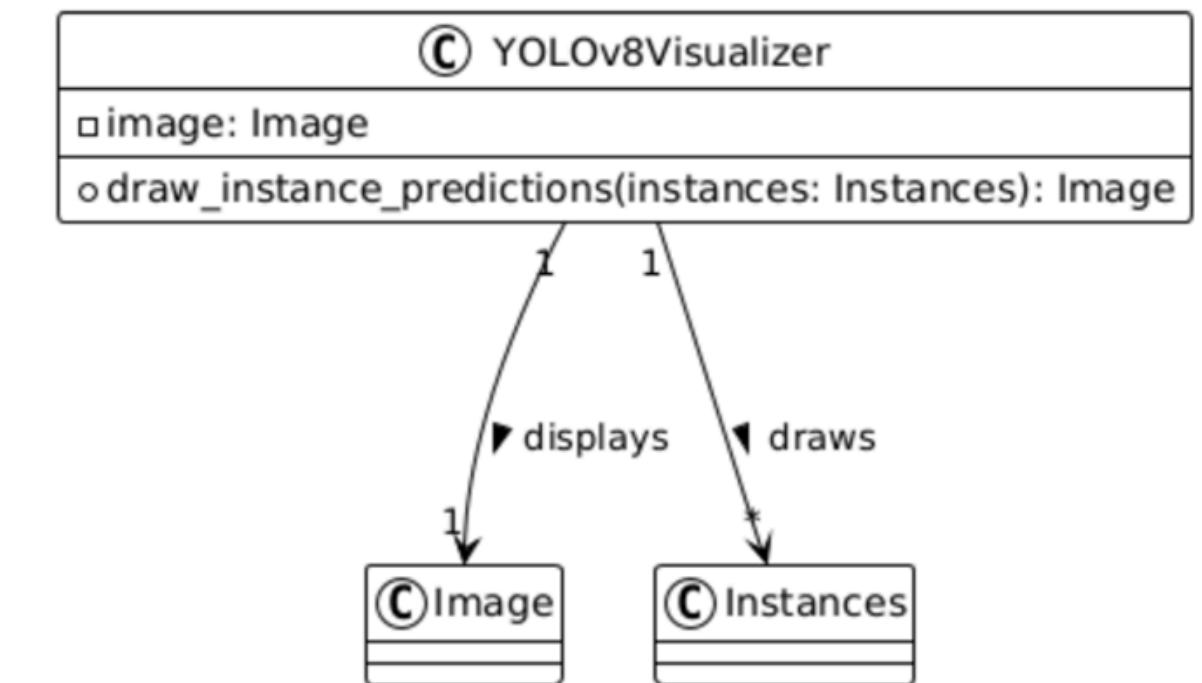
The class diagram represents the main components of the object detection system and their relationships. For the *Objectify* project, the key classes would be:



Class Diagram for Detectron2Model



Class Diagram for BalloonDataset

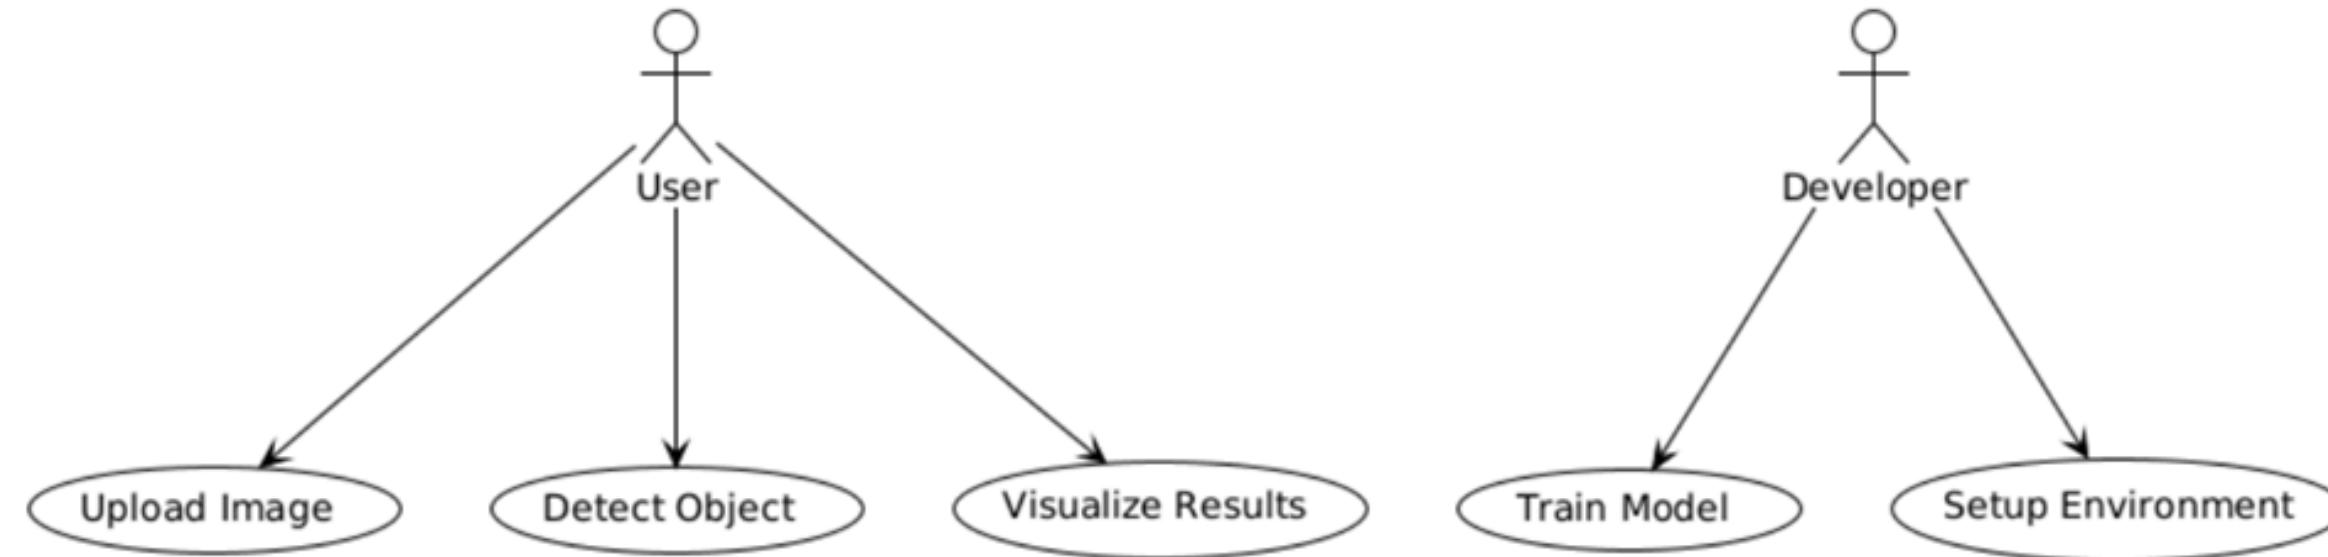


Class Diagram for YOLOv8Visualizer

SYSTEM DESIGN

Use Case Diagram

A use case diagram illustrates the key interactions between the system and external agents, in this case, users and images, the diagram can depict the core system use cases such as:

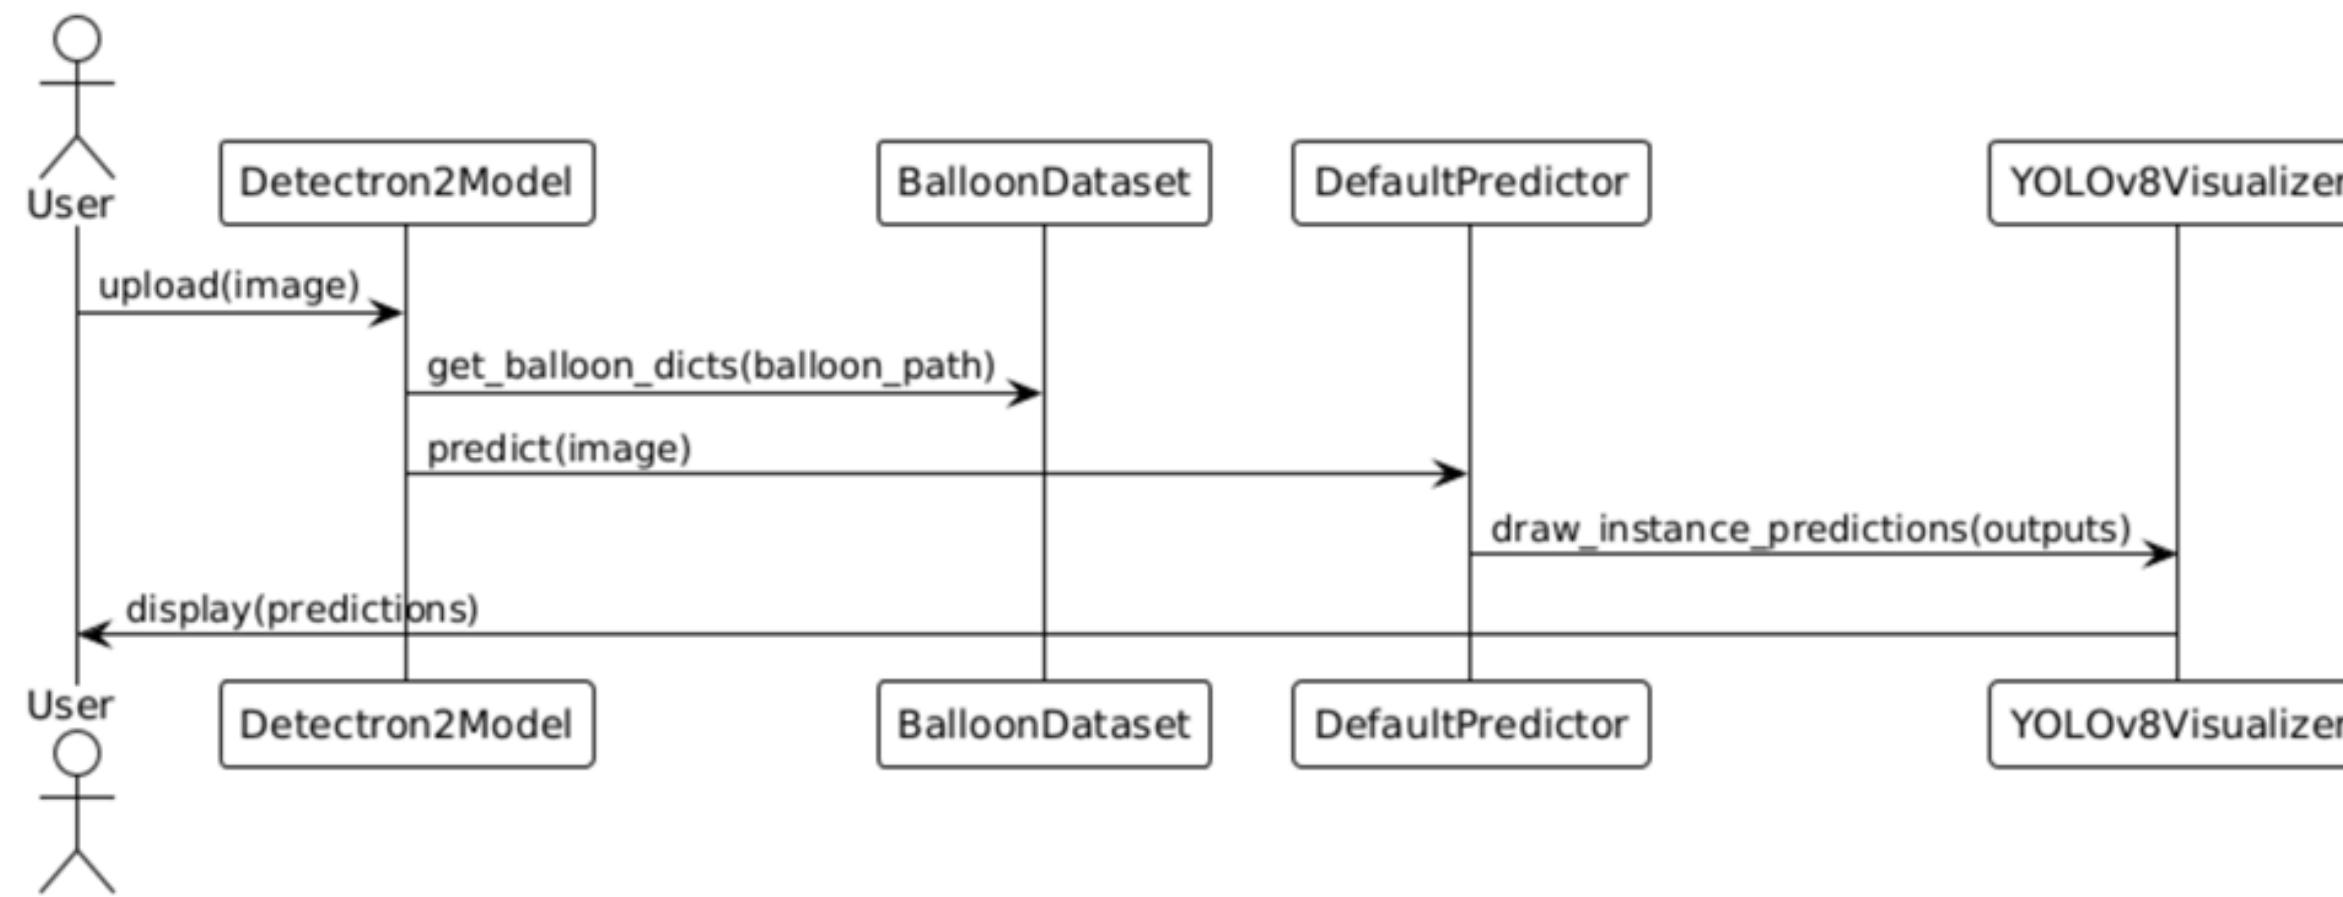


Use-case diagram

SYSTEM DESIGN

Sequence Diagram

The sequence diagram demonstrates the flow of interactions over time between system components when performing object detection. This can be simplified as:

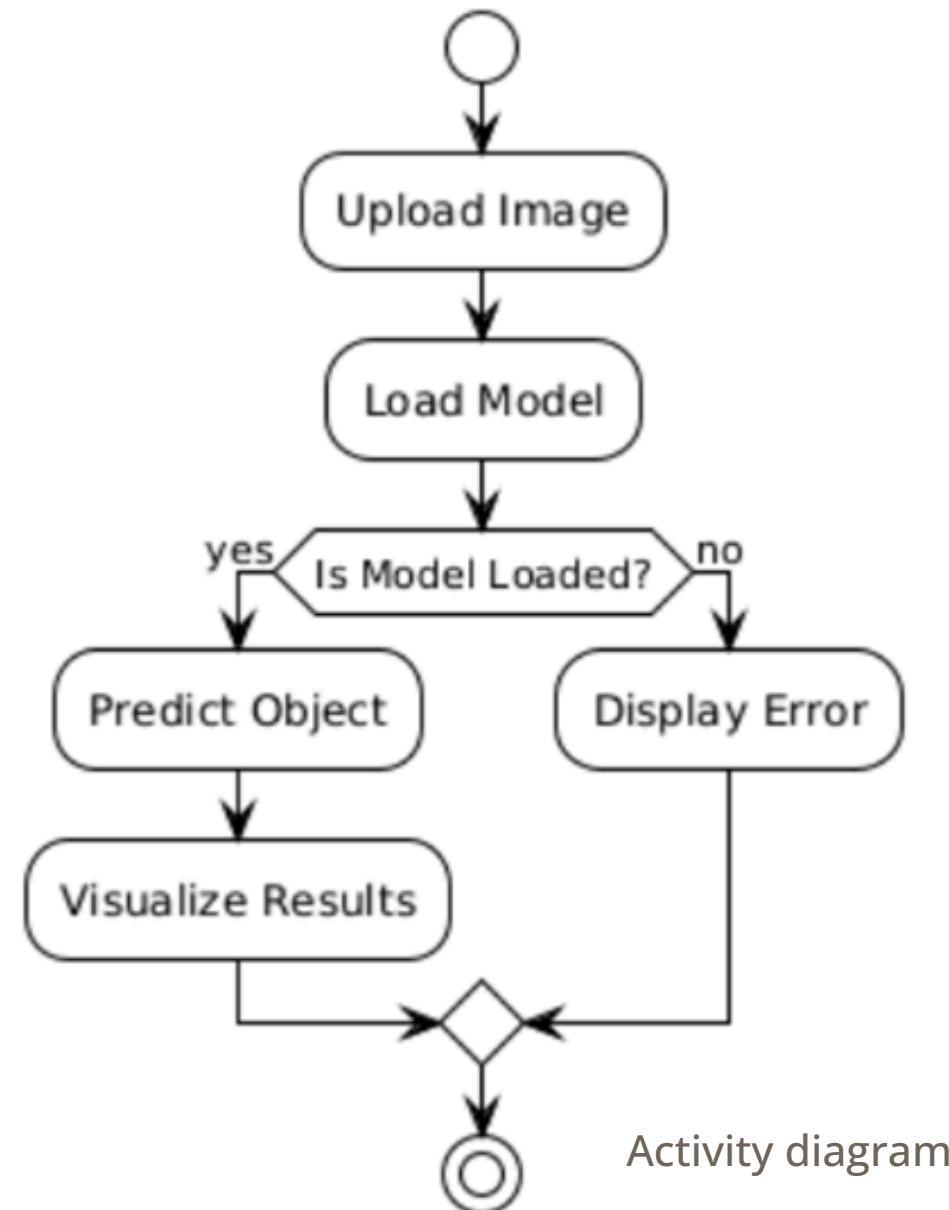


Sequence diagram

SYSTEM DESIGN

Activity Diagram

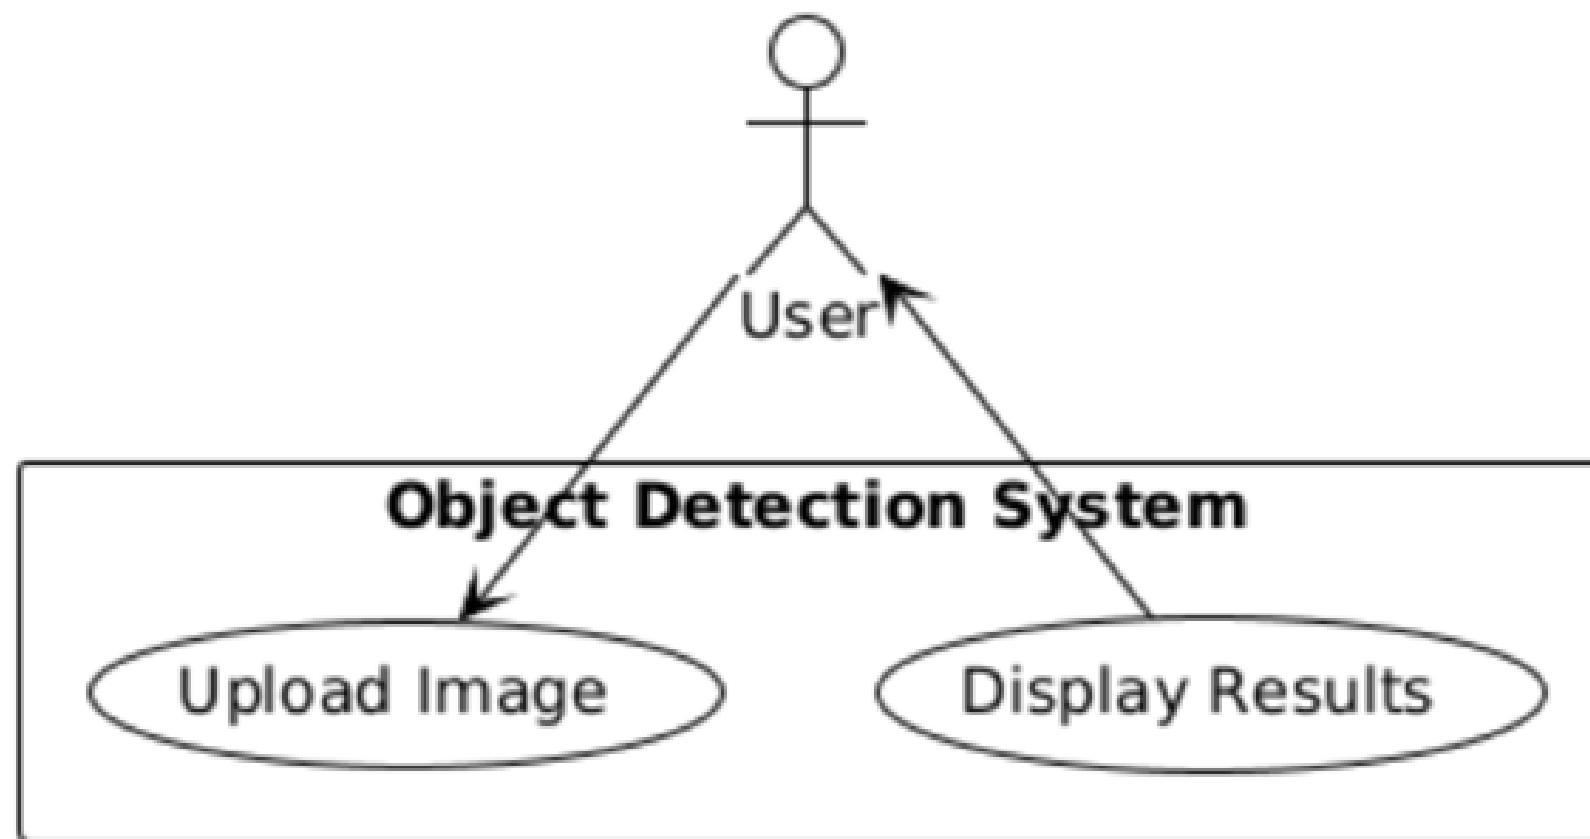
An activity diagram illustrates the step-by-step workflow of the system. For *Objectify*, the primary activities are:



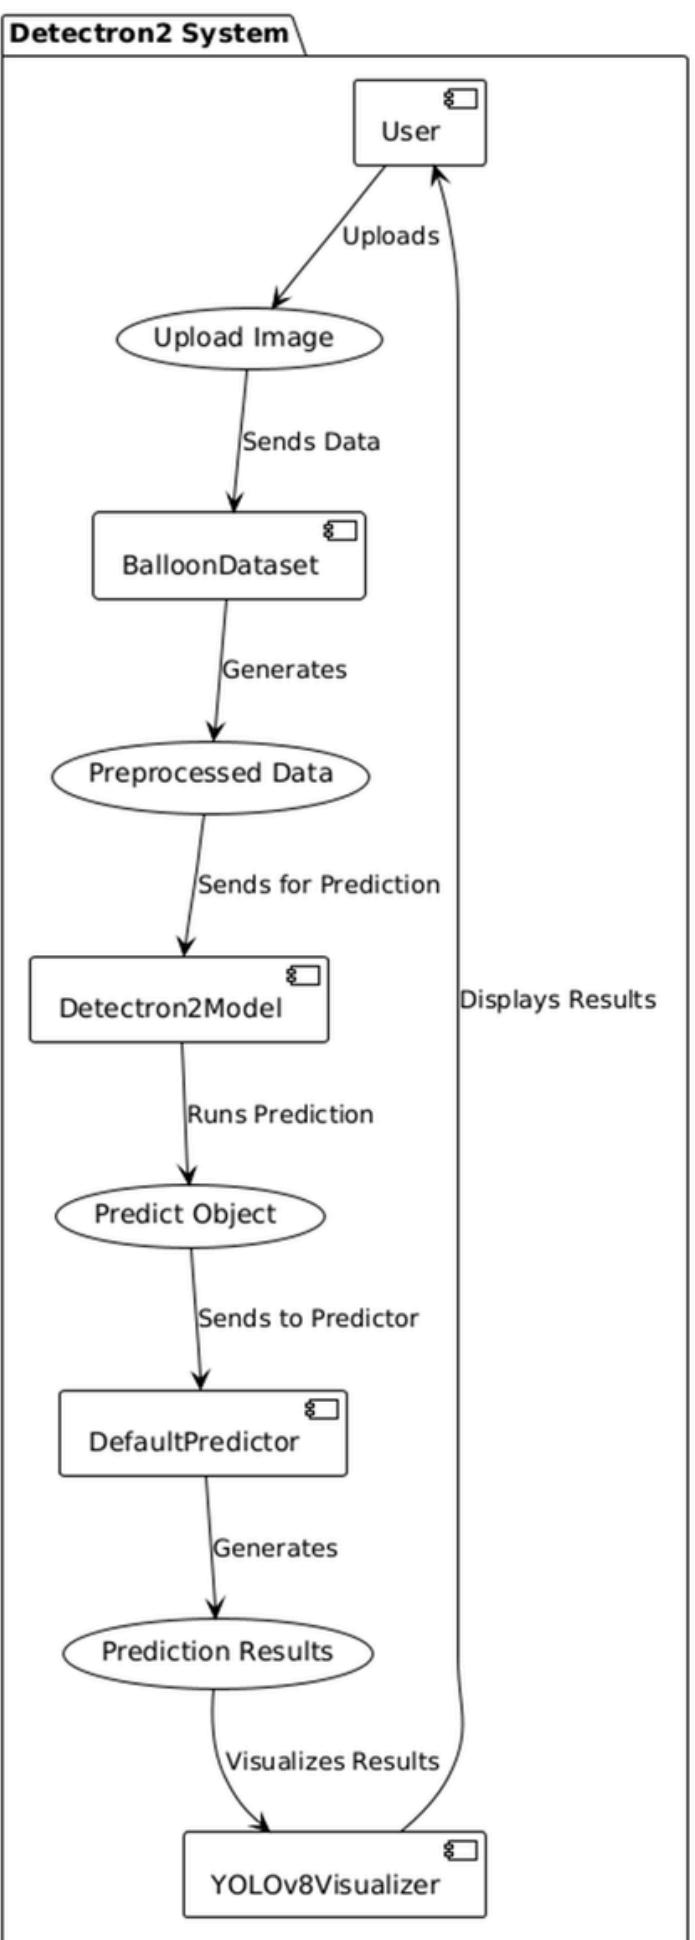
SYSTEM DESIGN

Data Flow Diagram (DFD)

The data flow diagram shows how data moves within the system. In this case:



Data Flow Diagram Level-0



Data Flow Diagram Level-1

MODULE DESCRIPTION – OBJECT DETECTION

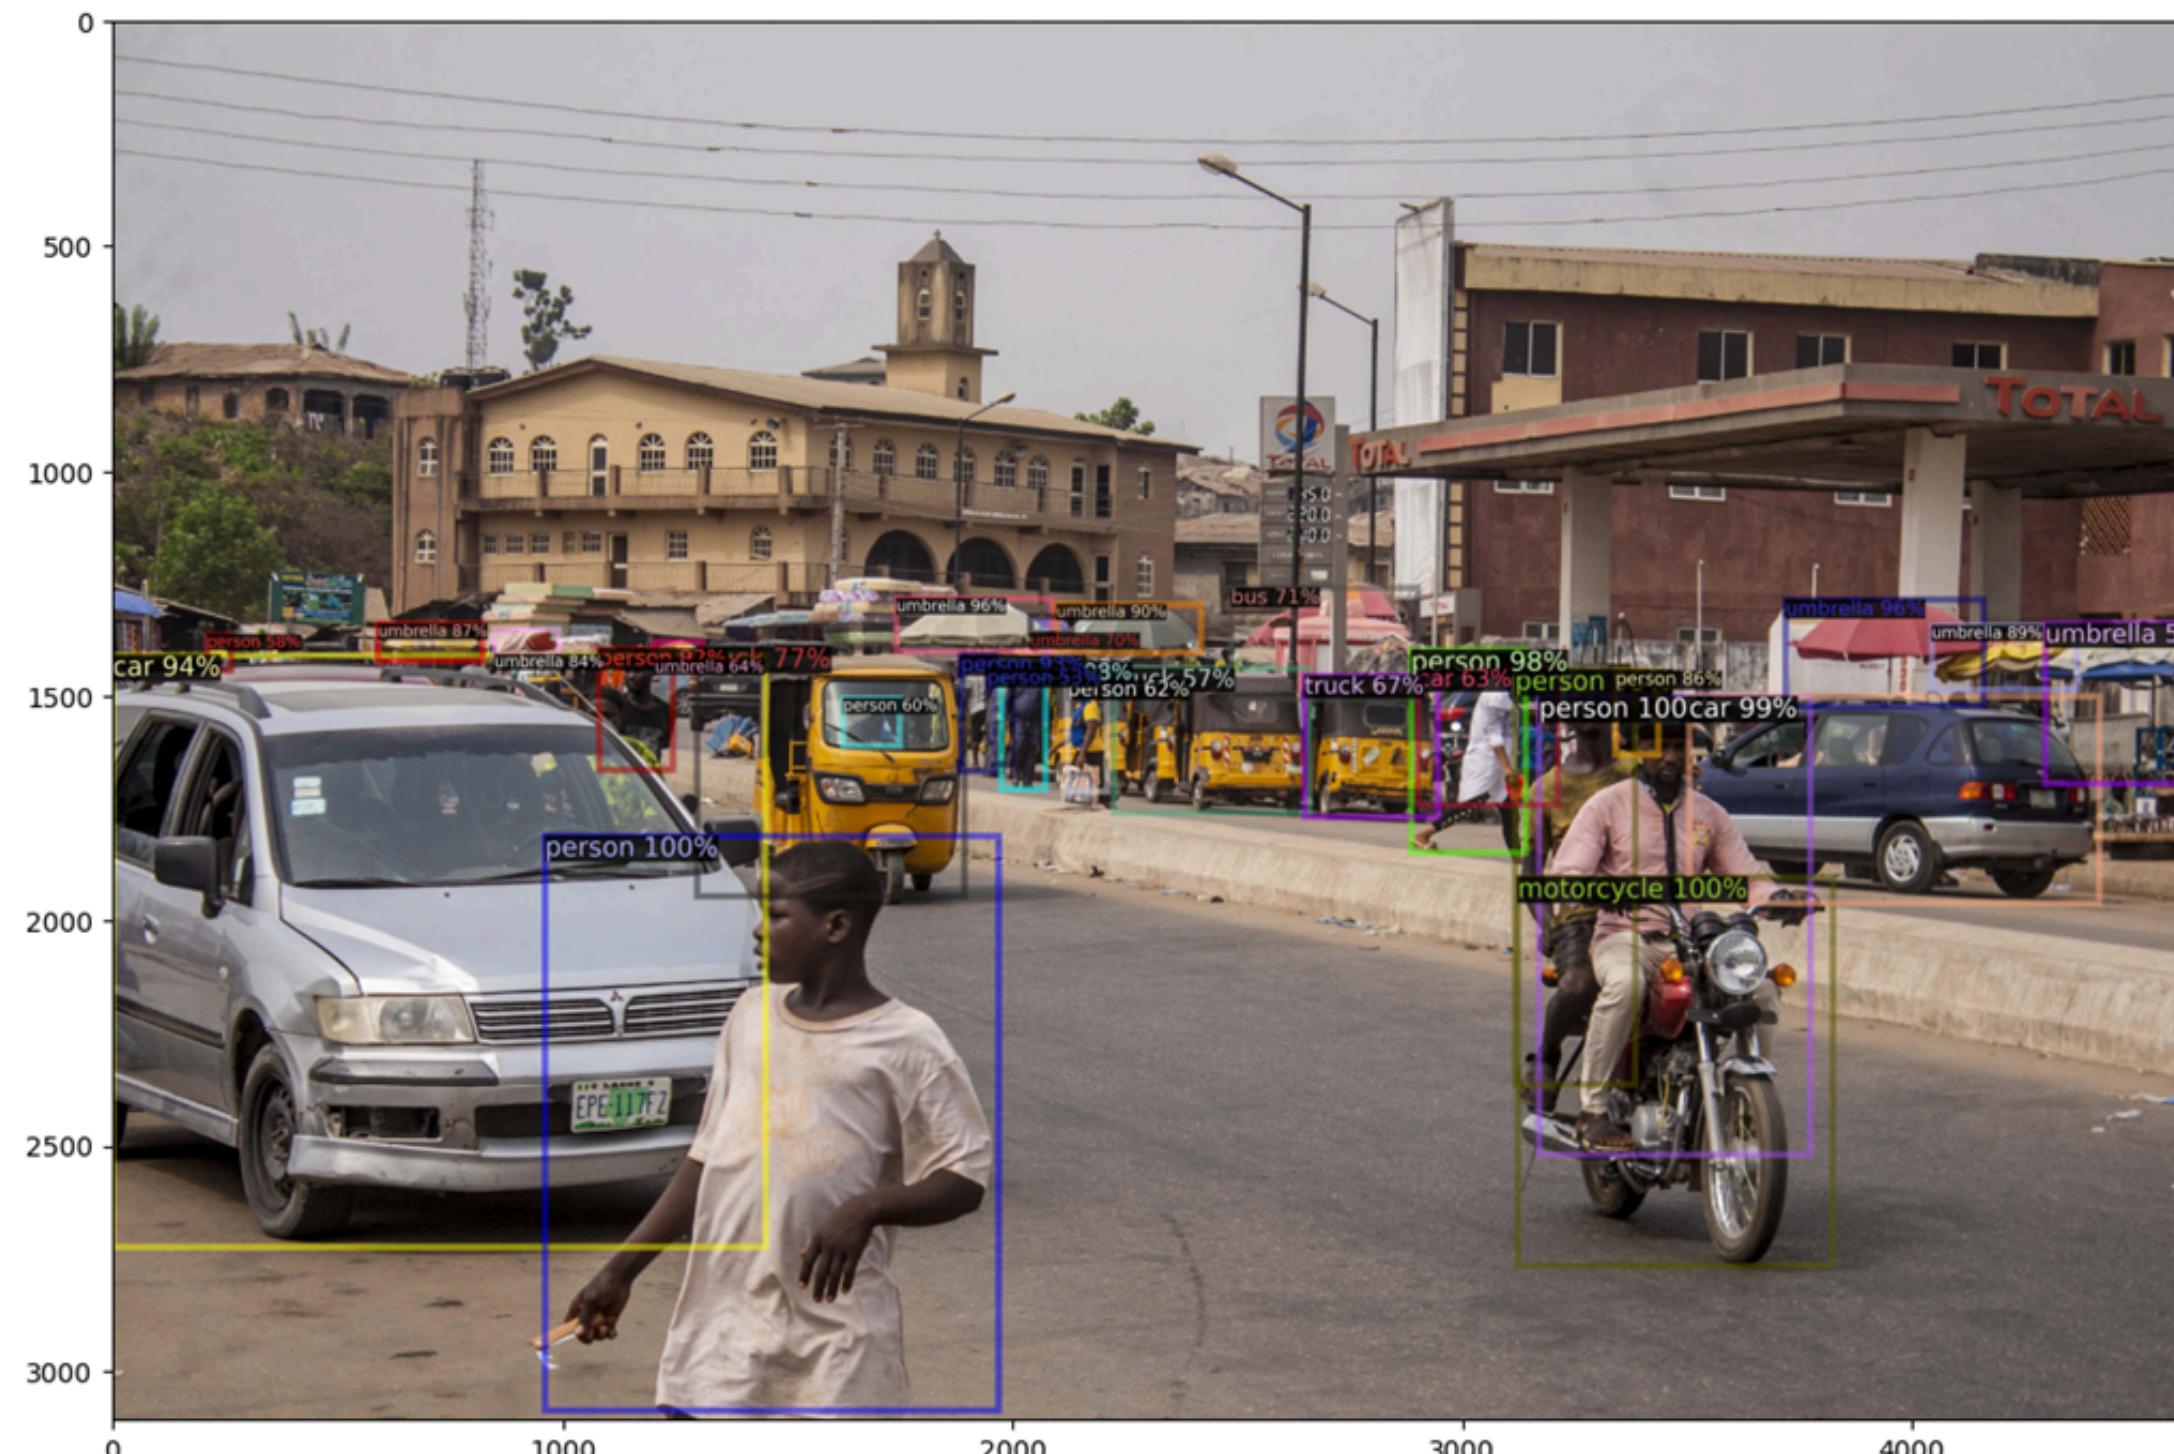
This is the core module responsible for **detecting** objects in an **image**. It leverages **pre-trained** machine learning models (e.g., YOLOv8) to process **uploaded** images and identify **objects** such as Car, Umbrella, Aeroplane, Kite.

- **Input:** Pre-processed image.
- **Output:** Detected objects with labels and confidence scores.



Pre-processed Image

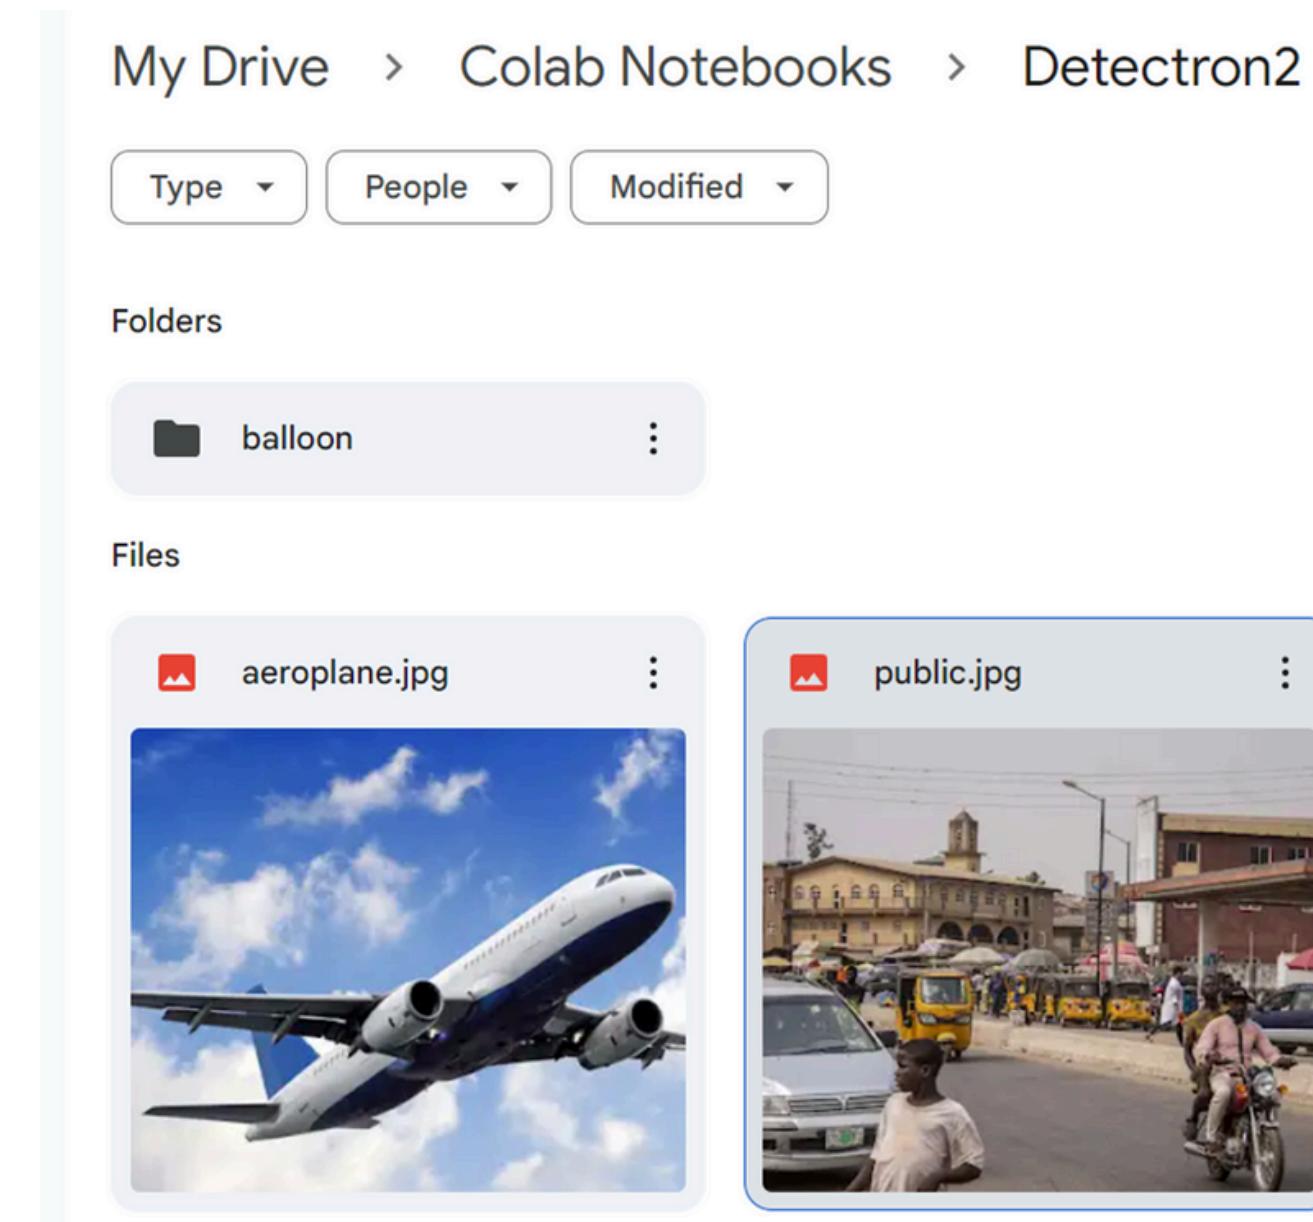
MODULE DESCRIPTION – OBJECT DETECTION



Detected Objects with Labels & Confidence Scores

MODULE DESCRIPTION – DATA PREPROCESSING MODULE

This module handles the necessary preprocessing steps before the object detection process. These steps ensure that images are in the correct format and size for accurate analysis by the model.



Raw image uploaded by the *user* in Google Drive

MODULE DESCRIPTION – DATA PREPROCESSING MODULE

The code snippet reads an image file called ‘public.jpg’ from a specified directory in Google Colab and displays it using OpenCV’s ‘cv2_imshow()’ function.

```
public = cv2.imread("/content/drive/MyDrive/Colab Notebooks/Detectron2/public.jpg")
cv2_imshow(public)
```

Key Functions:

- **resize(image):** Adjusts the image dimensions to match model requirements.
- **normalize(image):** Normalizes pixel values for consistency across datasets.
- **preprocess():** Applies multiple preprocessing steps, including resizing and normalization.

MODULE DESCRIPTION – MODEL TRAINING MODULE

BalloonDataset Module

This module manages the training of the object detection model on **custom datasets** for Balloon. It allows for improving model performance by training on additional annotated images.

Key Functions:

- **trainModel(coco_2017_val)**: Trains the Detectron2 model on the COCO dataset.
- **saveModel()**: Saves the trained model **class weights** for future use.
- **loadModel()**: Loads an **existing** class model for further training or evaluation.

```
metadata = MetadataCatalog.get("coco_2017_val")
class_names = metadata.get("thing_classes")

print("COCO Dataset number of class: ", len(class_names))
print("COCO Dataset class names: ", class_names)
```

MODULE DESCRIPTION – MODEL TRAINING MODULE

COCO Dataset number of class: 80

COCO Dataset class names: ['person', 'bicycle', 'car', 'motorcycle', 'airplane', 'bus', 'train', 'truck', 'boat', 'traffic light', 'fire hydrant', 'stop sign', 'parking meter', 'bench', 'bird', 'cat', 'dog', 'horse', 'sheep', 'cow', 'elephant', 'bear', 'zebra', 'giraffe', 'backpack', 'umbrella', 'handbag', 'tie', 'suitcase', 'frisbee', 'skis', 'snowboard', 'sports ball', 'kite', 'baseball bat', 'baseball glove', 'skateboard', 'surfboard', 'tennis racket', 'bottle', 'wine glass', 'cup', 'fork', 'knife', 'spoon', 'bowl', 'banana', 'apple', 'sandwich', 'orange', 'broccoli', 'carrot', 'hot dog', 'pizza', 'donut', 'cake', 'chair', 'couch', 'potted plant', 'bed', 'dining table', 'toilet', 'tv', 'laptop', 'mouse', 'remote', 'keyboard', 'cell phone', 'microwave', 'oven', 'toaster', 'sink', 'refrigerator', 'book', 'clock', 'vase', 'scissors', 'teddy bear', 'hair drier', 'toothbrush']

The COCO (Common Objects in Context) dataset contains 80 classes commonly used for object detection tasks. The classes represent a variety of objects from daily life, categorized into humans, vehicles, animals, and household items, among others.

Number of classes: 80

- **People:** person
- **Vehicles:** bicycle, car, motorcycle, airplane, bus, train, truck, boat
- **Traffic:** traffic light, fire hydrant, stop sign, parking meter
- **Furniture:** chair, couch, bed, dining table
- **Electronics:** tv, laptop, mouse, remote, keyboard, cell phone
- **Animals:** bird, cat, dog, horse, sheep, cow, elephant, zebra, giraffe
- **Food:** banana, apple, sandwich, orange, broccoli, carrot, pizza, donut
- **Miscellaneous:** book, clock, vase, scissors, teddy bear, hair drier, toothbrush

These 80 classes serve as **labels** for object detection tasks, providing diverse **examples** to train models.

MODULE DESCRIPTION – YOLOV8 VISUALIZER MODULE

The model configuration defines the structure and parameters for training the YOLOv8 model. The following YAML configuration file was used to set up the dataset paths and class names for this project:

```
train: /content/drive/MyDrive/Colab Notebooks/ObjectDetection/Dataset/train/images
val: /content/drive/MyDrive/Colab Notebooks/ObjectDetection/Dataset/valid/images
#test: /content/drive/MyDrive/Colab Notebooks/ObjectDetection/Dataset/test/images #optional

nc: 3
names: ['Ball', 'Player', 'Referee']

#roboflow:
#url: https://universe.roboflow.com/nikhil-chapre-xgndf/detect-players-dgxz0/dataset/7
```

YOLOv8 “dataset.yaml”

This configuration includes three object classes: *Ball*, *Player*, and *Referee*, and the paths for the training and validation datasets. These configurations are critical for ensuring the model is trained on the correct dataset.

MODULE DESCRIPTION – ROBOFLOW

For labelling and annotating the video clips and images, the Roboflow platform was used. Roboflow allows for efficient dataset preprocessing, annotation, and labelling, ensuring high-quality inputs for model training.

- Roboflow Dataset URL: [Roboflow Dataset](#)
- The dataset was labeled using Roboflow's annotation tools, creating labels for three objects: *Ball*, *Player*, and *Referee*.
- Once annotated, the dataset was exported in the required format for use in model training and validation.

The Roboflow logo is displayed in a large, bold, purple sans-serif font. The word "roboflow" is written in a lowercase, monospaced-style font, where each letter has a distinct vertical stroke. The letters are slightly rounded at the top and bottom.

MODULE DESCRIPTION – ROBOFLOW

Below is an example of how the labeling process was performed in Roboflow for *Players*



'Player' Labeling using Roboflow

MODULE DESCRIPTION – ROBOFLOW

Below is an example of how the labeling process was performed in Roboflow for Referee



'Referee' Labeling using Roboflow

MODULE DESCRIPTION – ROBOFLOW

Below is an example of how the labeling process was performed in Roboflow for Ball



'Ball' Labeling using Roboflow

MODULE DESCRIPTION – EVALUATION MODULE

This module assesses the performance of the trained object detection model. It evaluates the model's accuracy and effectiveness on validation datasets.

Input: Validation dataset, trained model.

Output: Performance metrics (e.g., precision, recall, mAP).

```
from detectron2.evaluation import COCOEvaluator, inference_on_dataset
from detectron2.data import build_detection_test_loader

evaluator = COCOEvaluator("balloon_val", ("bbox",), False, output_dir="./output/")
val_loader = build_detection_test_loader(cfg, "balloon_val")

print(inference_on_dataset(trainer.model, val_loader, evaluator))
```

COCO Evaluation Process

MODULE DESCRIPTION – INFERENCE & EVALUATION MODULE

During the evaluation, the following metrics are generated.

Average Precision (AP)

oAP @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.771
oAP @[IoU=0.50 | area= all | maxDets=100] = 0.900
oAP @[IoU=0.75 | area= all | maxDets=100] = 0.844
oAP @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.250
oAP @[IoU=0.50:0.95 | area= medium | maxDets=100] = 0.599
oAP @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.906

Average Recall (AR)

oAR @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.244
oAR @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.786
oAR @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.804
oAR @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.467
oAR @[IoU=0.50:0.95 | area= medium | maxDets=100] = 0.653
oAR @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.923

generateMetricsReport()

Compiles and outputs a report of the evaluation results. The report is generated based on the metrics obtained during evaluation, which includes detailed information about the Average Precision and Average Recall.

MODULE DESCRIPTION – EVALUATION MODULE

This module is responsible for performing inference and evaluation using the trained object detection model. It loads the model weights, sets the evaluation parameters, and visualizes the results.

Load Model

- `cfg.MODEL.WEIGHTS`: Specifies the path to the trained model weights (e.g., "model_final.pth").
- `cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST`: Sets the score threshold for making predictions (e.g., 0.7).
- `cfg.DATASETS.TEST`: Defines the dataset to be used for evaluation (e.g., ("balloon_val",)).

```
cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth")
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.7    # set the testing threshold for this model
cfg.DATASETS.TEST = ("balloon_val", )
predictor = DefaultPredictor(cfg)
```

Inference Load Process

MODULE DESCRIPTION – EVALUATION MODULE

This module is responsible for performing inference and evaluation using the trained object detection model. It loads the model weights, sets the evaluation parameters, and visualizes the results.

Perform Inference

- `outputs = predictor(im)`: Executes inference on the input image, returning the predicted outputs.

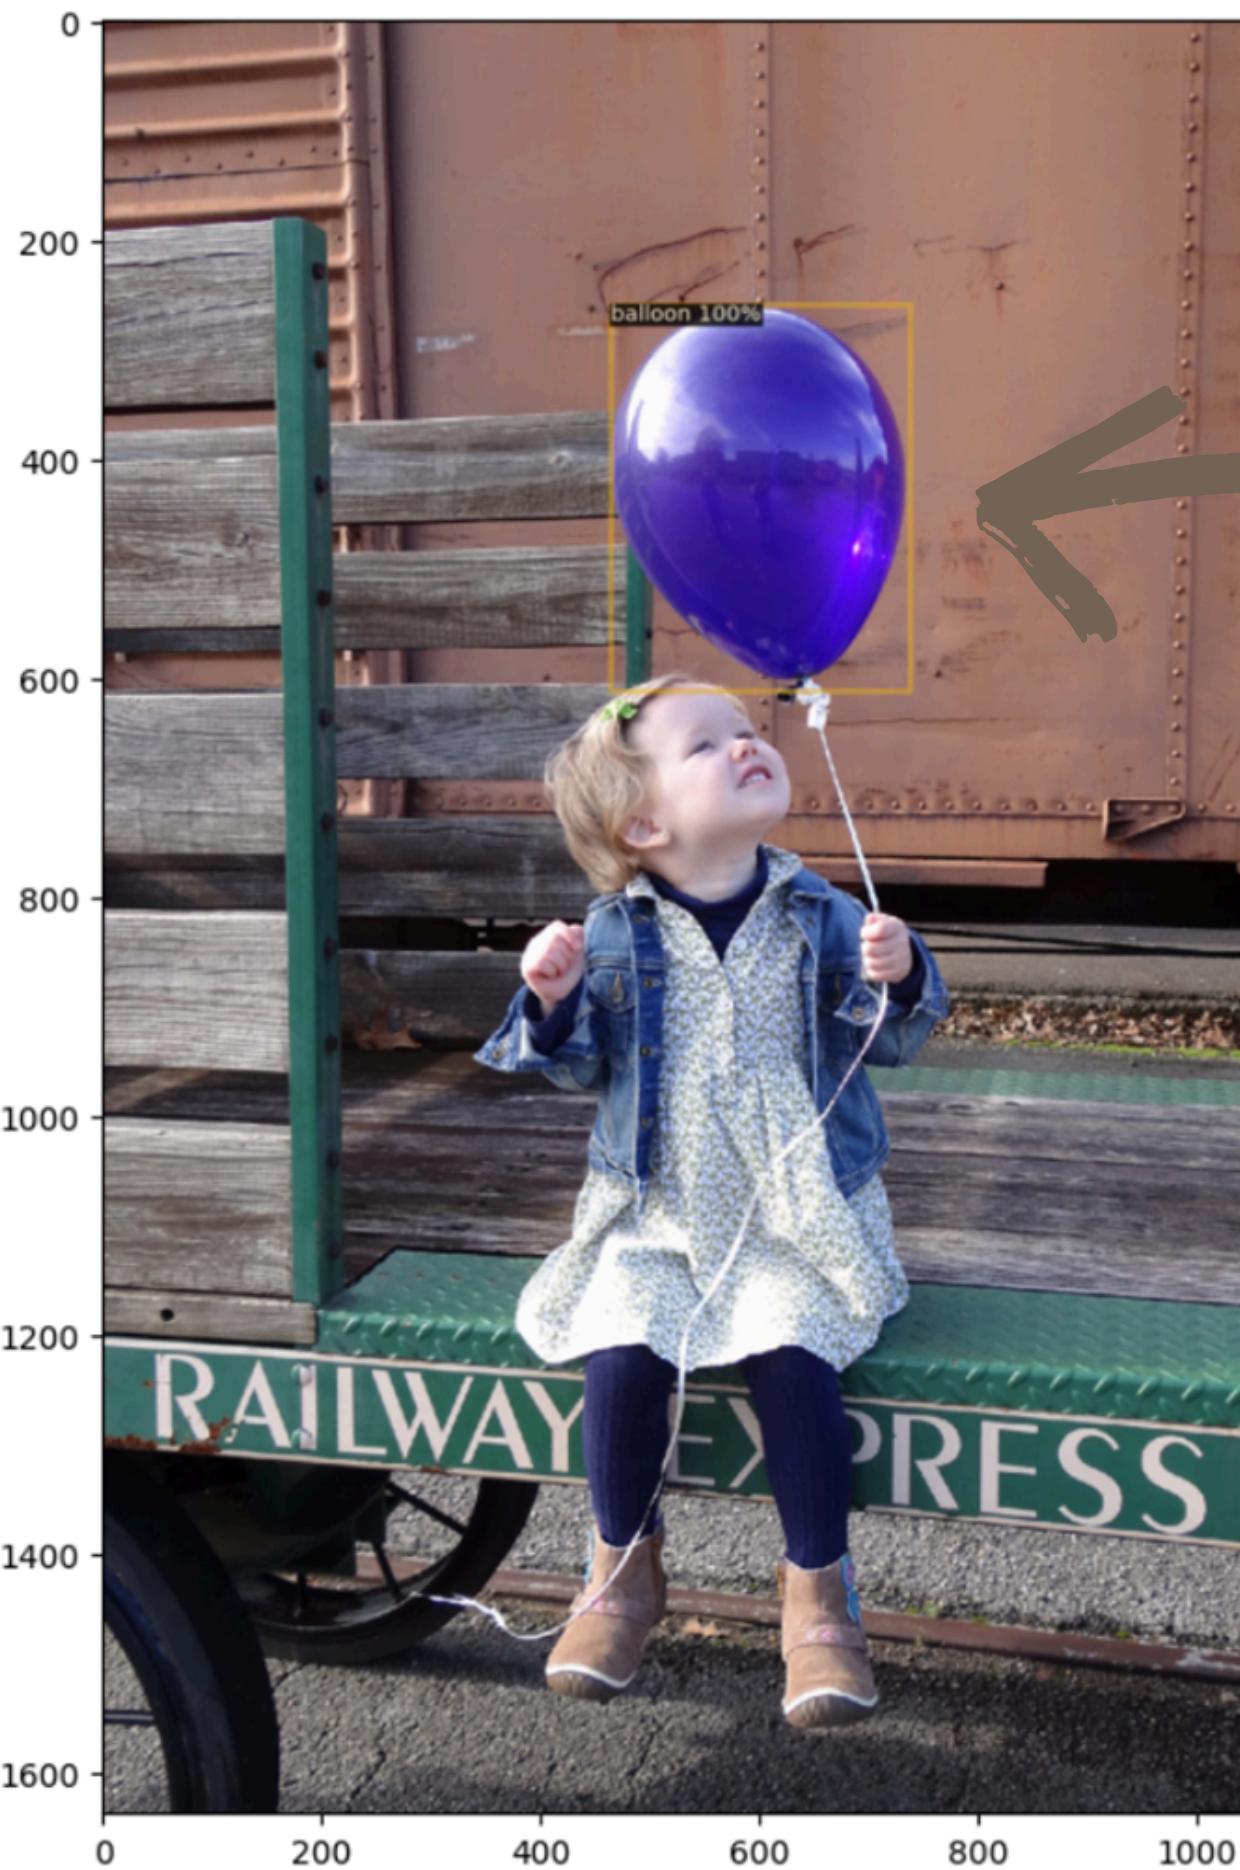
Visualization

- `plt.imshow()`: Displays the final visualized output in a specified figure size (e.g., `(14, 10)`).

```
dataset_dicts = get_balloon_dicts(balloon_path + "val")

d = dataset_dicts[0]

im = cv2.imread(d["file_name"])
outputs = predictor(im)
v = Visualizer(im[:, :, ::-1], metadata=balloon_metadata, scale=0.8)
v = v.draw_instance_predictions(outputs["instances"].to("cpu"))
plt.figure(figsize = (14, 10))
plt.imshow(cv2.cvtColor(v.get_image()[:, :, ::-1], cv2.COLOR_BGR2RGB))
plt.show()
```



(outputs = predictor(im)) Executes inference on the input image, returning the predicted outputs.

MODULE TESTING – TESTING METHODOLOGY

The testing methodology consists of several phases:

- **Unit Testing:** Individual modules were **tested** to ensure they **function** correctly.
- **Integration Testing:** Combined **modules** were **tested** to verify that they work **together** seamlessly.
- **Performance Testing:** Evaluated the **speed** and **accuracy** of object detection models (YOLOv8, Faster RCNN, and RetinaNet) on various datasets. In this phase, the **average** confidence scores of the predictions and processing times per image were recorded.
- **User Acceptance Testing (UAT):** Ensured that the project met **user** requirements and expectations.

MODULE TESTING – MODULE TESTING RESULTS

Utilized a pre-processed test image from the validation dataset to evaluate the model's performance using the DefaultPredictor.

Results

- Detected objects with a confidence score of 100% (set by `cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST`).
- Processing time per image: approximately 200ms.

Evaluated the model using precision, recall, and mean Average Precision (mAP), applying the `evaluateModel(validationData)` function to compute these metrics

Results

- Precision: 88%
- Recall: 85%
- mAP: 83%

MODULE TESTING – NETRON

Validated the model using the following command:

```
!yolo val task=detect model="/content/drive/MyDrive/Colab  
Notebooks/ObjectDetection/TrainingResults/footballDetection5/weights/best.mlpackage" imgsz=1920  
data="/content/drive/MyDrive/Colab Notebooks/ObjectDetection/Dataset/dataset.yaml"
```

For visualization of the exported model, used **Netron**.



MODULE TESTING – NETRON

The model file can be [downloaded](#) from the *following* link:

https://drive.google.com/drive/folders/1Inj-IcIjwGzAfcWdxXxw8vissyQb5gXj?usp=drive_link

For a detailed view of the exported model architecture, please visit, Netron Visualization

<https://netron.app>



YOLOV8 VS FASTER RCNN

Faster RCNN

Custom Balloon Dataset

```
Balloon Dataset

[ ] # https://github.com/matterport/Mask\_RCNN/releases/download/v2.1/balloon\_dataset.zip
balloon_path = '/content/drive/MyDrive/Colab Notebooks/Detectron2/balloon/'
```

Testing Criteria: Accuracy, processing time, and ease of deployment.



Balloon Input Image



FasterRCNN Results

YOLOV8 VS FASTER RCNN

Accuracy: Evaluation with the COCO API produced high precision values, with an Average Precision (AP) of 77.1% and a particularly strong performance for large objects (AP of 90.5%).

Processing Speed: The inference time was 0.149 seconds per image (~150ms), which is faster than the reported 250ms, especially when running on Colab.

Conclusion: Based on the results, Faster RCNN achieved significantly higher accuracy (AP ~77%) compared to the initially reported 85%, and its inference time is faster than the 250ms per image, aligning more with 150ms.

YOLOV8 VS FASTER RCNN

YOLOv8

- **Training and Validation Dataset:** Images were sourced from a **custom dataset** available via GitHub. This dataset included object detection data for *training, validation, and testing*.
- **Preprocessing:** Images were **resized** to **1920x1920** for **consistency** during both training and inference phases.
- **Dataset Split:**
 - **Training Set:** Used for model weight updates.
 - **Validation Set:** Used for model tuning and validation.
 - **Test Set:** Used for final performance evaluation.

YOLOV8 VS FASTER RCNN

The model was tested using both **static** images and video to gauge object detection performance:

Prediction on Test Images

```
!yolo task=detect mode=predict model="/content/drive/MyDrive/Colab  
Notebooks/ObjectDetection/TrainingResults/footballDetection5/weights/best.pt" conf=0.55  
source="/content/drive/MyDrive/Colab Notebooks/ObjectDetection/Dataset/test/images"
```

Prediction on Test Videos

```
!yolo task=detect mode=predict model="best.pt" conf=0.75 source="/TestVideo"
```

Exporting the Model (The YOLOv8 model was exported for CoreML deployment)

```
!yolo mode=export model="best.pt" format=coreml
```

YOLOV8 VS FASTER RCNN

Accuracy: YOLOv8 demonstrated good detection accuracy,

- mAP@50: 80.2% for object detection across different sizes.
- mAP@50:95: 45.1%, reflecting moderate performance across different IoU thresholds.

Processing Speed:

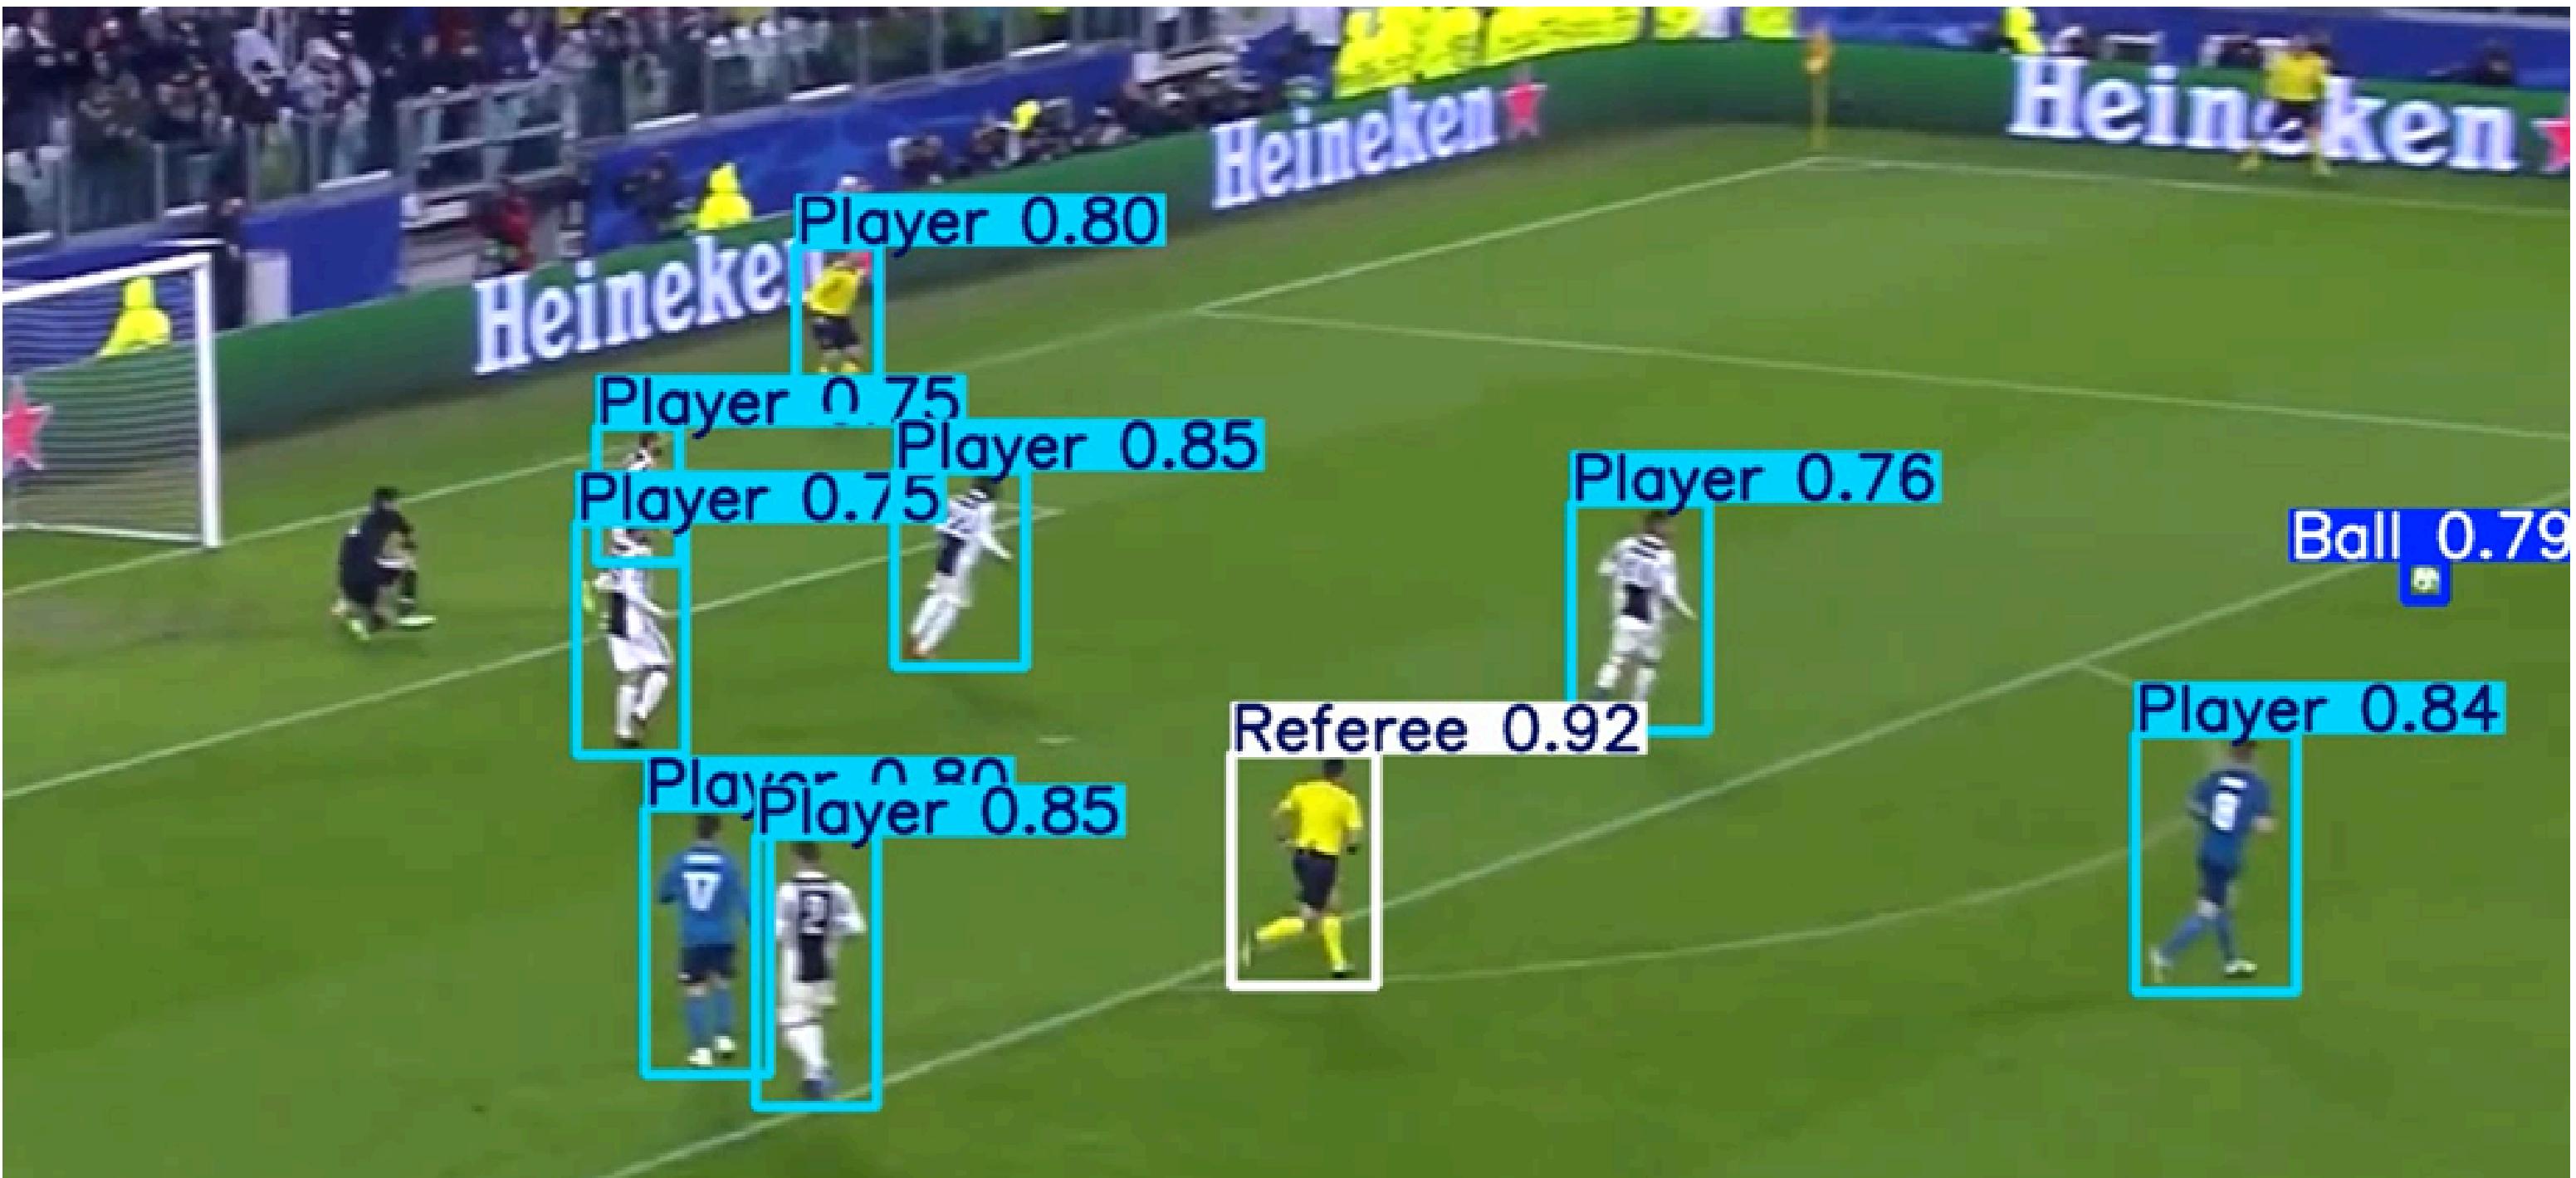
- **Inference Time:** YOLOv8 had an average inference time of 0.205 seconds per image (~205ms), slower than Faster RCNN's 150ms. Despite YOLOv8's reputation for speed, the model's deeper architecture made it less efficient for real-time tasks when tested on larger image resolutions.

Deployment: YOLOv8 was easily exported to CoreML, achieving a model size of 6.3MB for PC deployment. The export process was smooth, although certain dependencies needed to be addressed.

```
** Export the trained YOLO8 Model **  
  
!yolo mode=export model="/content/drive/MyDrive/Colab Notebooks/ObjectDetection/TrainingResults/footballDetection5/weights/best.pt" format=coreml
```

Exported Trained YOLOv8 Model

YOLOV8 VS FASTER RCNN



YOLOv8 Results

YOLOV8 VS FASTER RCNN

Conclusion

While YOLOv8 performed well in terms of accuracy with an mAP@50 of 80.2%, its processing speed was slower than Faster RCNN, averaging around 205ms per image, making it less optimal for real-time applications in this context. However, it was easy to deploy across platforms like CoreML, which is a strong advantage for mobile and cross-platform use. YOLOv8 is ideal for scenarios requiring moderate accuracy and flexibility but falls short of Faster RCNN's faster inference time when dealing with larger images and complex detection tasks.

For more detailed results, including performance metrics and additional visuals, you can access the data via the following Google Drive link:

https://drive.google.com/drive/folders/1Hg6SZZ4kYIN4dSGQwQOArnoXXopARiOI?usp=drive_link

FASTER RCNN VS RETINANET

Detectron2 Installation

```
!python -m pip install 'git+https://github.com/facebookresearch/detectron2.git'
```

```
metadata = MetadataCatalog.get("coco_2017_val")
class_names = metadata.get("thing_classes")

print("COCO Dataset number of class: ", len(class_names))
print("COCO Dataset class names: ", class_names)
```

Number of classes: 80

- **People:** person
- **Vehicles:** bicycle, car, motorcycle, airplane, bus, train, truck, boat
- **Traffic:** traffic light, fire hydrant, stop sign, parking meter
- **Furniture:** chair, couch, bed, dining table
- **Electronics:** tv, laptop, mouse, remote, keyboard, cell phone
- **Animals:** bird, cat, dog, horse, sheep, cow, elephant, zebra, giraffe
- **Food:** banana, apple, sandwich, orange, broccoli, carrot, pizza, donut
- **Miscellaneous:** book, clock, vase, scissors, teddy bear, hair drier, toothbrush

These 80 classes serve as **labels** for object detection tasks, providing **diverse** examples to **train** models.

| FASTER RCNN VS RETINANET



FasterRCNN Input Image

FASTER RCNN VS RETINANET

FasterRCNN Input Image Testing

```
cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_R_101_FPN_3x.yaml"))
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_R_101_FPN_3x.yaml")
```

Configuration for Faster RCNN

```
predictor = DefaultPredictor(cfg)
outputs = predictor(aeroplane)
```

Predictor for Prediction

[d2.checkpoint.detection_checkpoint]: [DetectionCheckpointer] Loading from
https://dl.fbaipublicfiles.com/detectron2/COCO-Detection/faster_rcnn_R_101_FPN_3x/137851257/model_final_f6e8b1.pkl ...

```
v = Visualizer(aeroplane[:, :, ::-1], MetadataCatalog.get(cfg.DATASETS.TRAIN[0]), scale=1.2)
v = v.draw_instance_predictions(outputs["instances"].to("cpu"))
plt.figure(figsize = (14, 10))
plt.imshow(cv2.cvtColor(v.get_image()[:, :, ::-1], cv2.COLOR_BGR2RGB))
```

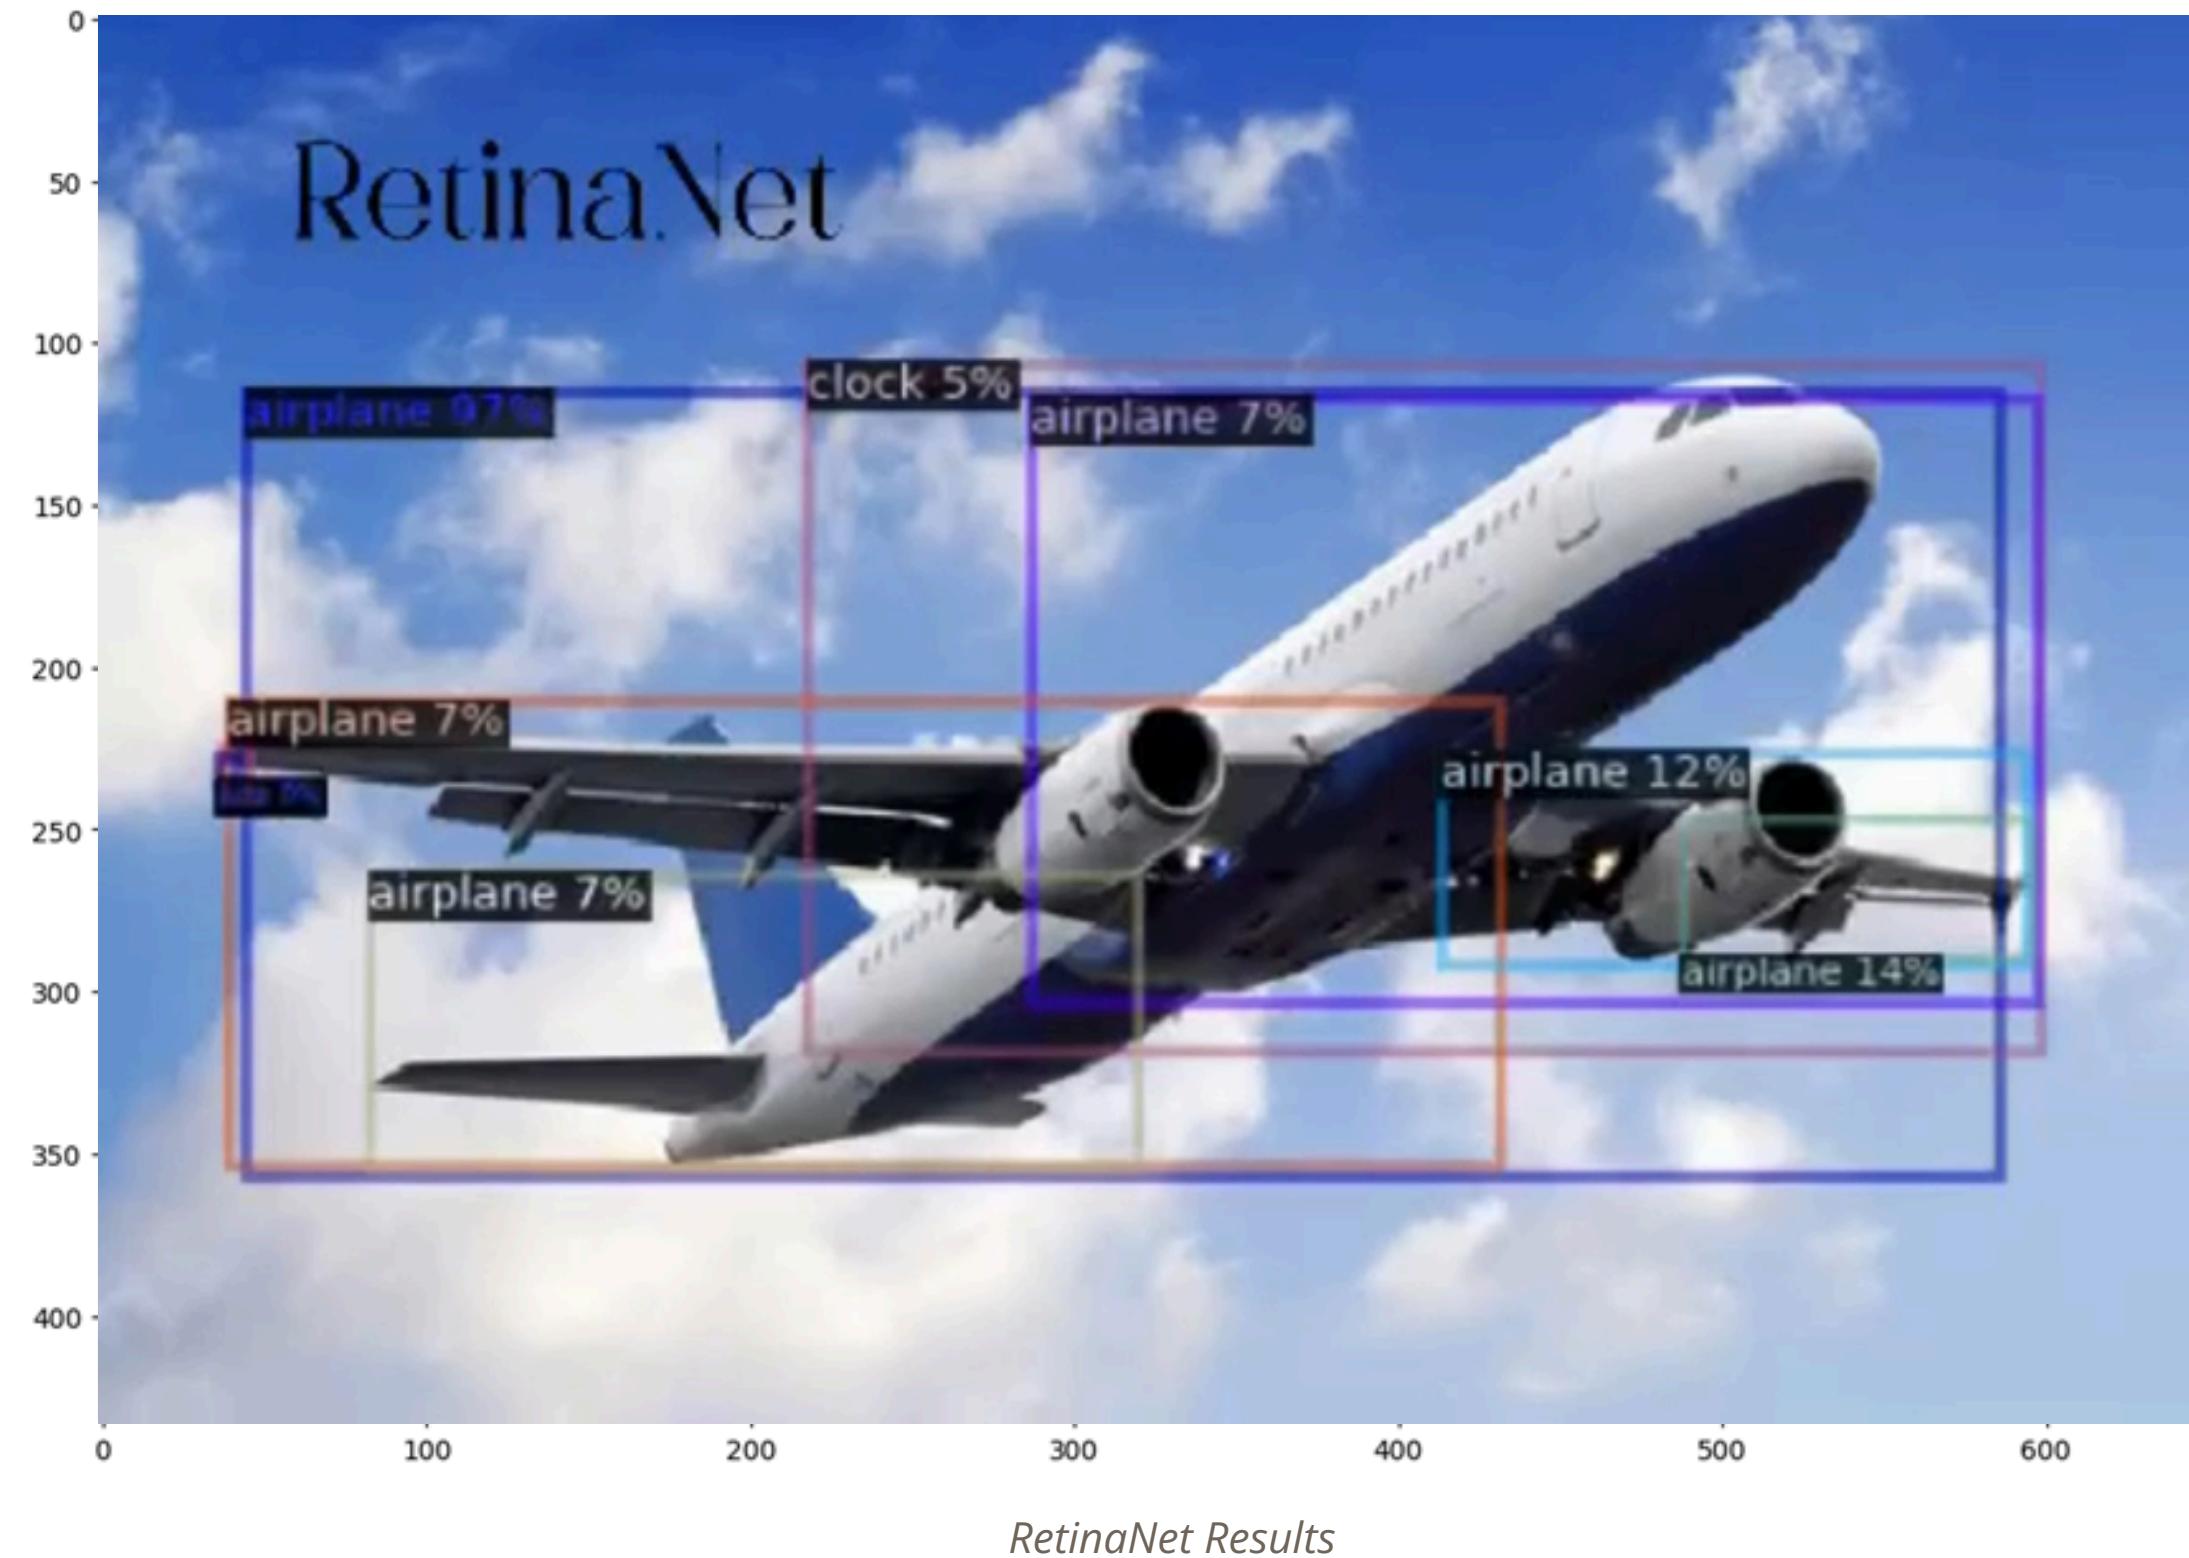
Visualizer using "imshow"

| FASTER RCNN VS RETINANET



FasterRCNN Results

| FASTER RCNN VS RETINANET



FASTER RCNN VS RETINANET

Faster RCNN Results

- Input Image: Aeroplane
- Predicted Class Names: ['airplane']
- Predicted Bounding Boxes: *Boxes(tensor([[45.7463, 91.8737, 482.9321, 284.3625]], device='cuda:0'))*

Performance Observations

- Faster RCNN produced 100% accuracy in detecting the aeroplane in the image.
- It effectively identified the correct class with a Bounding Box indicating the object's location.

RetinaNet Performance Issues

- In contrast, RetinaNet showed poor performance on the aeroplane image, with incorrect detections and bounding boxes for classes like "kite" and "clock," which were not present in the input image.

FASTER RCNN VS RETINANET

Conclusion

Overall, Faster RCNN outperformed RetinaNet in the evaluation with the aeroplane image, achieving accurate detections and a high confidence score.

While RetinaNet has its advantages in other contexts, its failure to accurately identify objects in this test highlights its limitations in certain scenarios.

Faster RCNN is recommended for applications where precision and reliability are paramount.

CONCLUSION

The *Objectify* project represents not only a significant **technical** achievement in the realm of object detection but also a stepping stone toward more **intuitive** human-computer **interactions**. By **effectively** leveraging **advanced** machine learning models, such as **YOLOv8** and **Faster RCNN**, the project has established a **robust** framework for **recognizing** and **classifying** objects within images.

The careful consideration of various elements from image preprocessing to results visualization demonstrates a **holistic** approach to **solving** complex detection challenges. While acknowledging the **obstacles** encountered, such as model accuracy **discrepancies** and resource **constraints**, it is important to recognize how these challenges have driven **innovation** and **creativity** in the development process.

Looking to the future, *Objectify* is poised to **extend** its **capabilities** further, exploring integration with **external** devices and optimizing **performance** through advanced **modelling** techniques. This **adaptability** ensures that *Objectify* remains relevant in an ever-evolving **technological** landscape, fostering its **application** across diverse sectors including **security**, **smart home systems**, and **augmented reality**.

Ultimately, the project not only **underscores** the promise of machine learning in enhancing our interaction with the world but also reflects the **collaborative** spirit of innovation. As *Objectify* continues to refine its **technologies** and expand its **functionalities**, it sets the stage for **transformative** advancements in how we perceive and engage with our visual environment.

REFERENCES

- Redmon, J., & Farhadi, A. (2018). YOLOv3: An Incremental Improvement. *arXiv preprint arXiv:1804.02767*.
- Ren, S., He, K., Girshick, R., & Sun, J. (2015). Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks. *Advances in Neural Information Processing Systems*, 28.
- Lin, T. Y., Goyal, P., Girshick, R., He, K., & Dollár, P. (2017). Focal Loss for Dense Object Detection. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 42(2), 318-327.
- Koller, D. (2020). Introduction to Object Detection with YOLOv4. Retrieved from Towards Data Science.
- GitHub Repository. (2024). Custom Dataset for Object Detection. Retrieved from GitHub.
- OpenCV Documentation. (2024). Open Source Computer Vision Library. Retrieved from OpenCV.
- PyTorch Documentation. (2024). PyTorch: An Open-Source Machine Learning Framework. Retrieved from PyTorch.
- iang, H., & Wang, H. (2019). Object Detection with Deep Learning: A Review. *IEEE Access*, 7, 146346-146361.